

# Towards First Class References as a Security Infrastructure in Dynamically-Typed Languages

Jean-Baptiste Arnaud

Supervisor: Stéphane Ducasse  
Co-Supervisor: Marcus Denker  
RMoD team



February 18, 2013

# Dynamically-typed languages



*JavaScript*



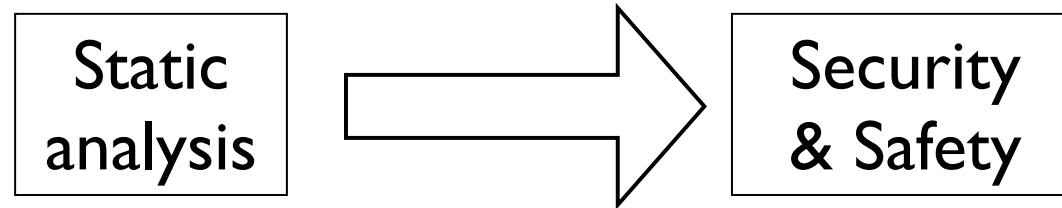
**Ruby**  
*A Programmer's Best Friend*

Dynamically-typed languages are becoming popular

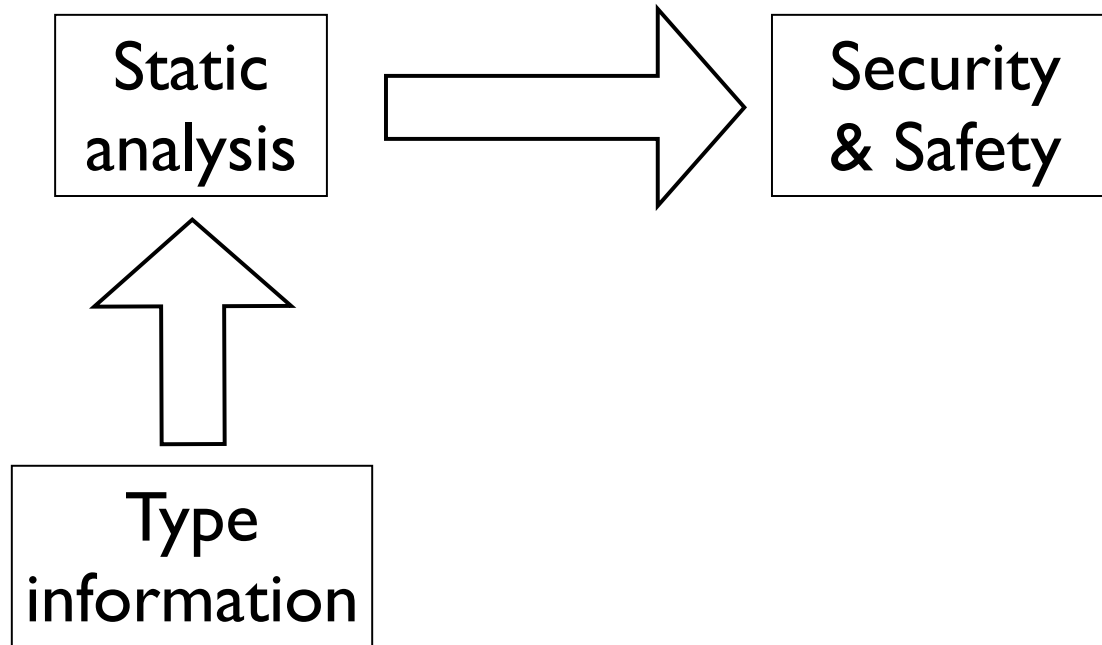
Need

Security & Safety

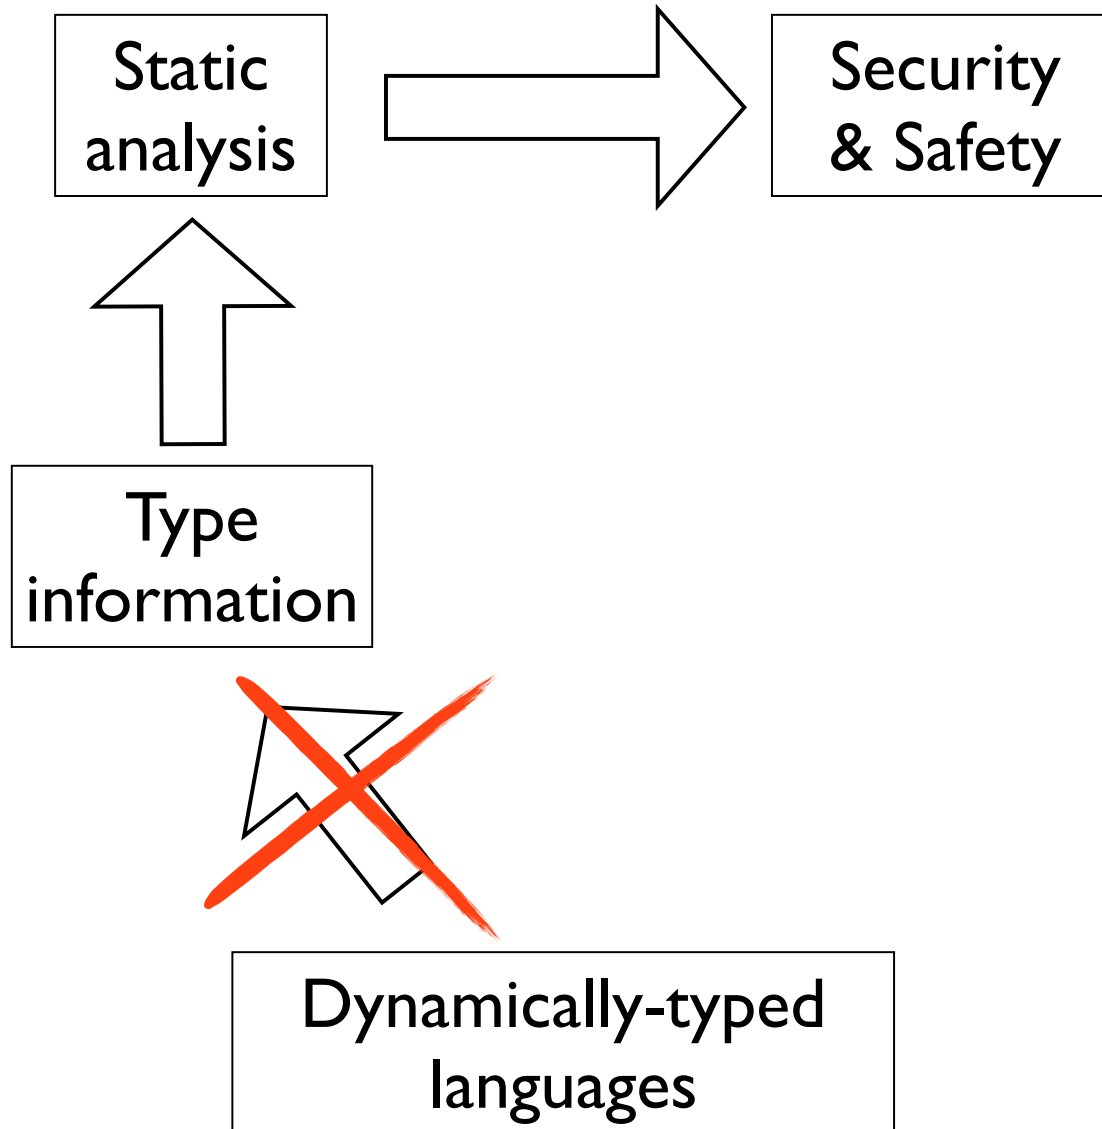
# Security and Safety



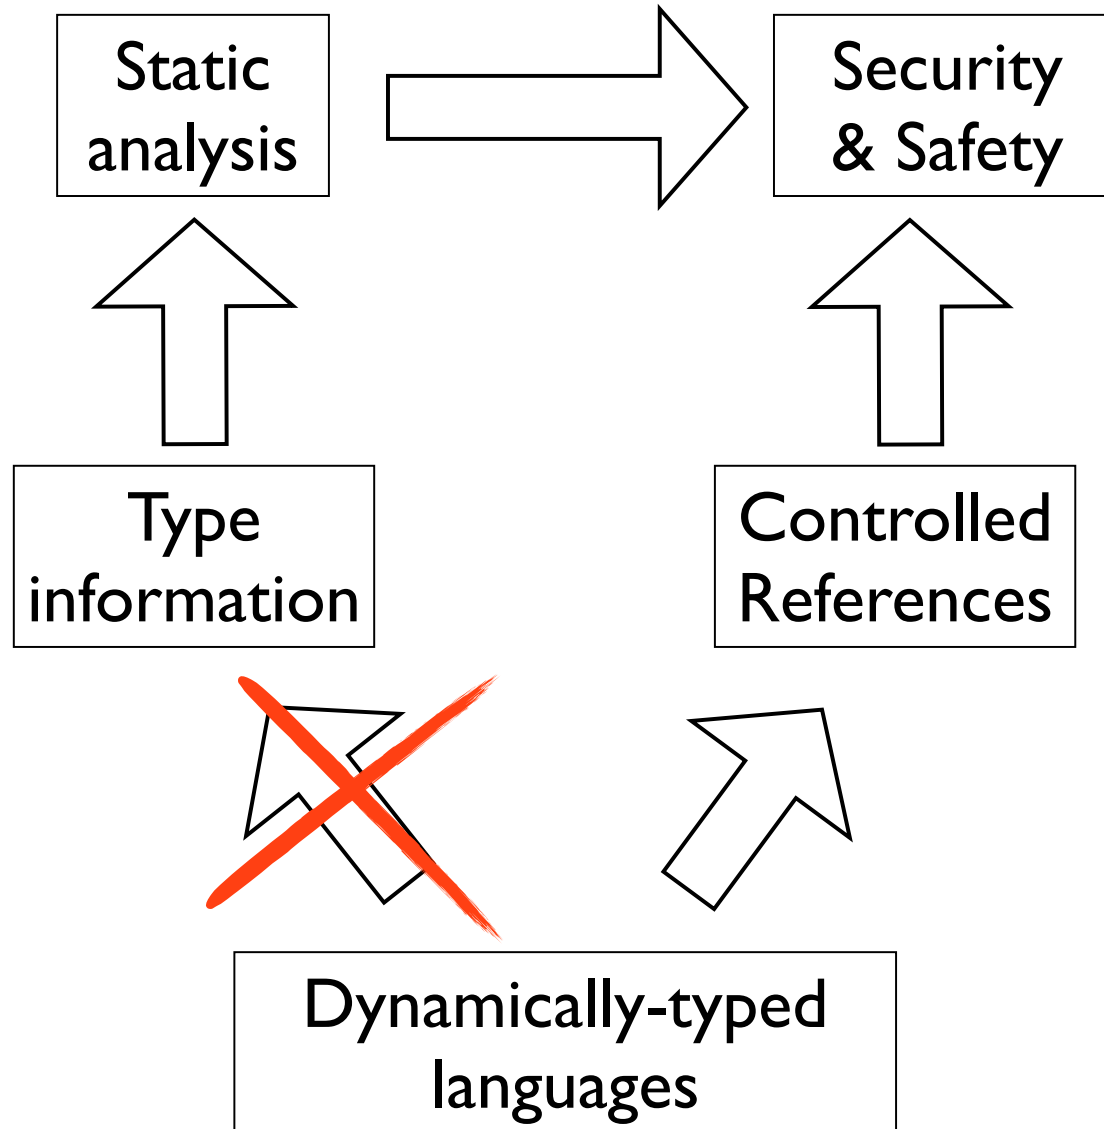
# Security and Safety



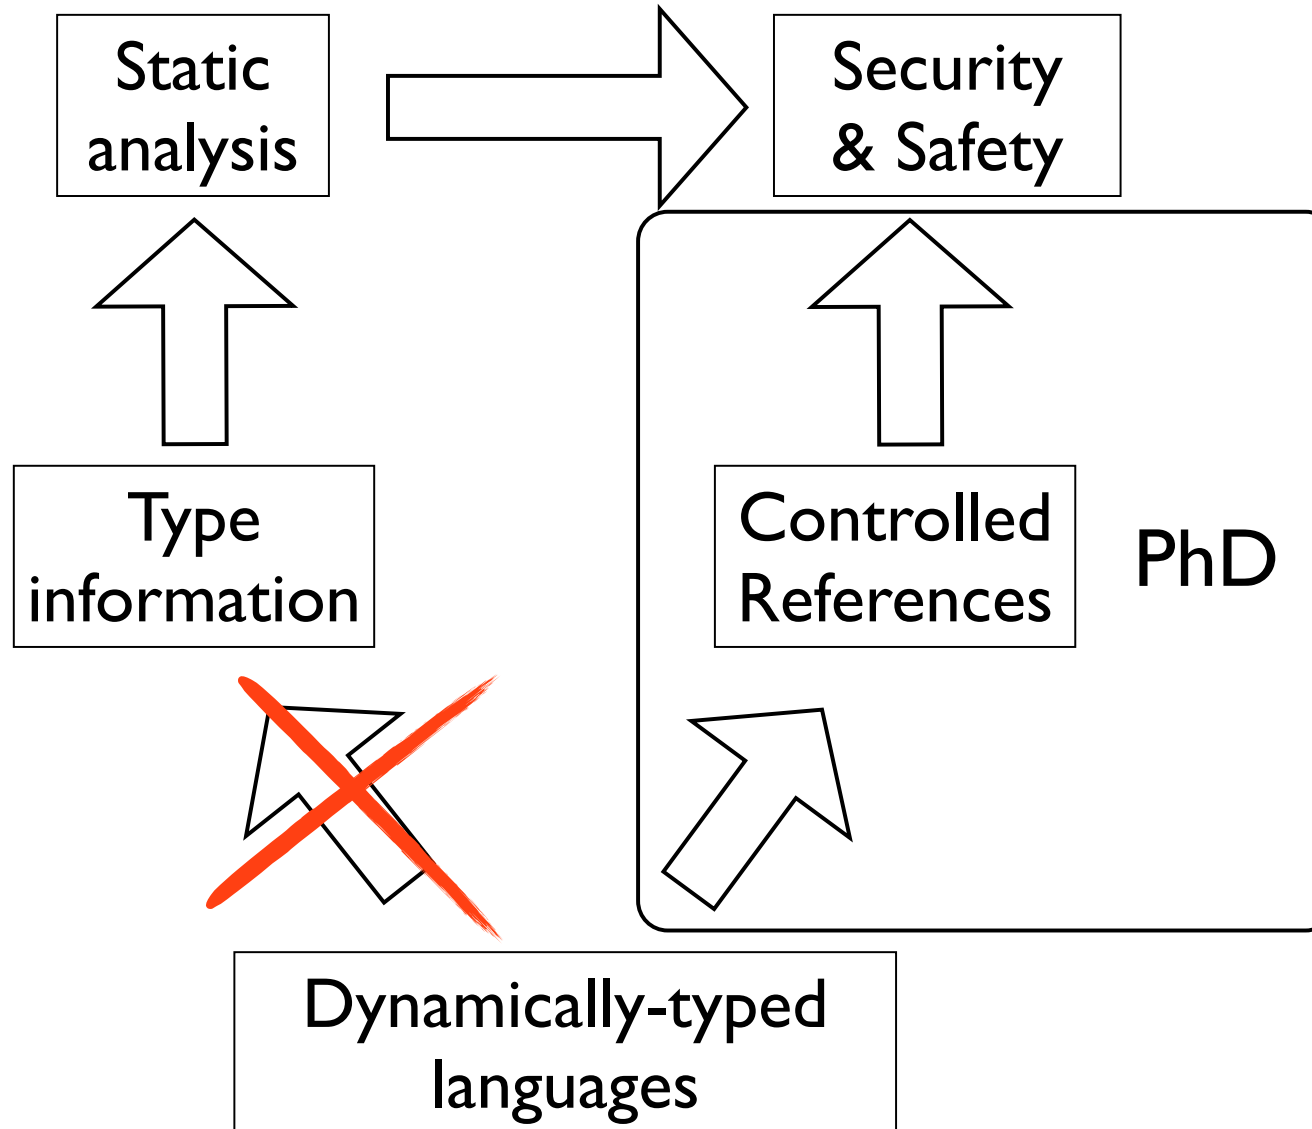
# Security and Safety



# Security and Safety



# Security and Safety



# Problem

How can we provide a *general-purpose* way to control references in dynamically-typed languages?



# Outline

I. Motivation

II. Approaches

III. Solution

IV. Consequences

V. Usages

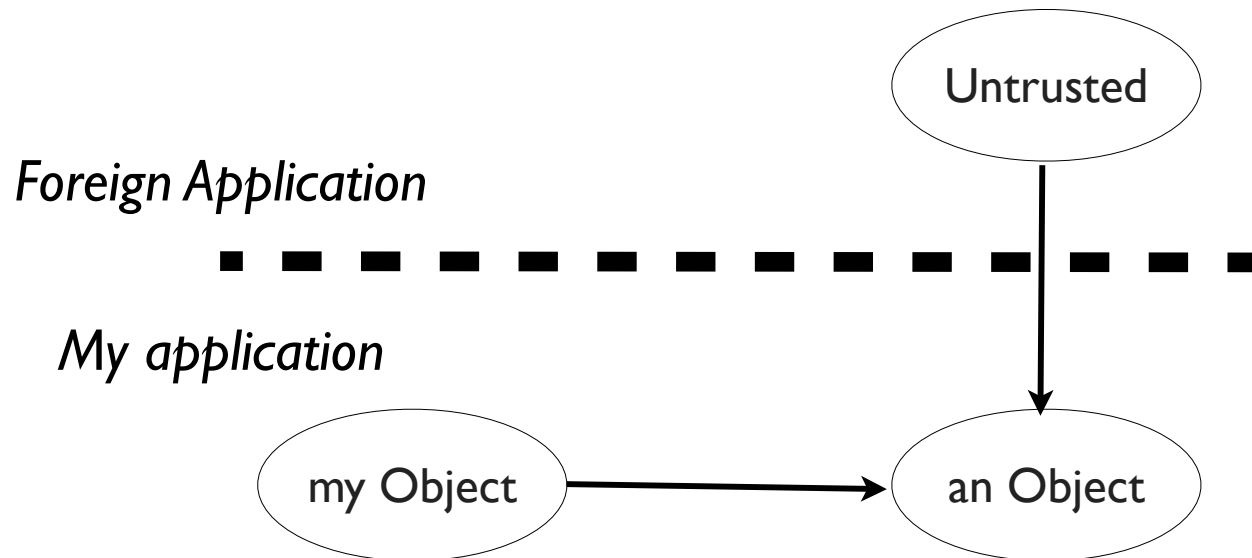
VI. Validation

VII. Conclusion

# Use cases for controlled references

## Side effect management:

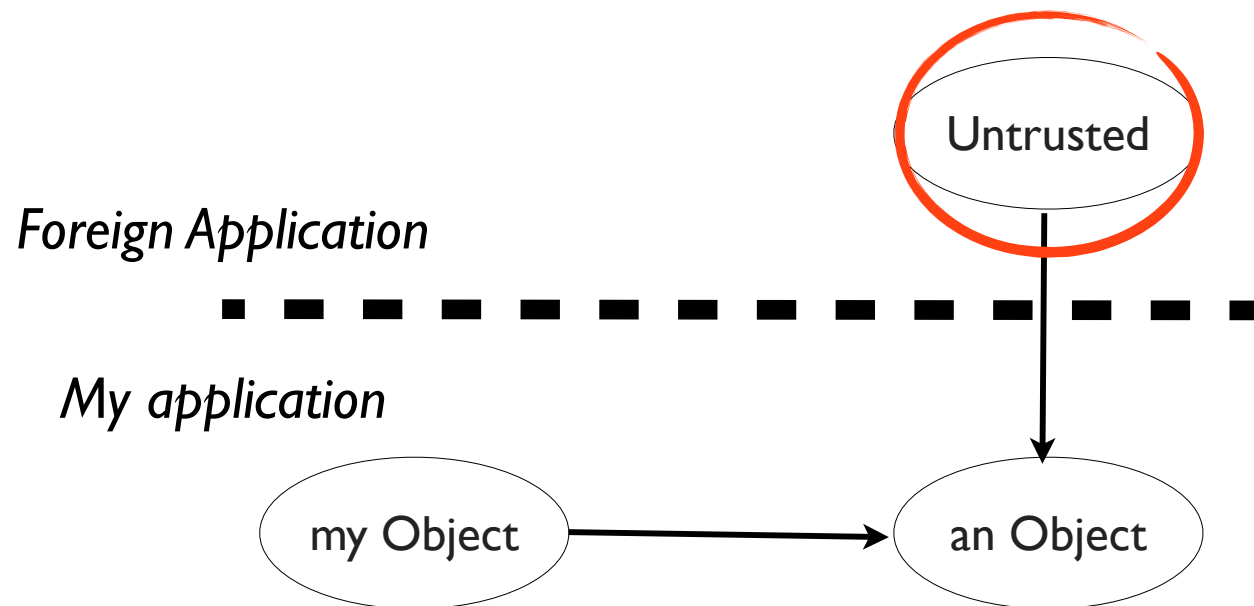
- Prevent side effect from occurring
- Accept side effect and check them



# Use cases for controlled references

## Side effect management:

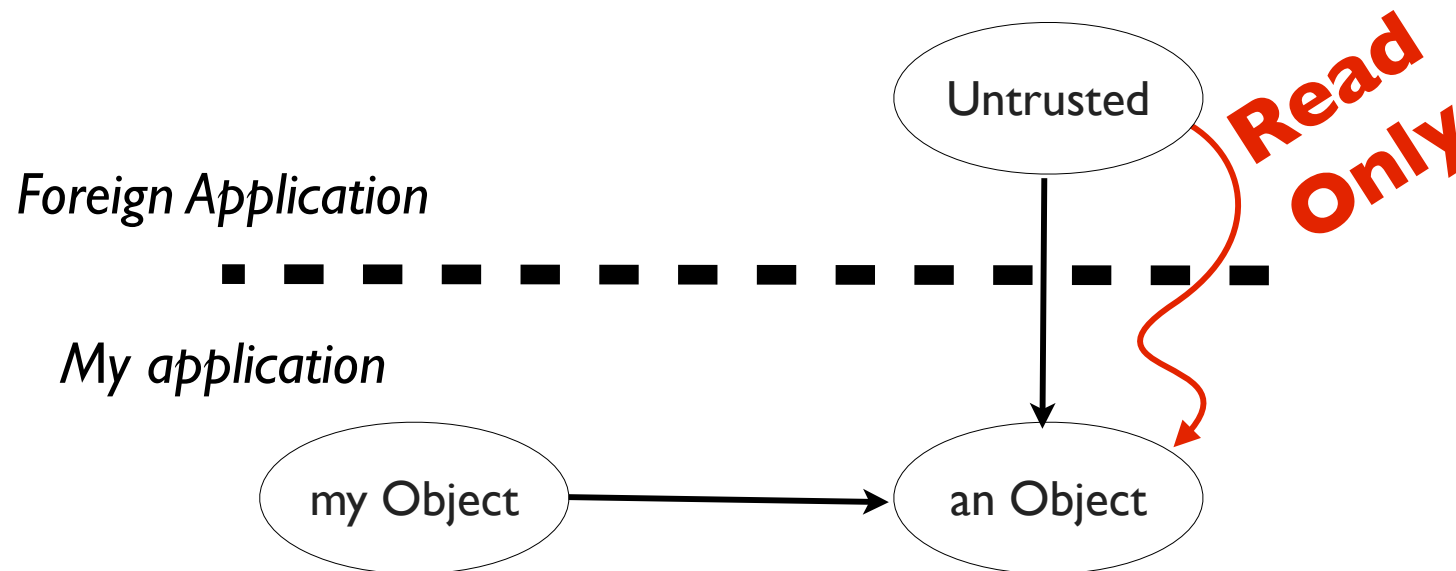
- Prevent side effect from occurring
- Accept side effect and check them



# Use cases for controlled references

## Side effect management:

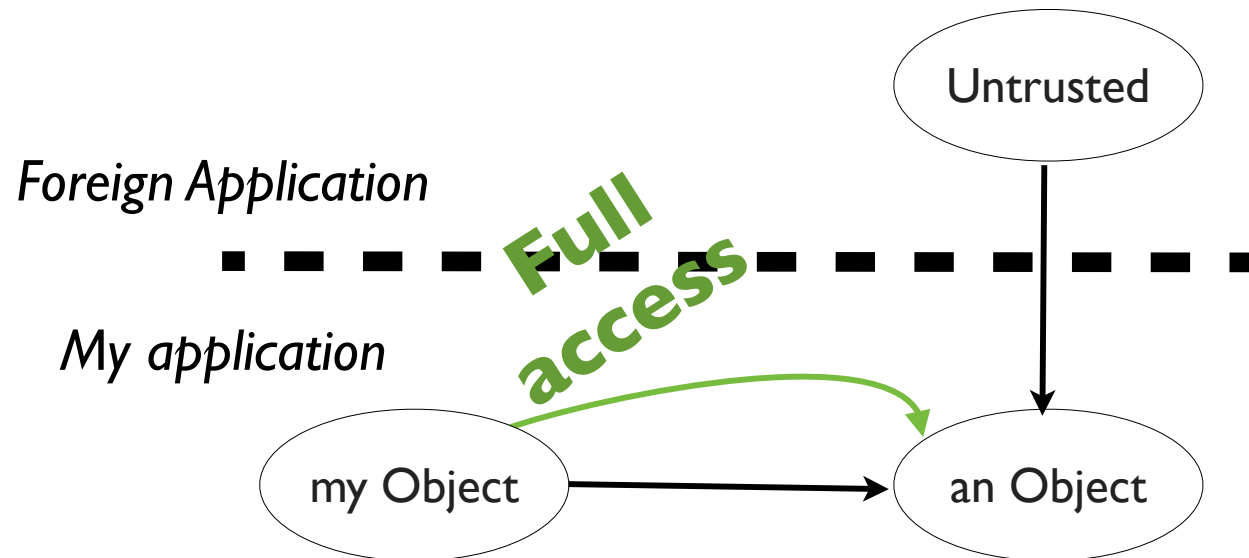
- Prevent side effect from occurring: **ReadOnly Execution**
- Accept side effect and check them



# Use cases for controlled references

## Side effect management:

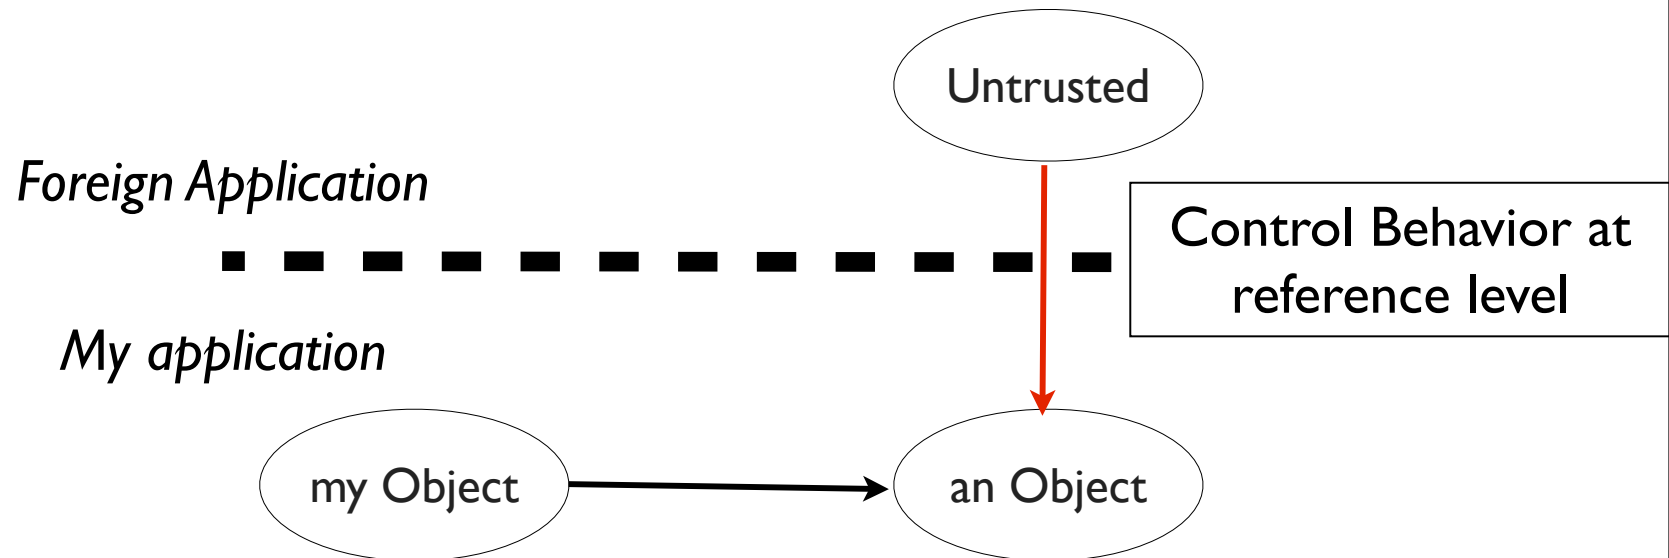
- Prevent side effect from occurring: ReadOnly Execution
- Accept side effect and check them



# Use cases for controlled references

## Side effect management:


- Prevent side effect from occurring: ReadOnly Execution
- Accept side effect and check them



# Use cases for controlled references

## Side effect management:

- Prevent side effect from occurring
- Accept side effect and check them

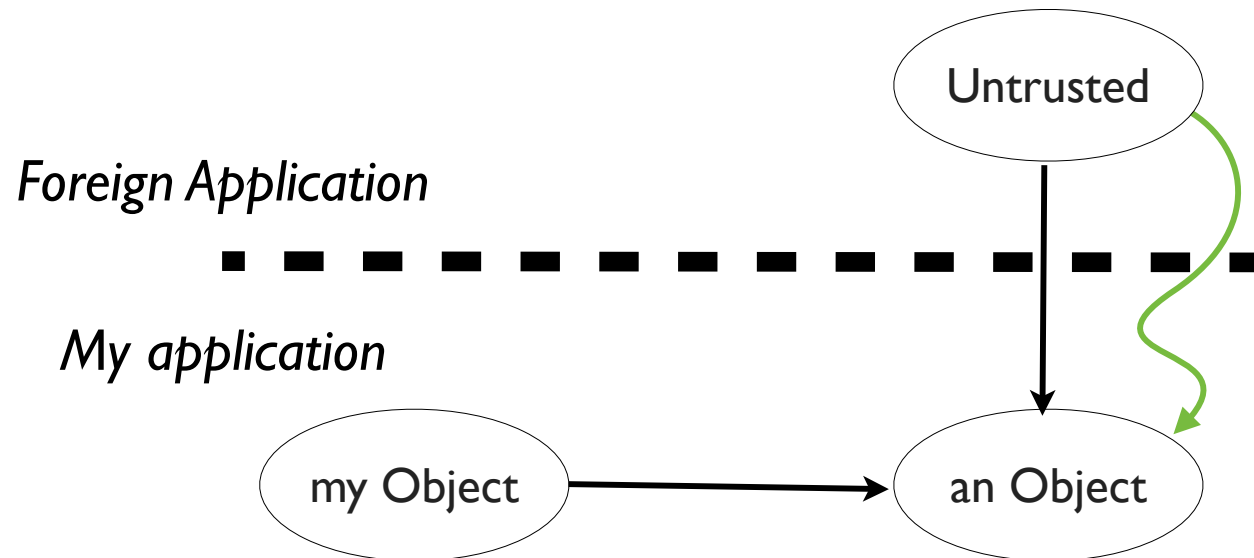


Controlled references ?  
~~X~~ Controlled behavior

# Use cases for controlled references

## Side effect management:

- Prevent side effect from occurring
- Accept side effect and check them: Isolated Execution

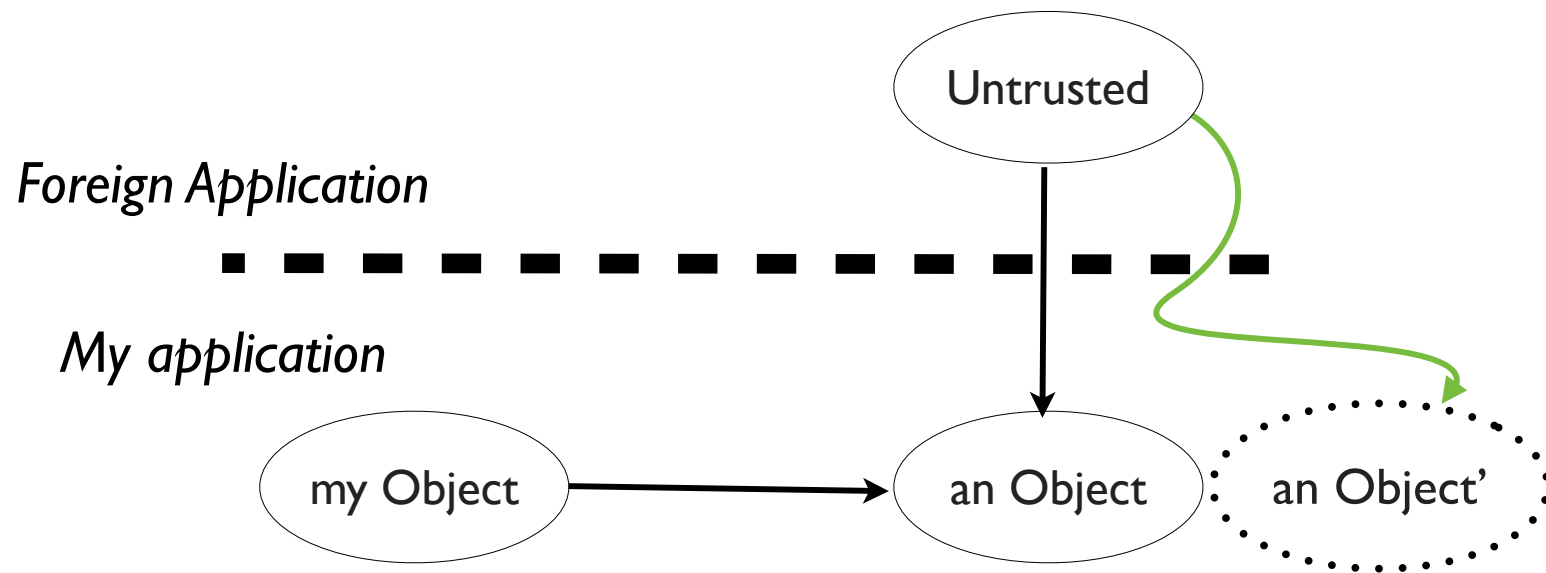




# Use cases for controlled references

## Side effect management:

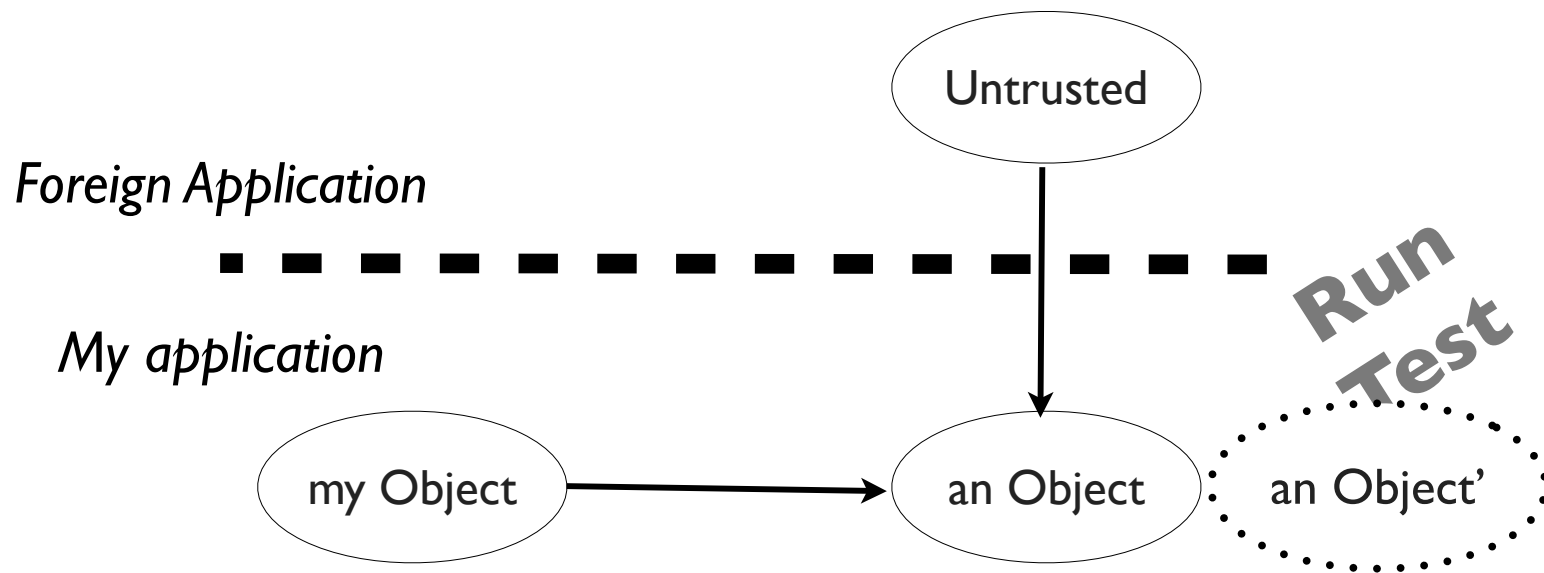
- Prevent side effect from occurring
- Accept side effect and check them: Isolated Execution



# Use cases for controlled references

## Side effect management:

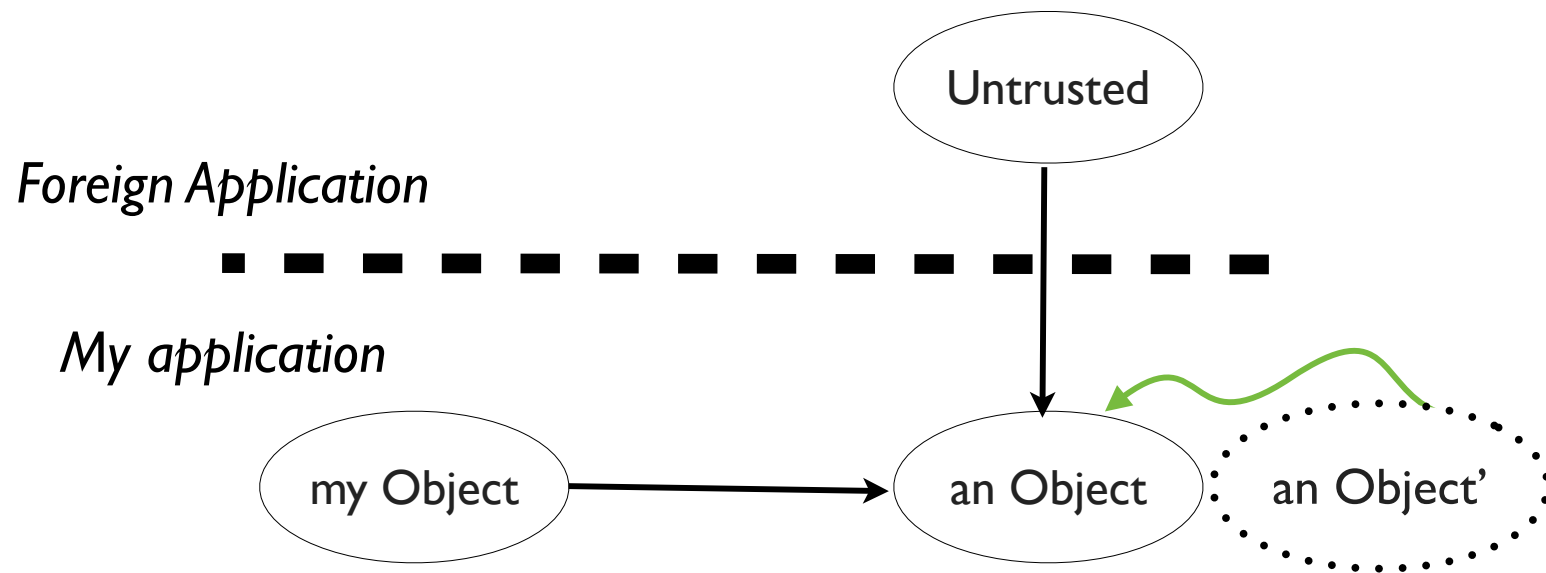
- Prevent side effect from occurring
- Accept side effect and check them: Isolated Execution



# Use cases for controlled references

## Side effect management:

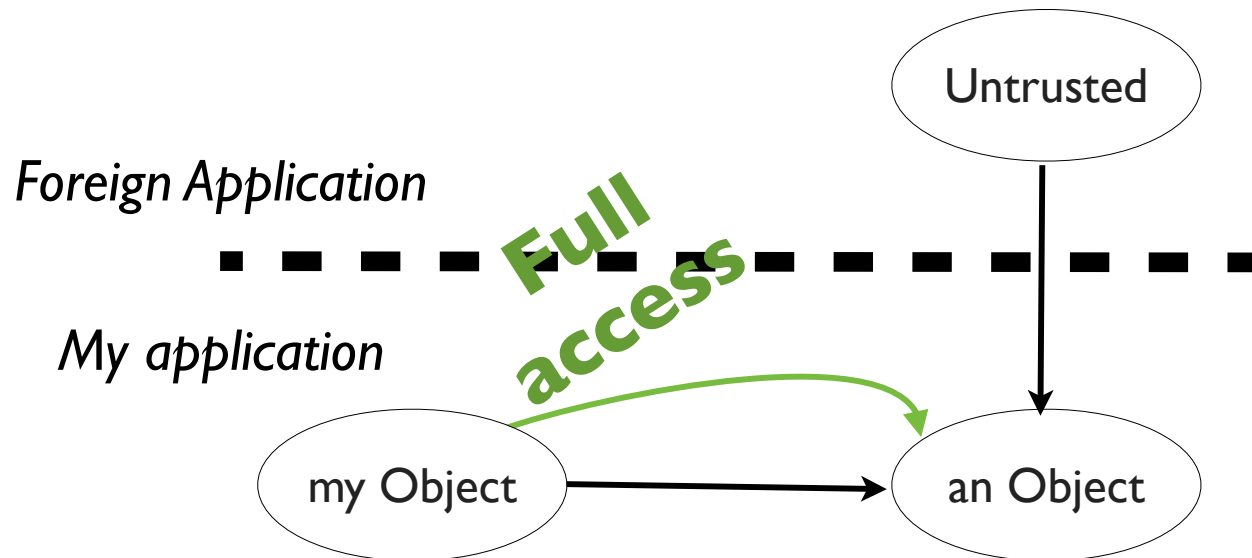
- Prevent side effect from occurring
- Accept side effect and check them: Isolated Execution



# Use cases for controlled references

## Side effect management:

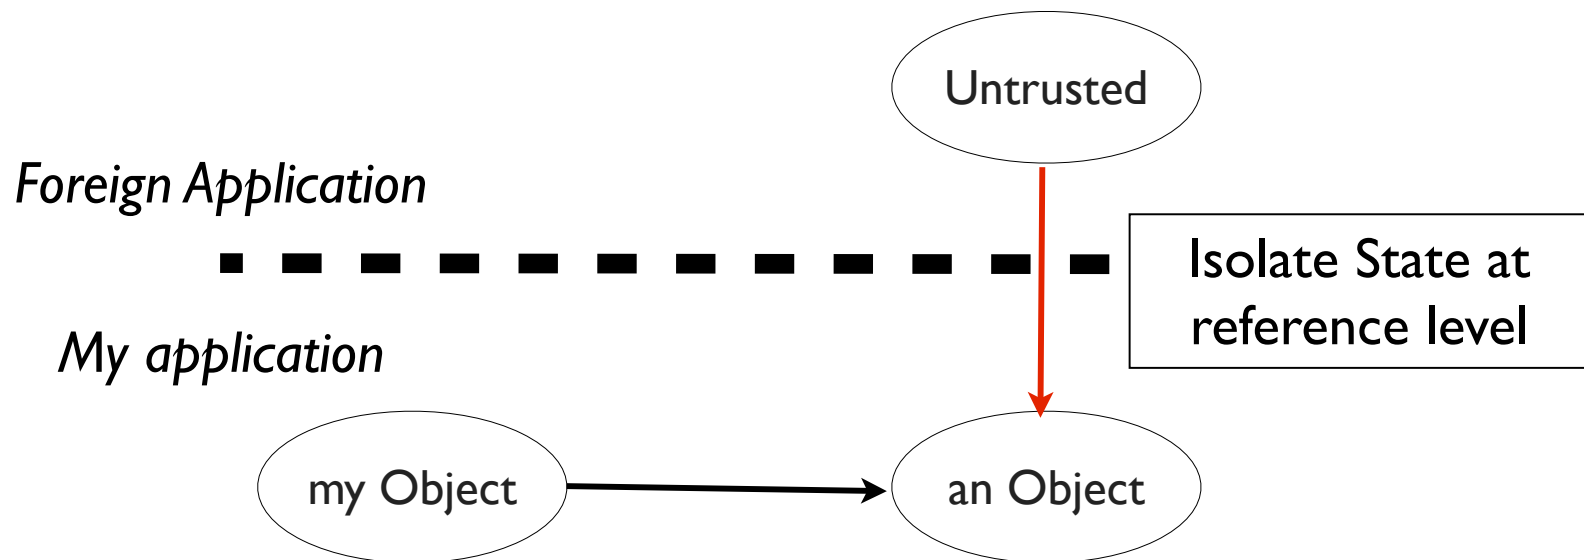
- Prevent side effect from occurring
- Accept side effect and check them: Isolated Execution



# Use cases for controlled references

## Side effect management:

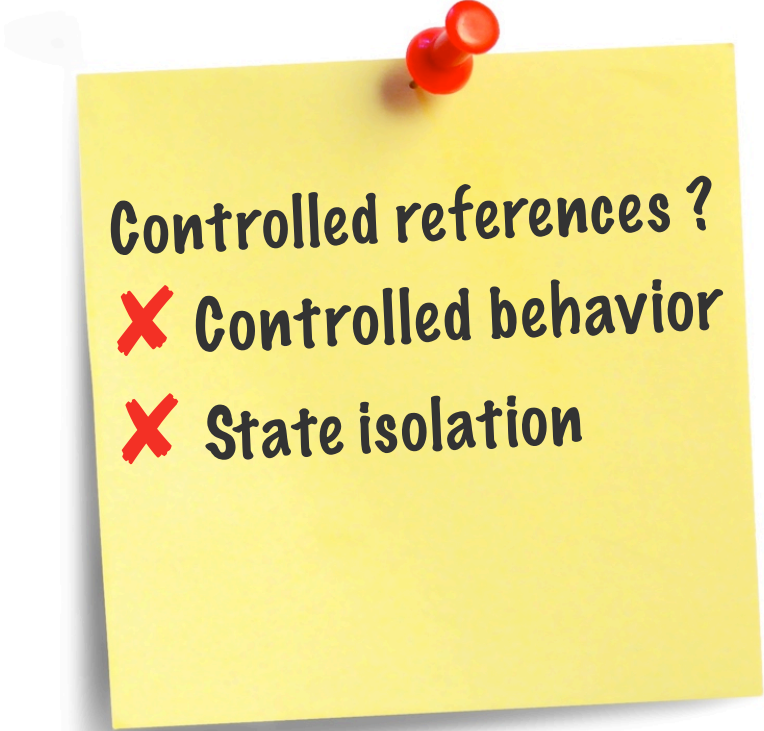
- Prevent side effect from occurring
- Accept side effect and check them: Isolated Execution



# Use cases for controlled references

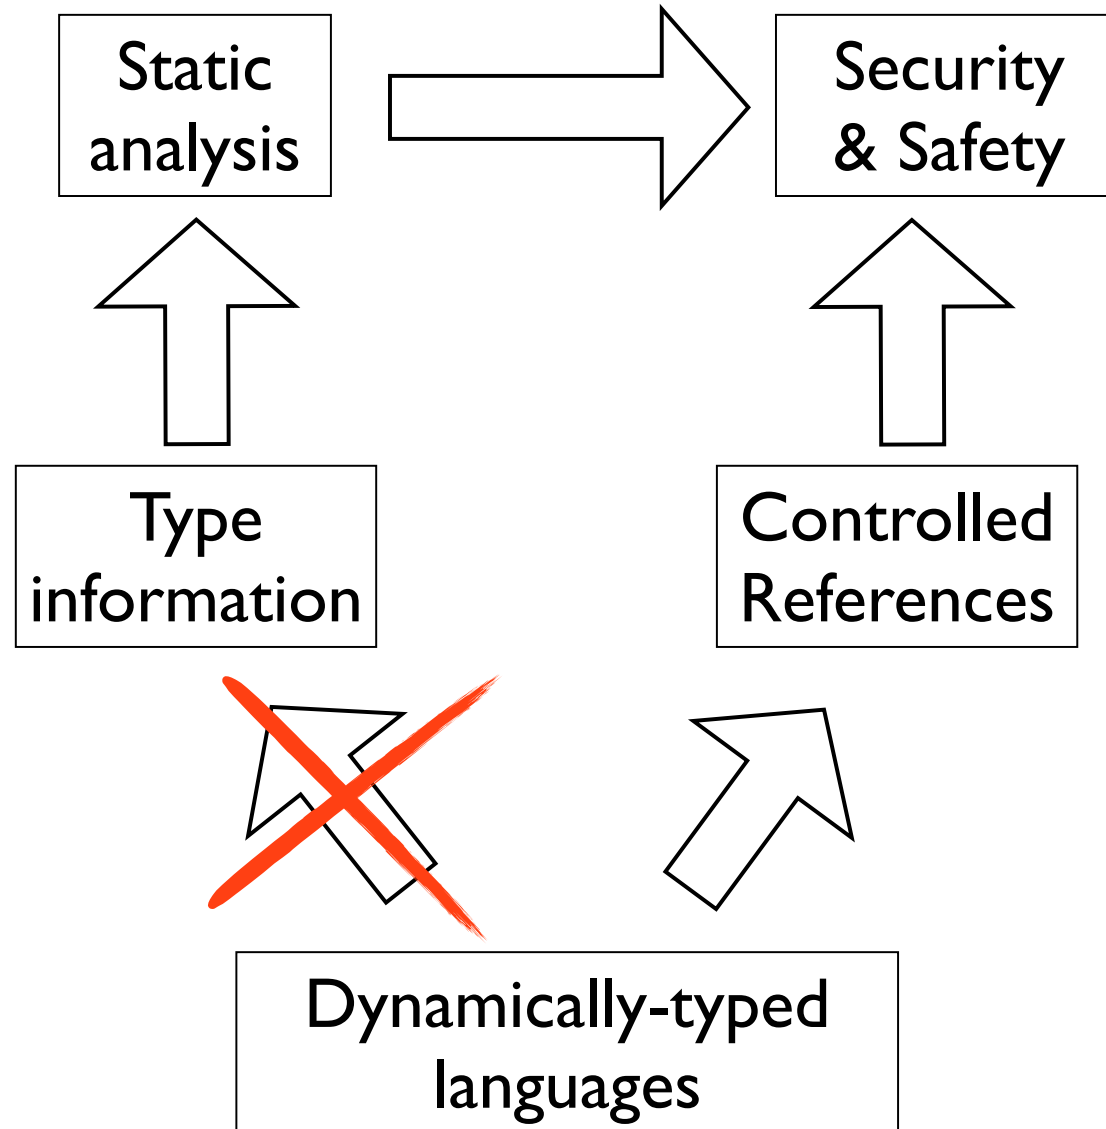
## Side effect management:

- Prevent side effect from occurring
- Accept side effect and check them

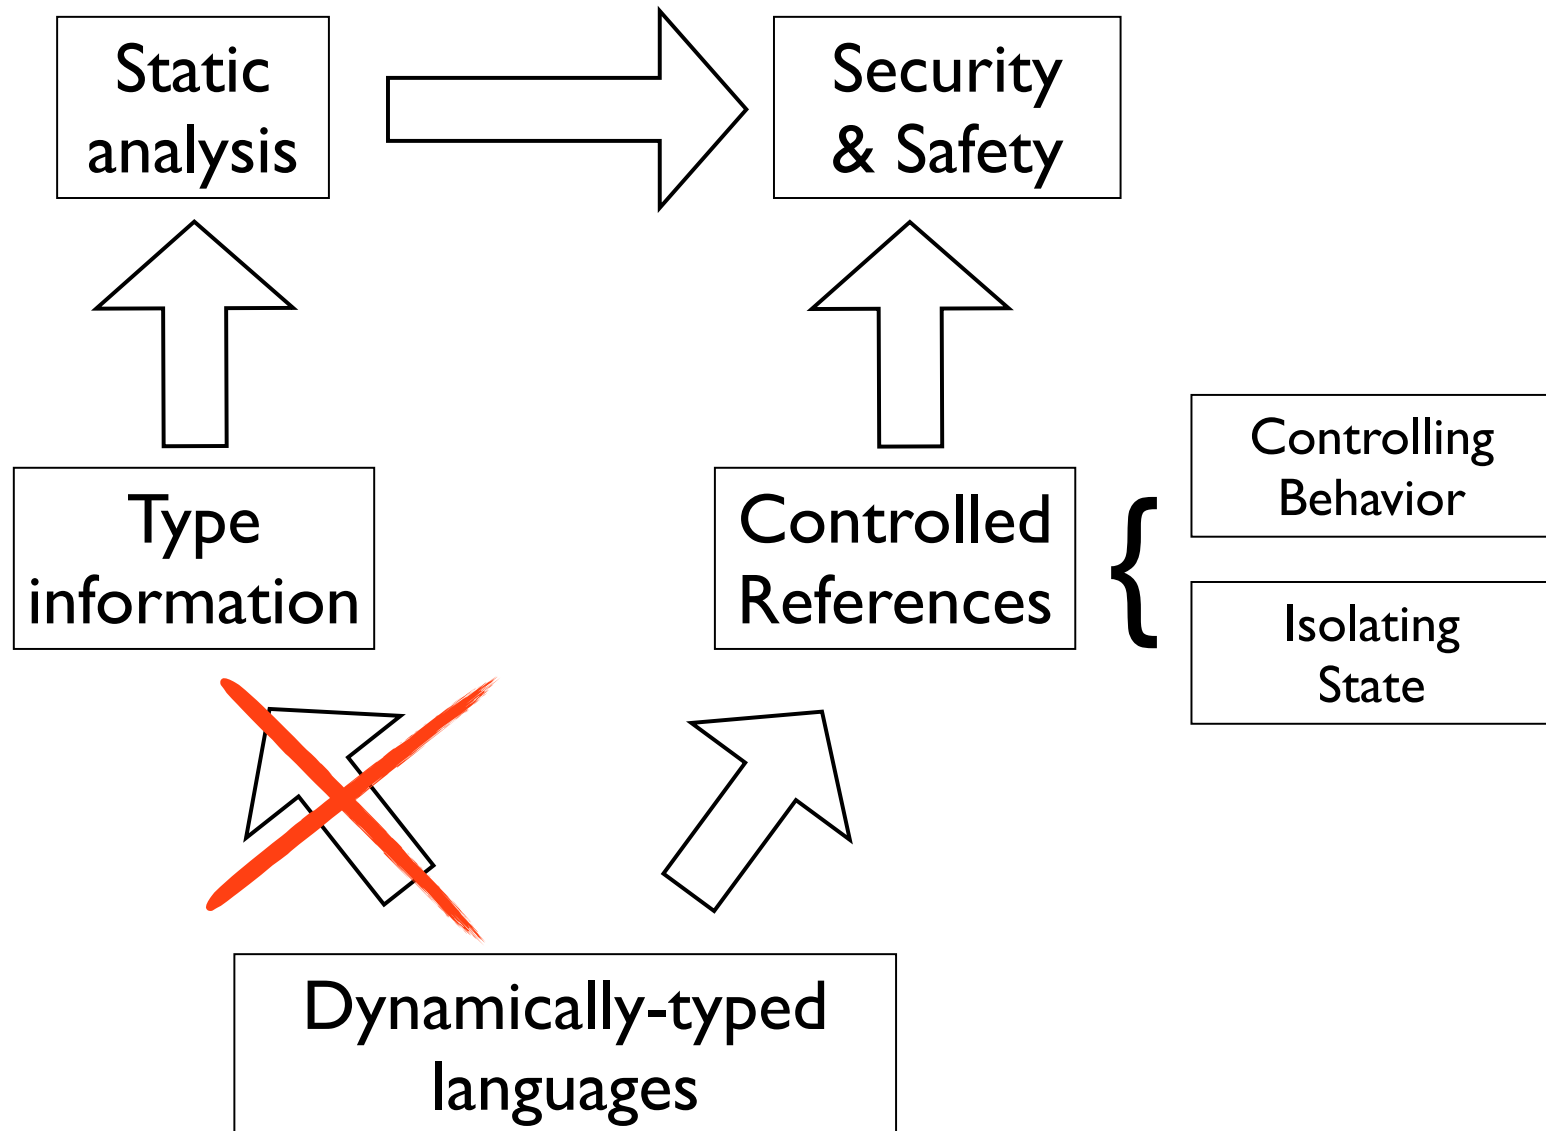


Controlled references ?  
**X** Controlled behavior  
**X** State isolation

# Security and Safety



# Security and Safety





# Controlled references


Use references to control the semantics of target object

Family's solution	Behavior control	State isolation
Object-oriented capabilities	Only Restrict behavior	No
Encapsulation	Only Restrict behavior	No
Dynamic ownership	Only Restrict behavior	No
Contextual value	No	Yes per thread

# Controlled references

Use references to control the semantics of target object

Family's solution	Behavior control	State isolation
Object-oriented capabilities	Only Restrict behavior	No
Encapsulation	Only Restrict behavior	No
Dynamic ownership	Only Restrict behavior	No
Contextual value	No	Yes per thread

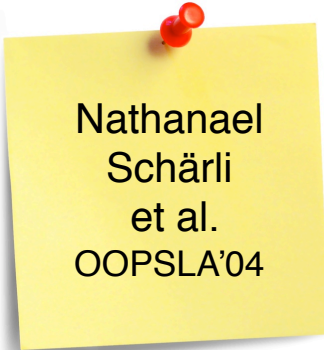


Mark Samuel  
Miller et al.  
IAFOR'2003

# Controlled references

Use references to control the semantics of target object

Family's solution	Behavior control	State isolation
Object-oriented capabilities	Only Restrict behavior	No
Encapsulation	Only Restrict behavior	No
Dynamic ownership	Only Restrict behavior	No
Contextual value	No	Yes per thread




Nathanael  
Schärli  
et al.  
OOPSLA'04

# Controlled references

Use references to control the semantics of target object

Family's solution	Behavior control	State isolation
Object-oriented capabilities	Only Restrict behavior	No
Encapsulation	Only Restrict behavior	No
Dynamic ownership	Only Restrict behavior	No
Contextual value	No	Yes per thread




Gordon et al.  
DLS'2007

# Controlled references

Use references to control the semantics of target object

Family's solution	Behavior control	State isolation
Object-oriented capabilities	Only Restrict behavior	No
Encapsulation	Only Restrict behavior	No
Dynamic ownership	Only Restrict behavior	No
Contextual value	No	Yes per thread



Eric Tanter  
DLS'2008


# Problem

How can we provide a *general-purpose* way to control references in dynamically-typed languages?

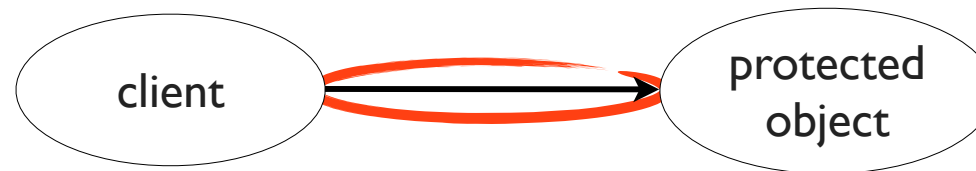
# Thesis

In the context of dynamically-typed languages, reifying references, controlling behavior, and isolating state via such references, is a practical way to control references

# Solution

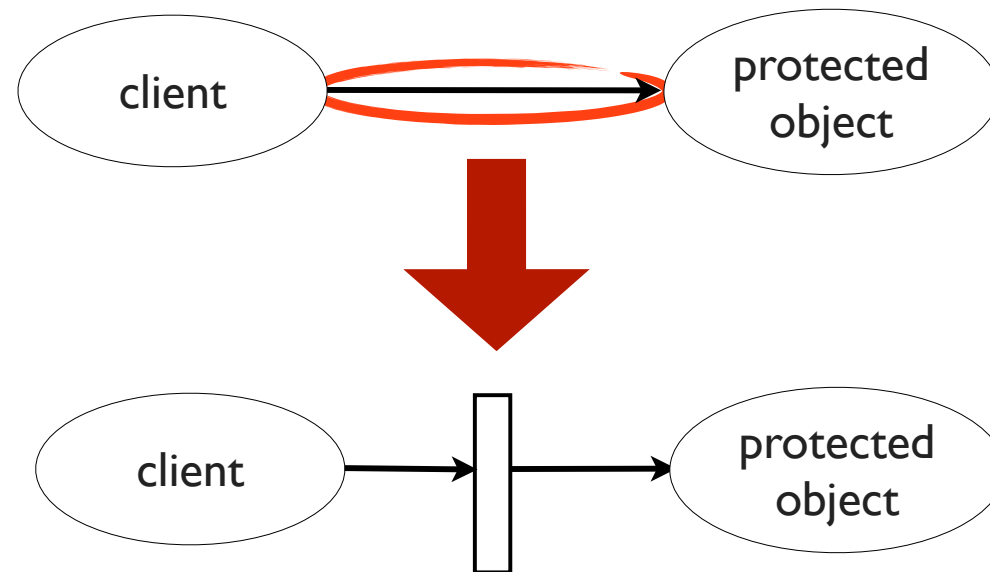
- 
- ✗ Reifying references**
    - ▶ **✗ Controlling behavior**
    - ▶ **✗ Isolating state**

# Reifying references



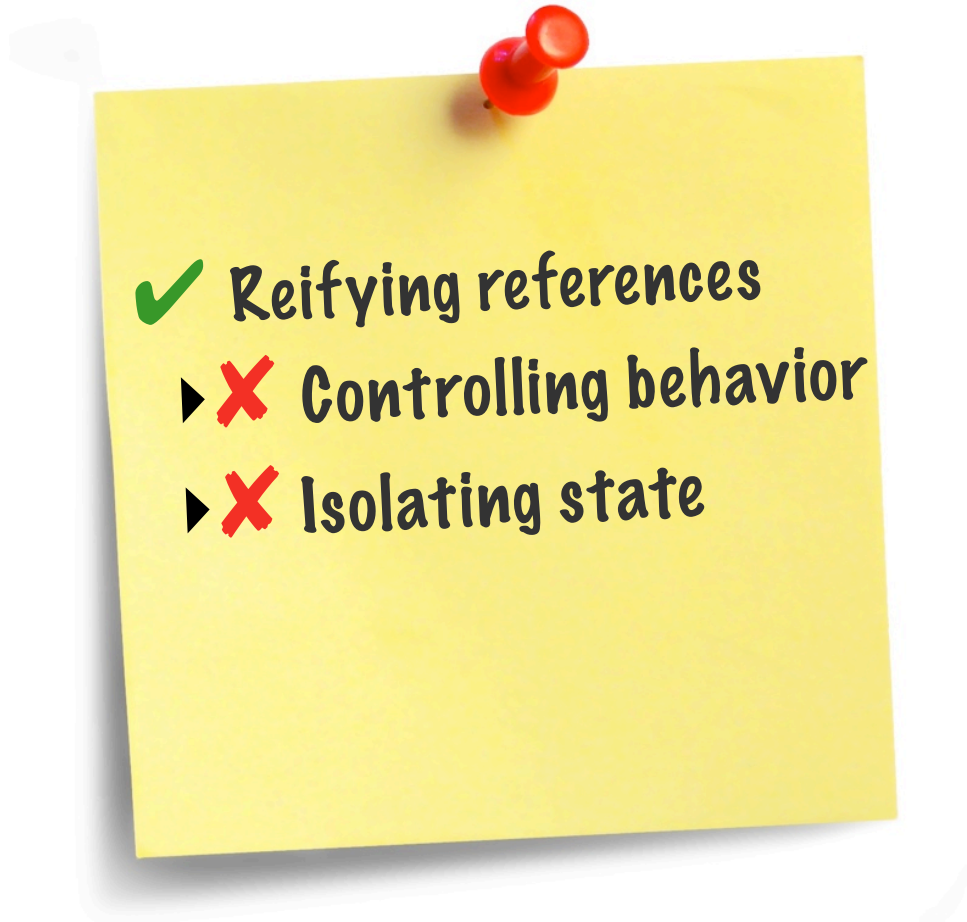


# Reifying references



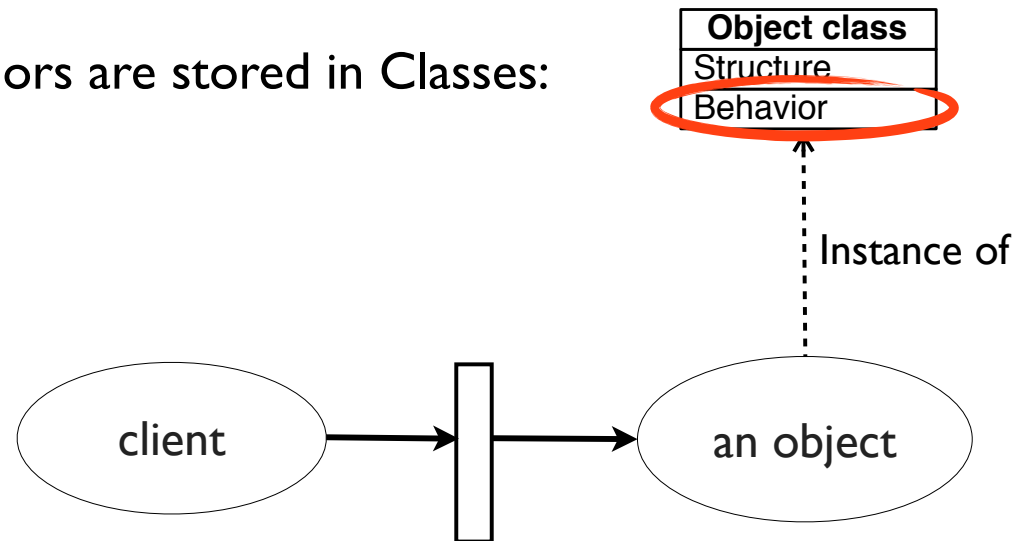
***Handle***

# Solution

- 
- ✓ Reifying references
  - ▶ ✗ Controlling behavior
  - ▶ ✗ Isolating state

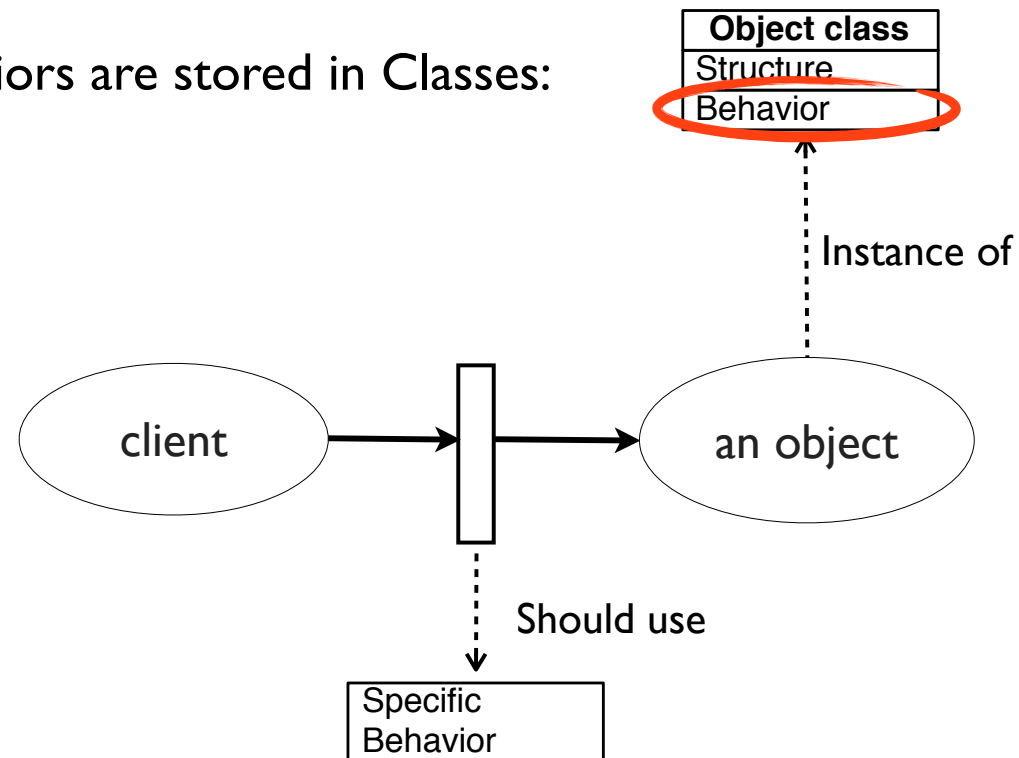
# Controlling behavior

Behaviors are stored in Classes:

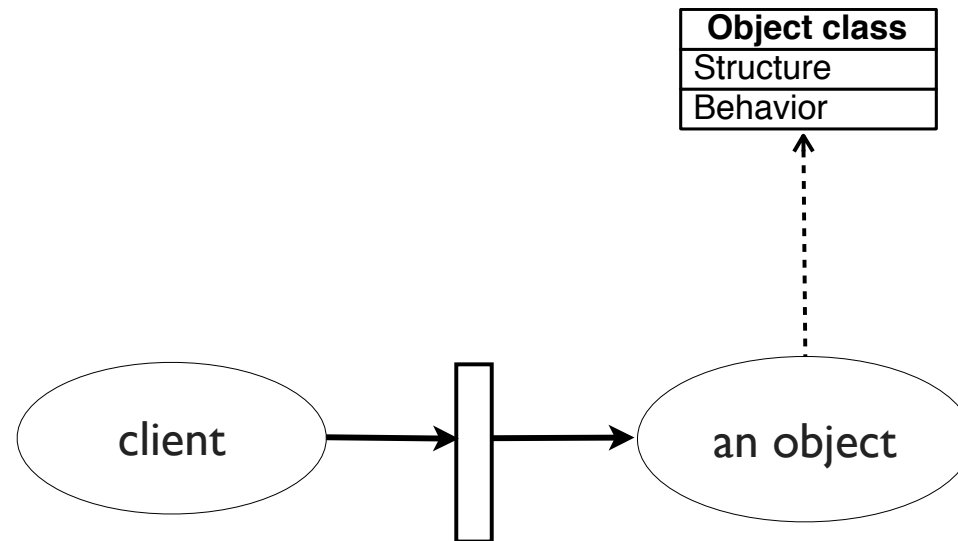


# Controlling behavior

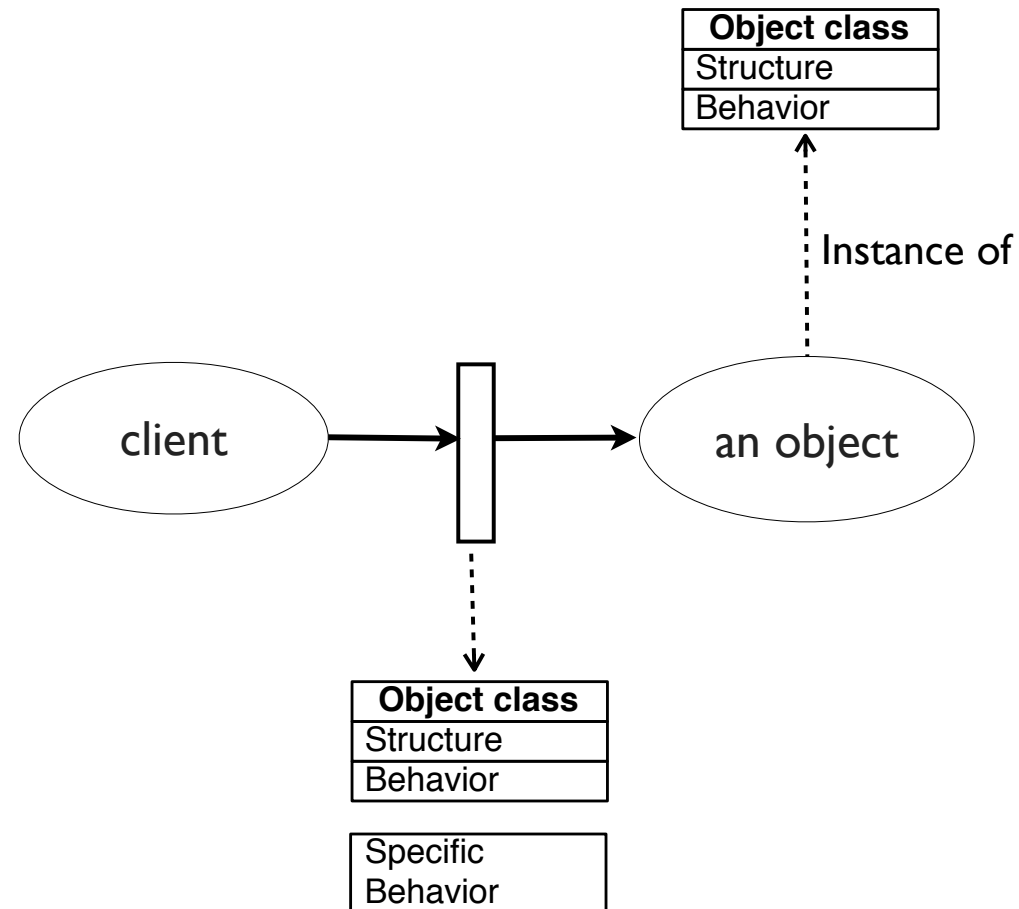
Behaviors are stored in Classes:



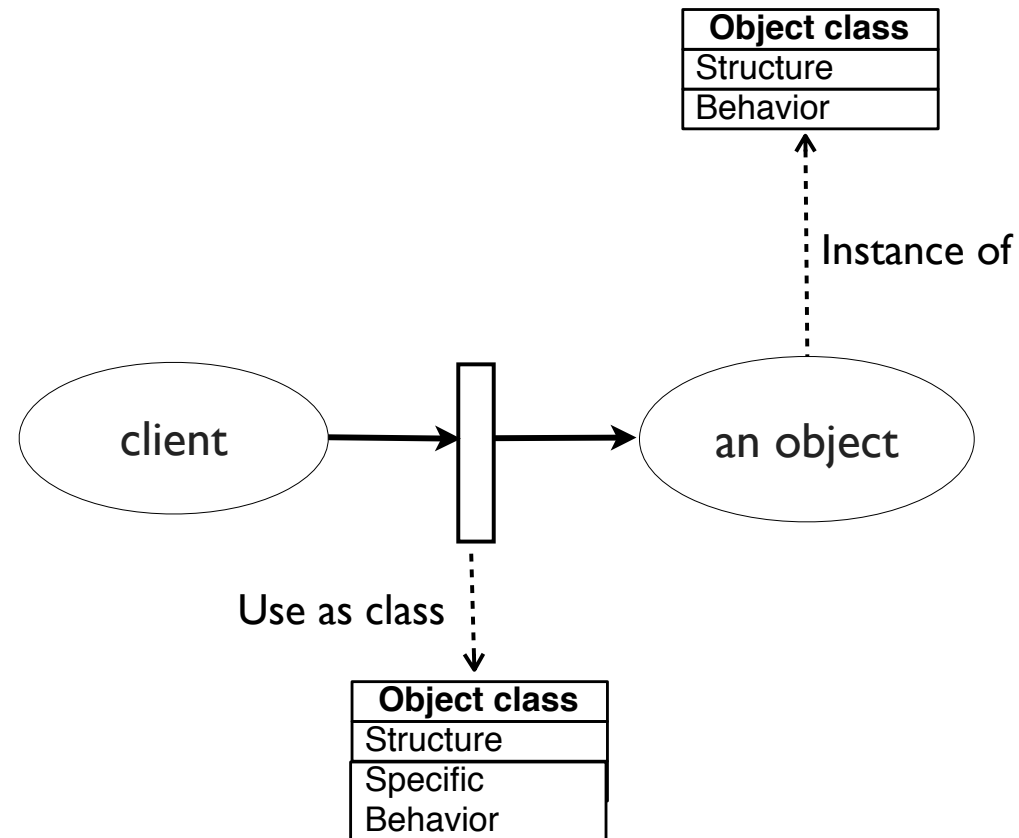
# Controlling behavior



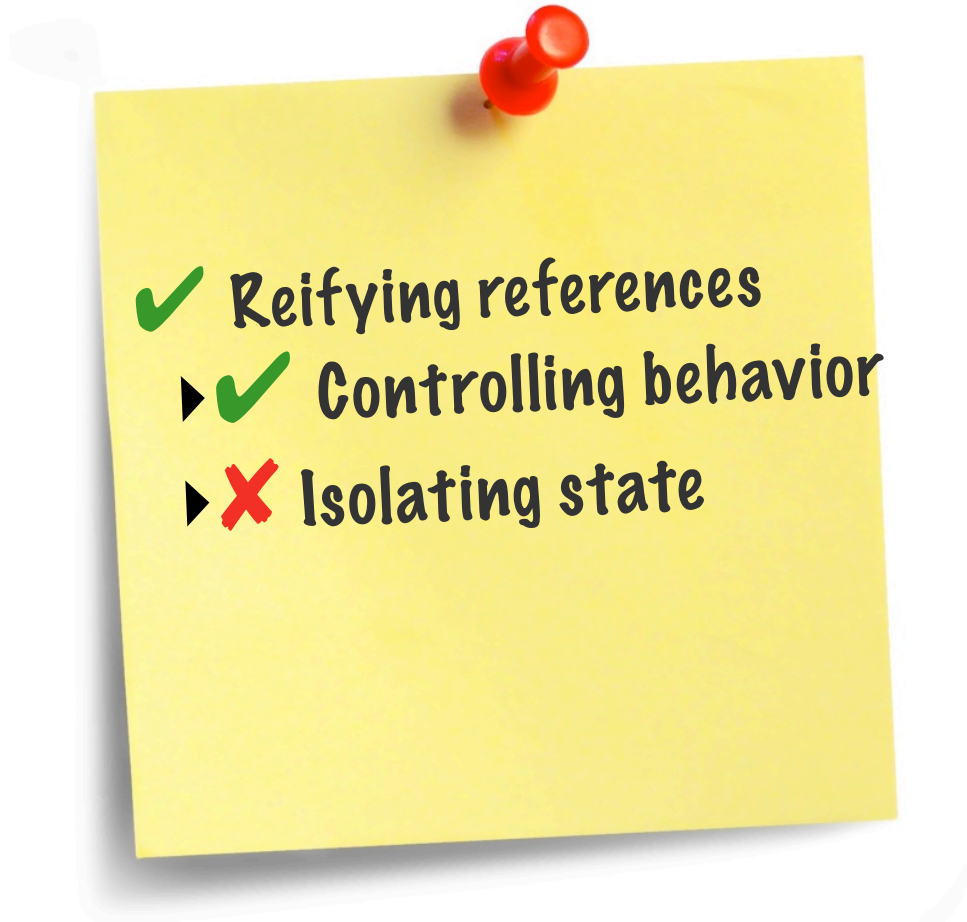
# Controlling behavior



# Controlling behavior

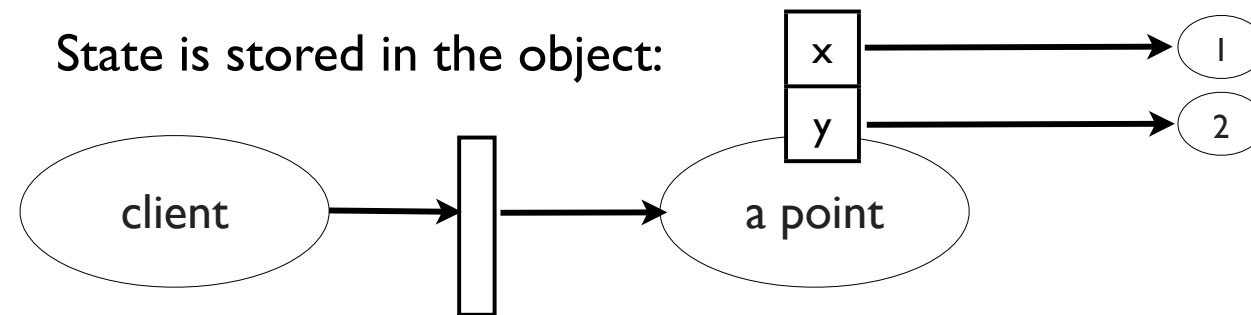


# Solution

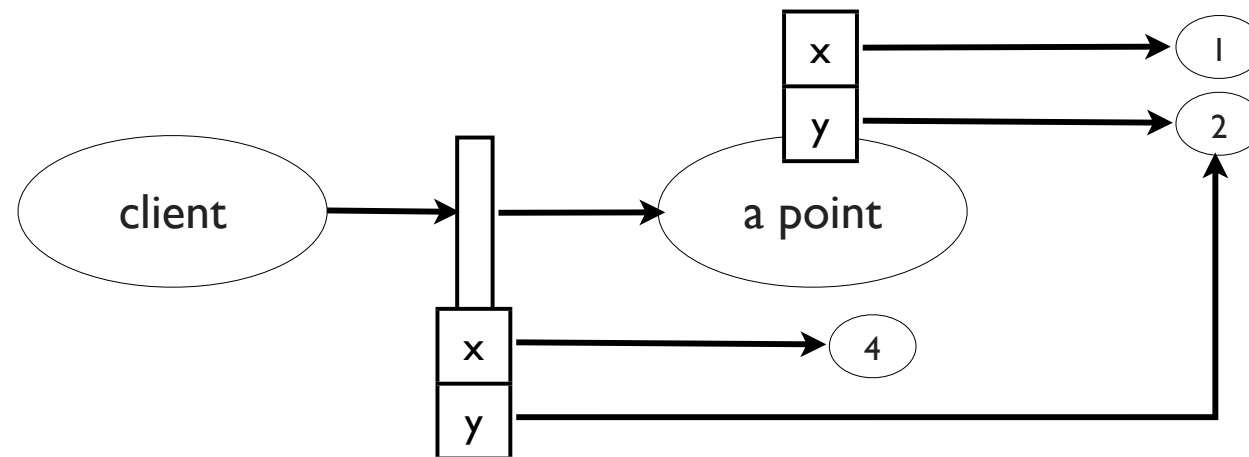
- 
- ✓ Reifying references
  - ▶ ✓ Controlling behavior
  - ▶ ✗ Isolating state



# Isolating State



# Isolating State

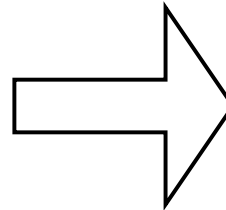


## Core

- ✓ Reifying references
  - ▶ ✓ Controlling behavior
  - ▶ ✓ Isolating state

## Core

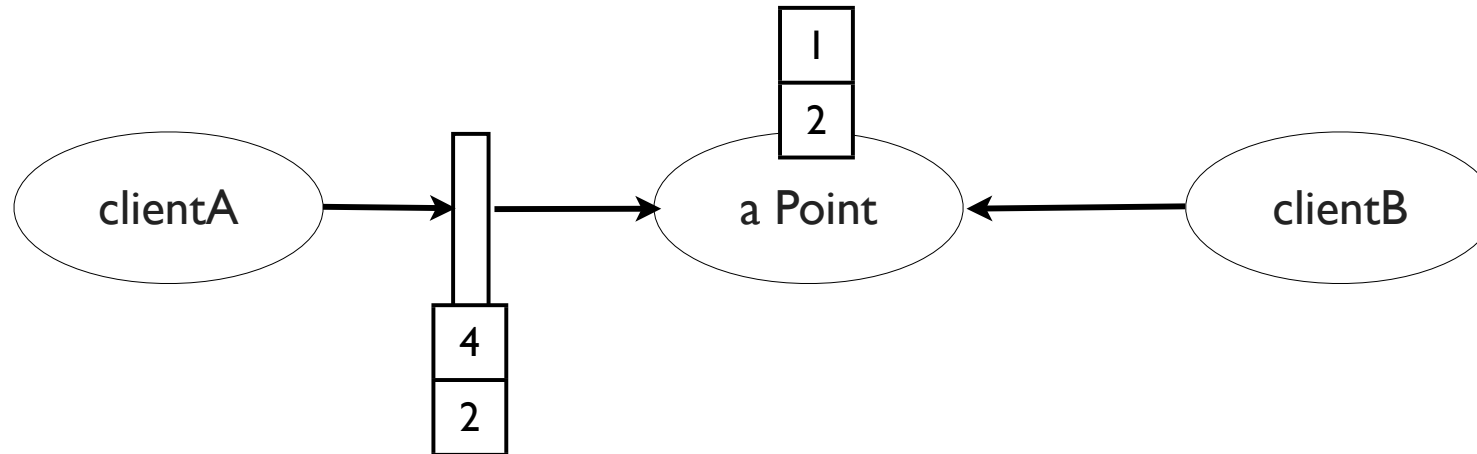
- ✓ Reifying references
- ▶ ✓ Controlling behavior
- ▶ ✓ Isolating state



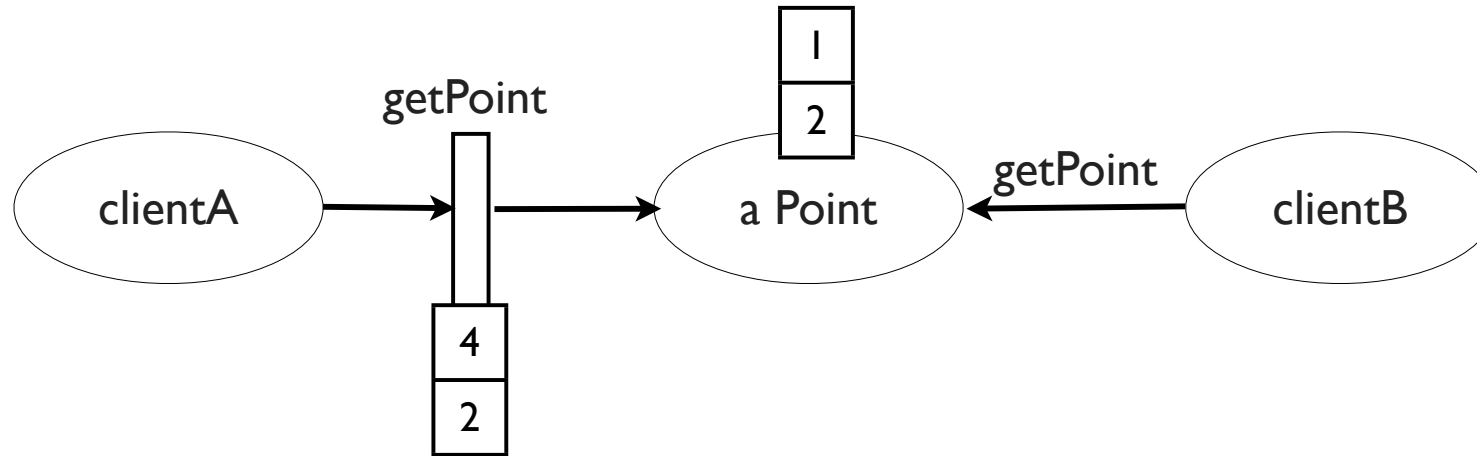
## Consequences

- ✗ Identity is broken
- ✗ Graph control is absent
- ✗ Handle creation

# Identity problem



# Identity problem



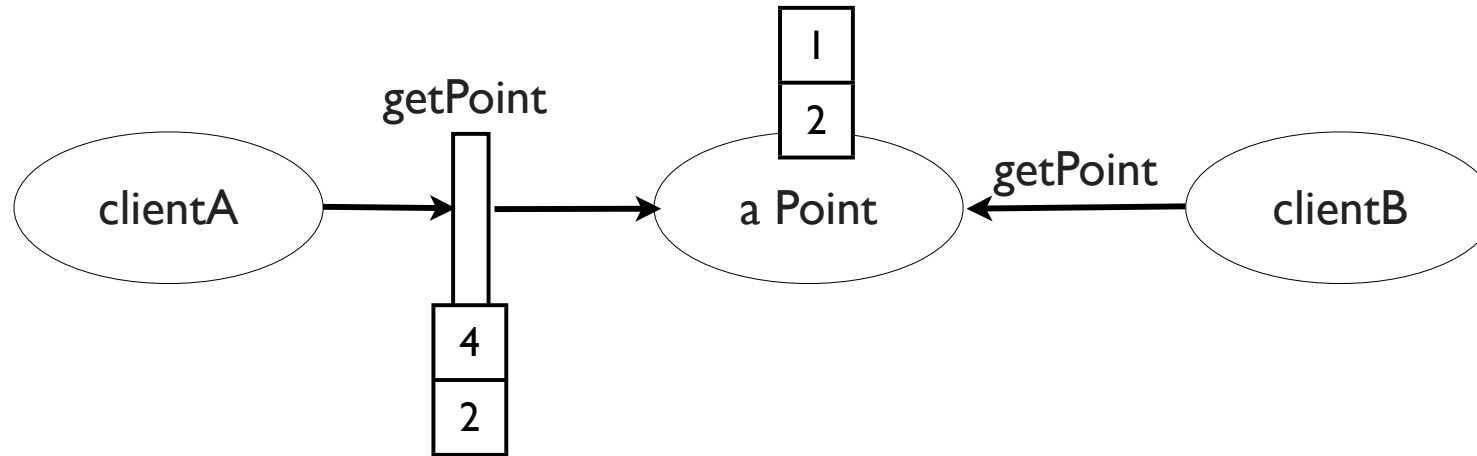
clientA getPoint

identity

clientB getPoint

True

# Identity problem



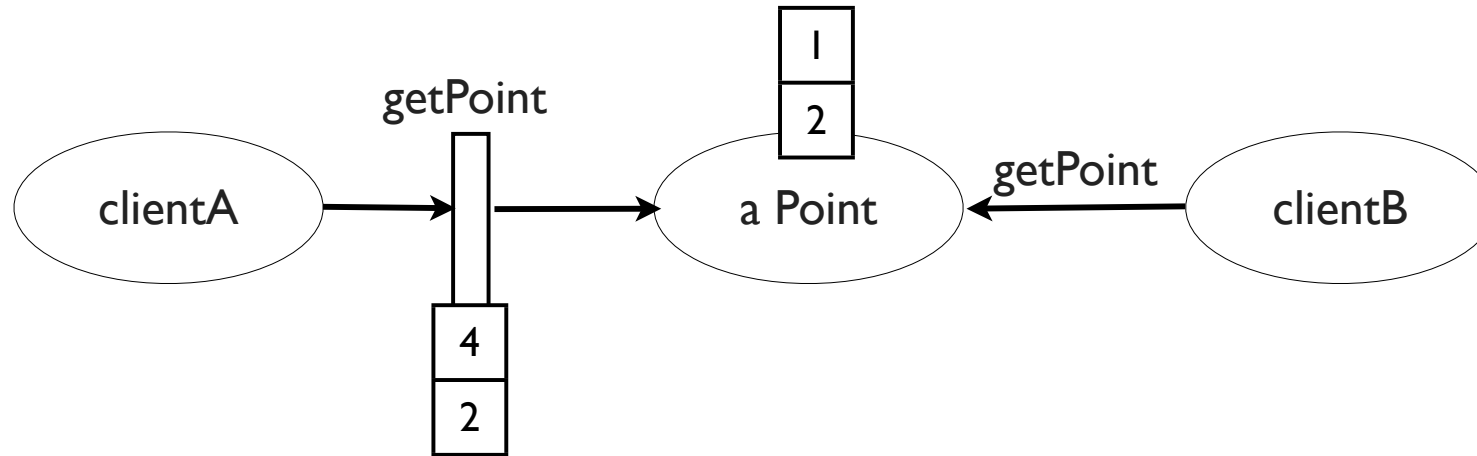
clientA getPoint

equality

clientB getPoint

**False**

# Identity problem



clientA `getPoint`

equality

clientB `getPoint`

**False**



Identity does not imply equality anymore



# Rewrite identity model

## Equality

clientA getPoint

equality

clientB getPoint

**False**

## Similarity (object identity)

clientA getPoint

similarity

clientB getPoint

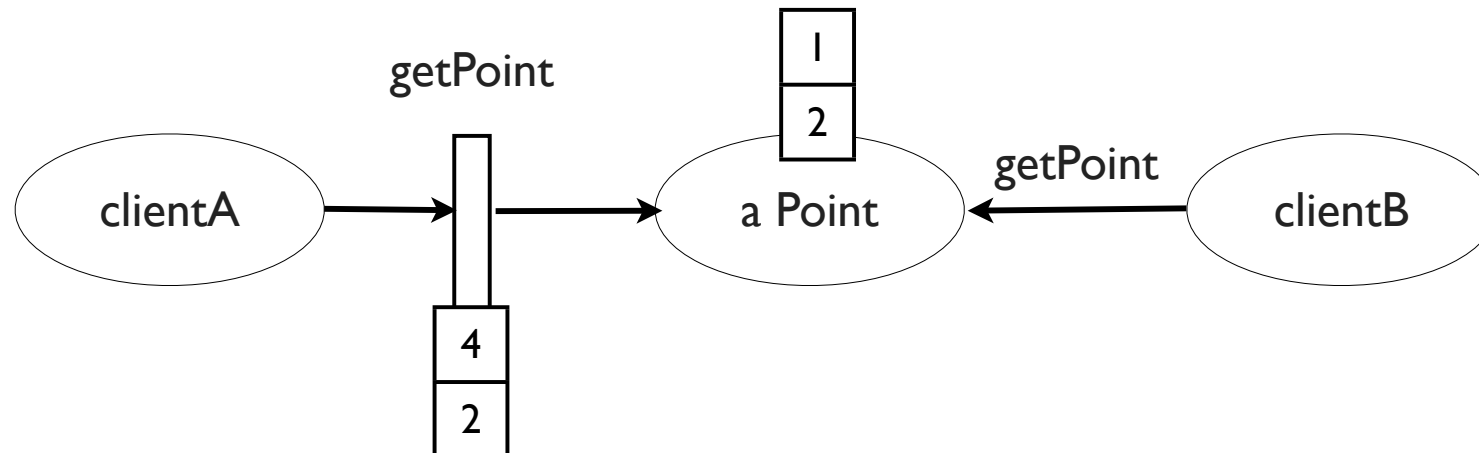
**True**

## Identity (reference identity)

clientA getPoint

identity

clientB getPoint

**False**

# Rewrite identity model

## Equality

clientA getPoint

equality

clientB getPoint

False

## Similarity (object identity)

clientA getPoint

similarity

clientB getPoint

True

## Identity (reference identity)

clientA getPoint

identity

clientB getPoint

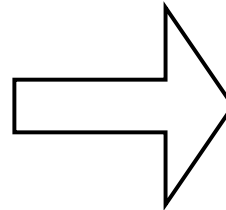
False

We can distinguish objects and references  
Identity still means equality

The identity invariant is maintained

## Core

- ✓ Reifying references
- ▶ ✓ Controlling behavior
- ▶ ✓ Isolating state

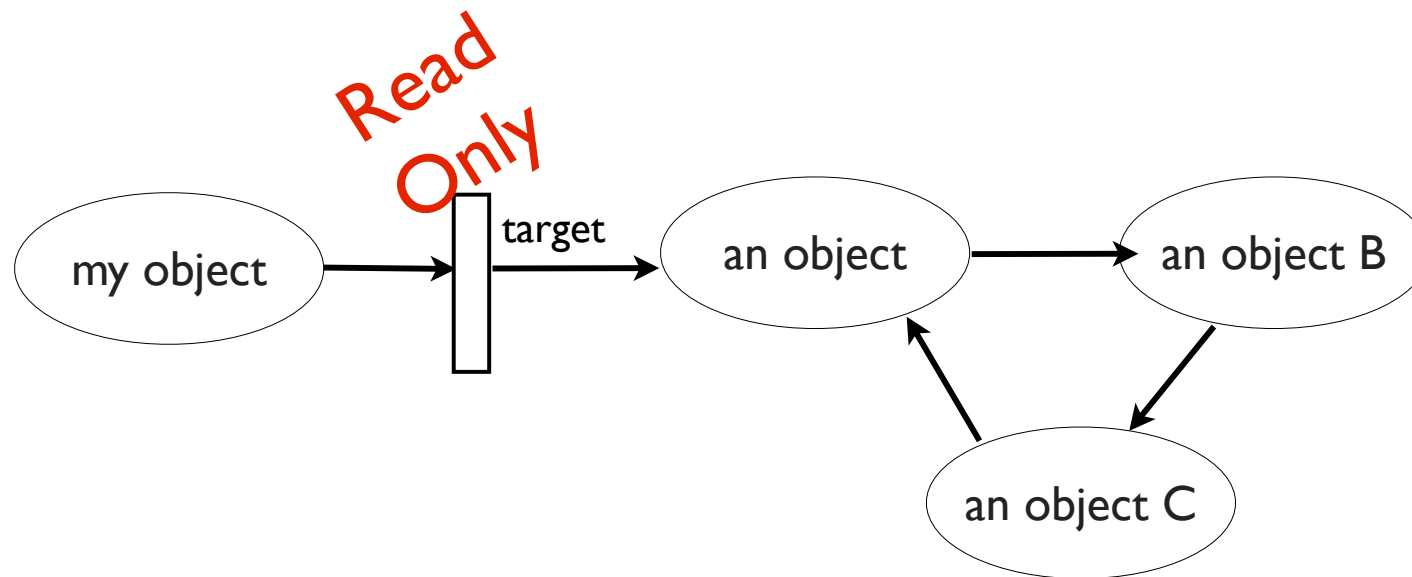


## Consequences

- ✓ Identity is preserved
- ✗ Graph control is absent
- ✗ Handle creation

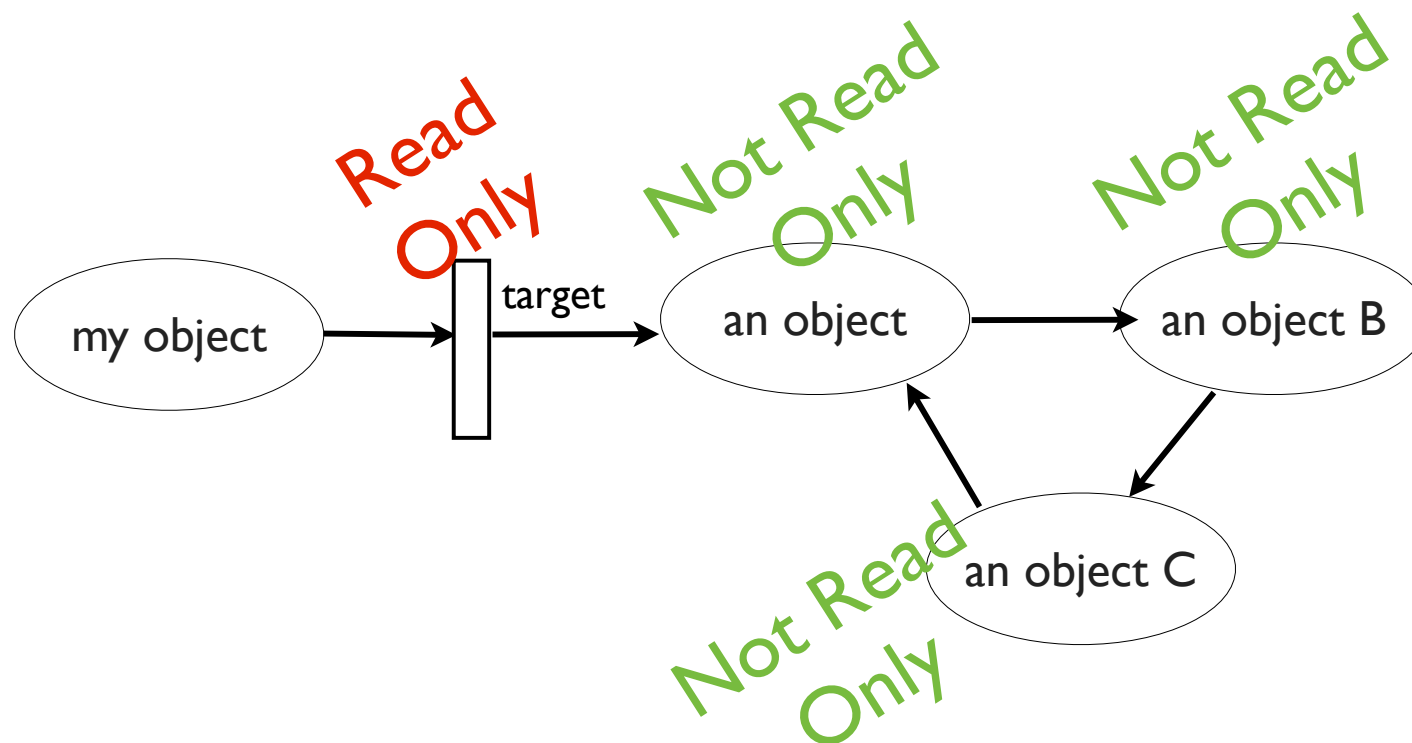
# Controlling graph

Controlling one atomic reference is not enough



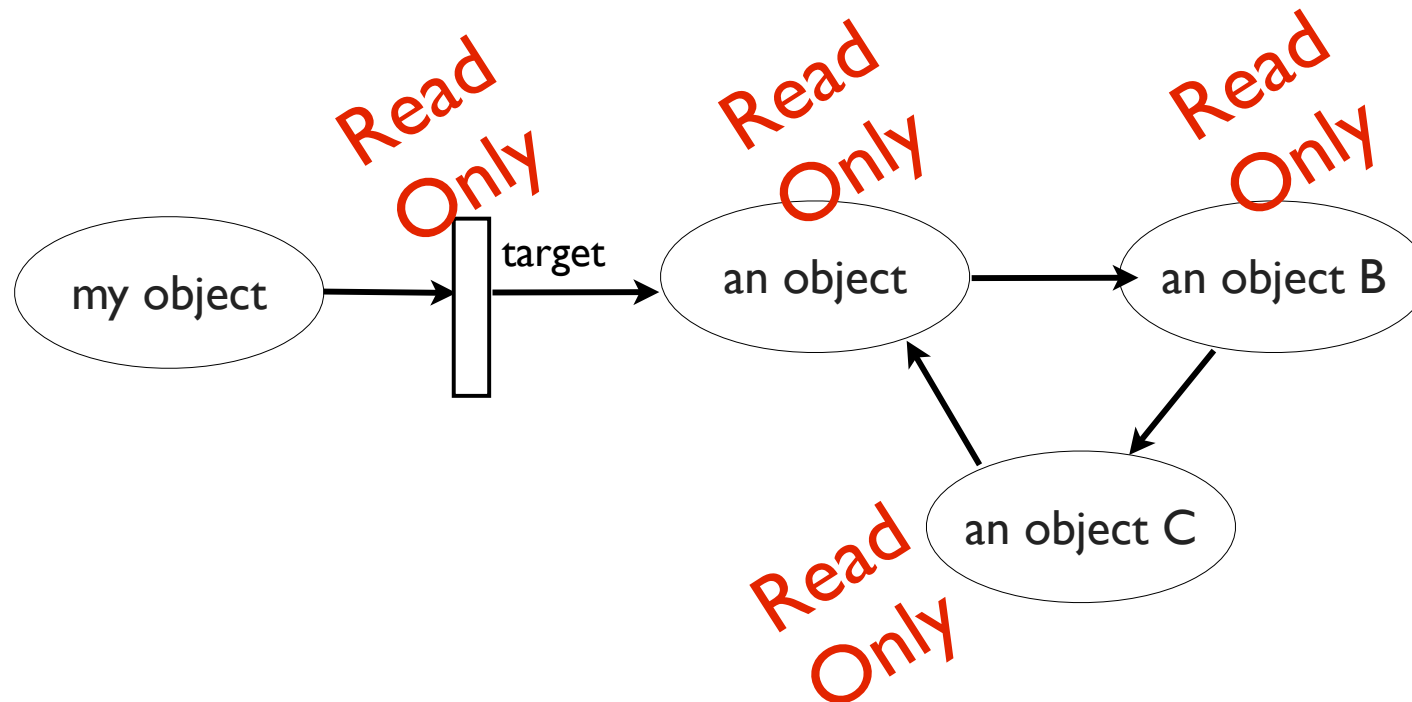
# Controlling graph

Controlling one atomic reference is not enough

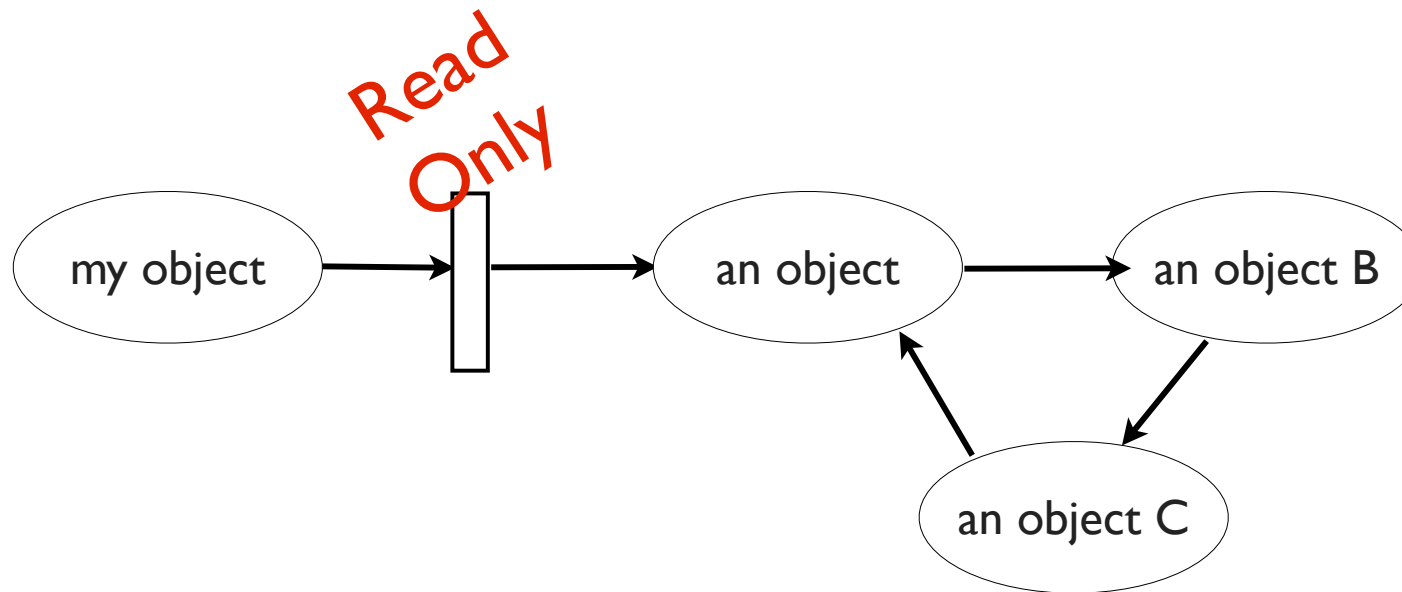


# Controlling graph

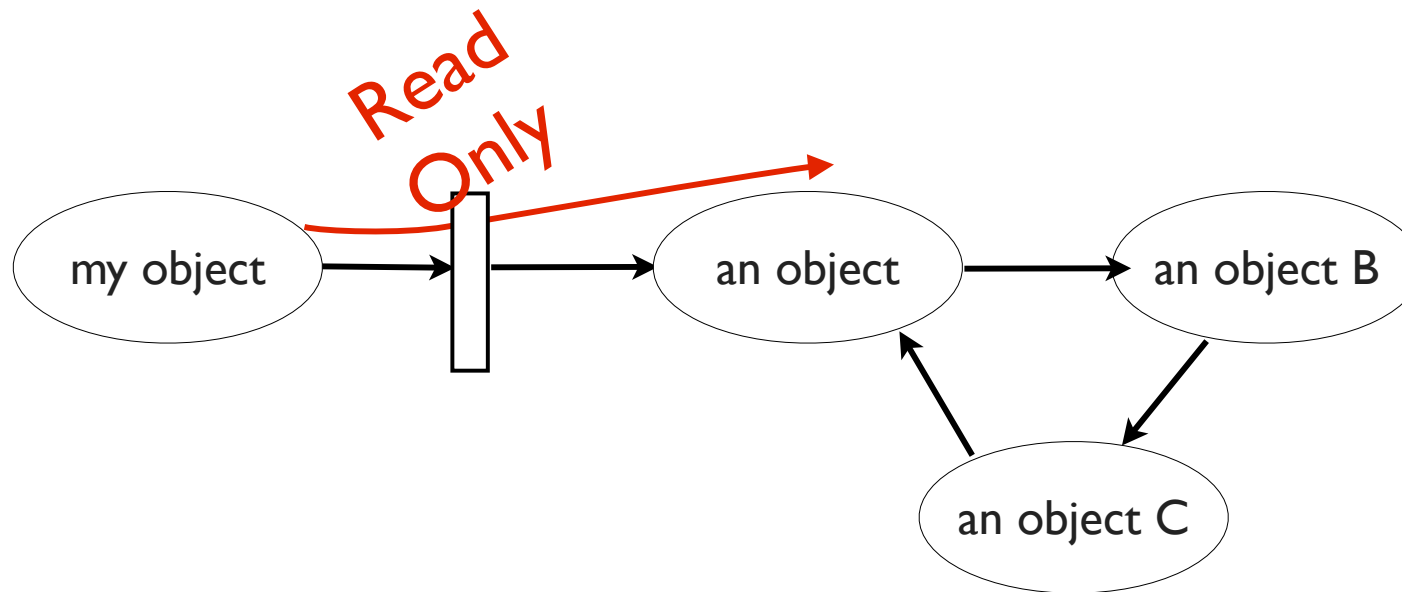
Controlling one atomic reference is not enough



# Propagation Semantics

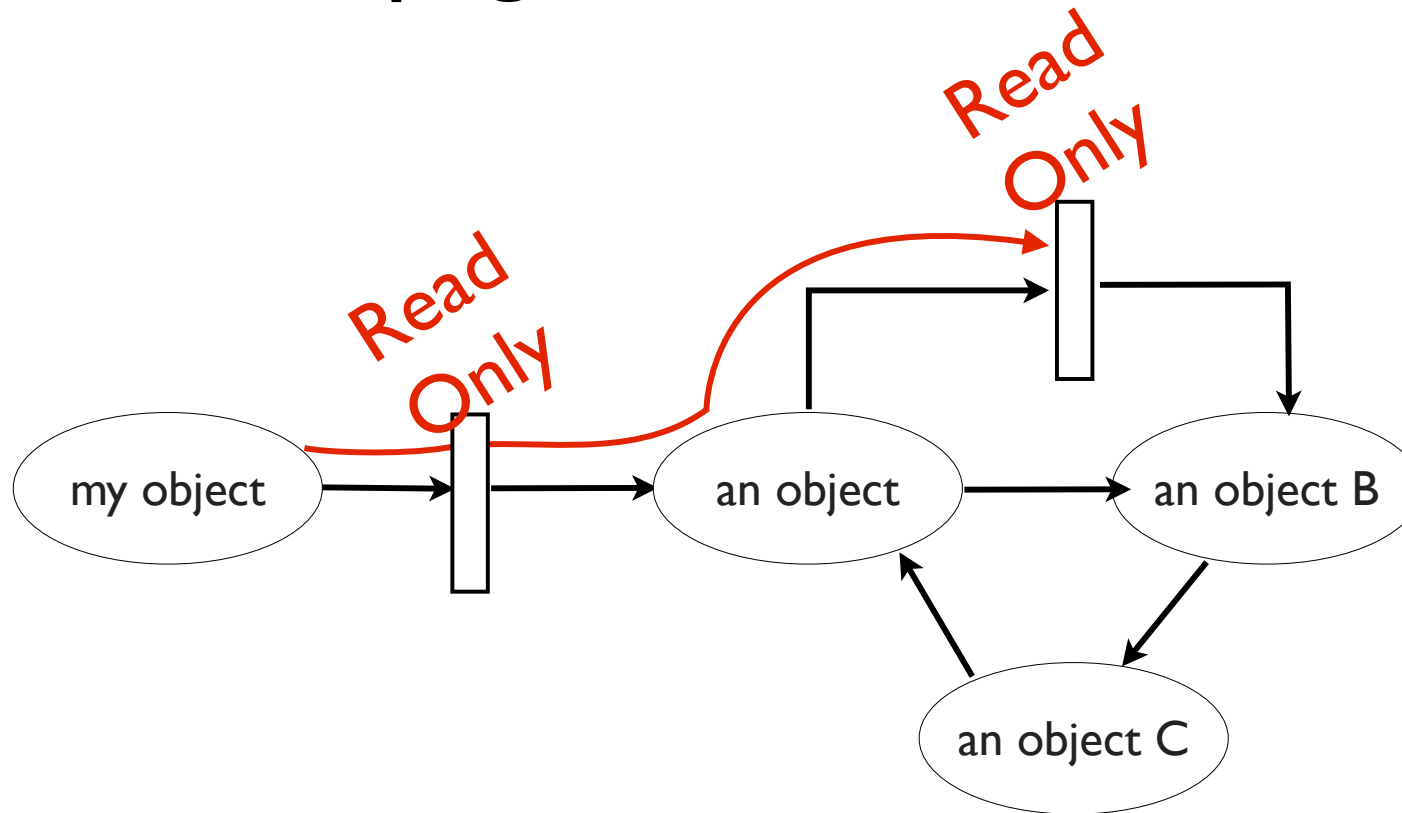


# Propagation Semantics

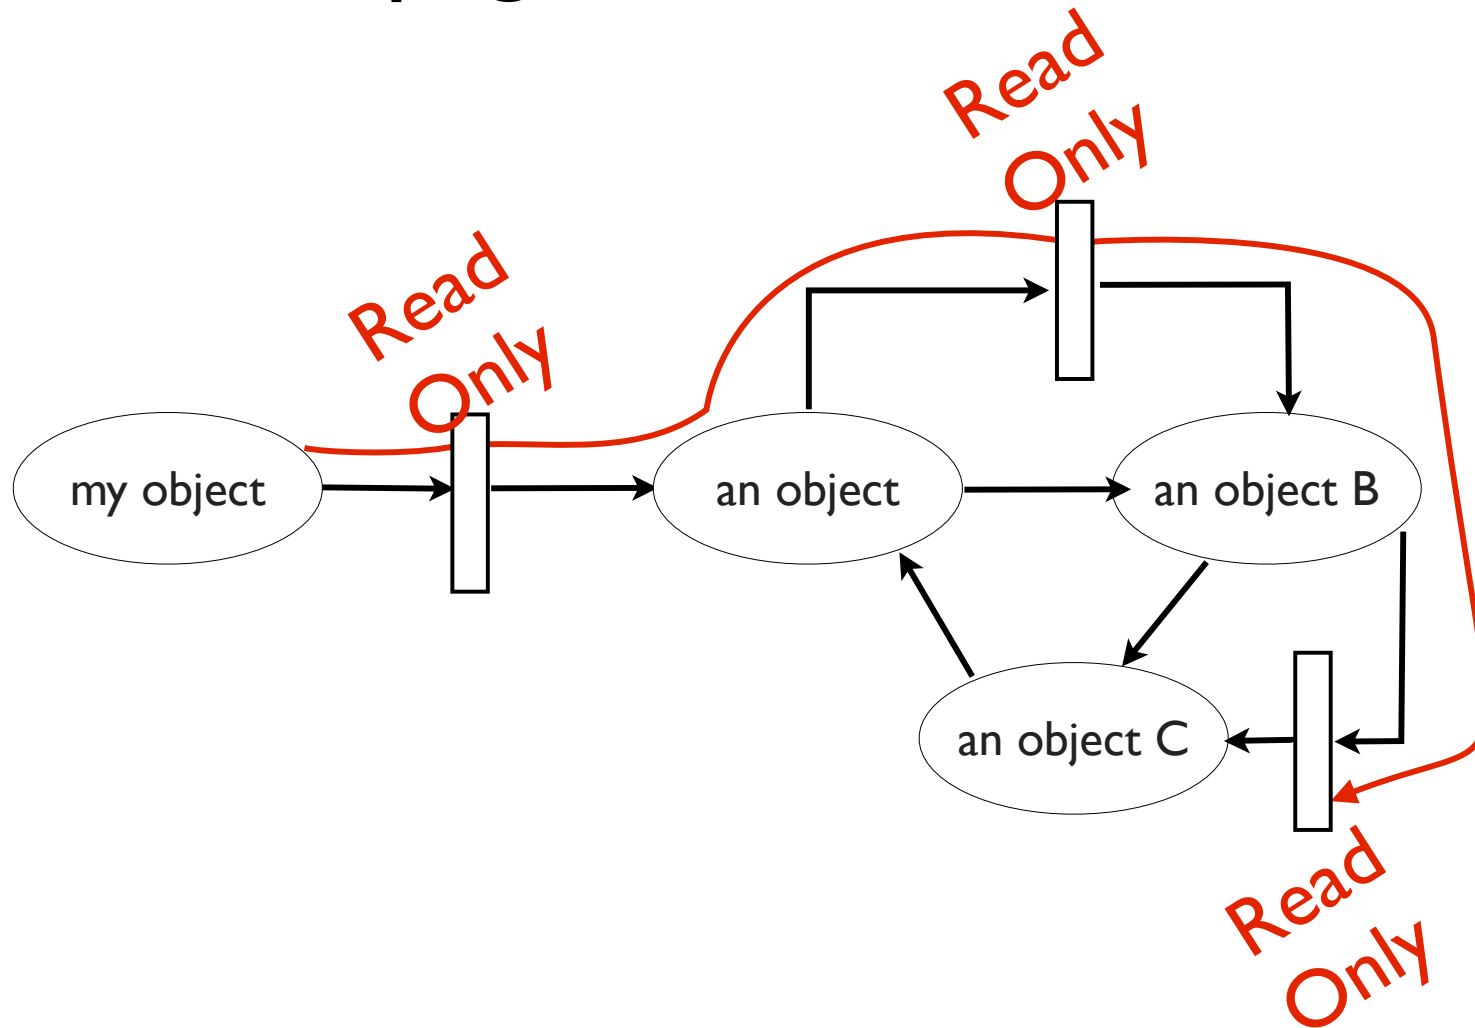




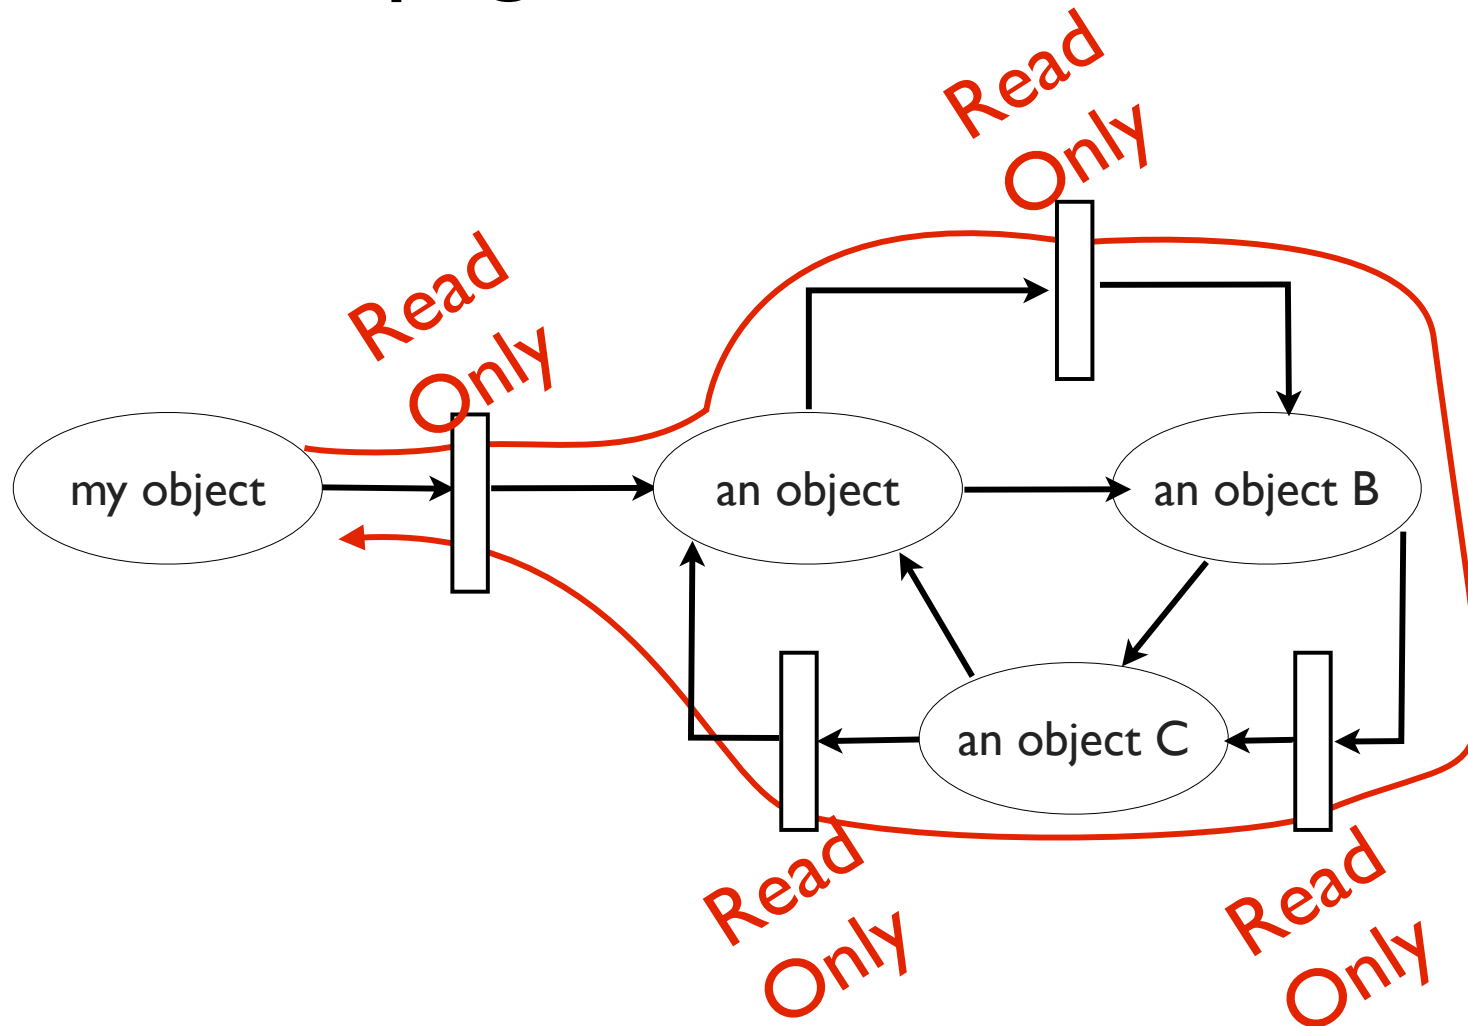
# Propagation Semantics



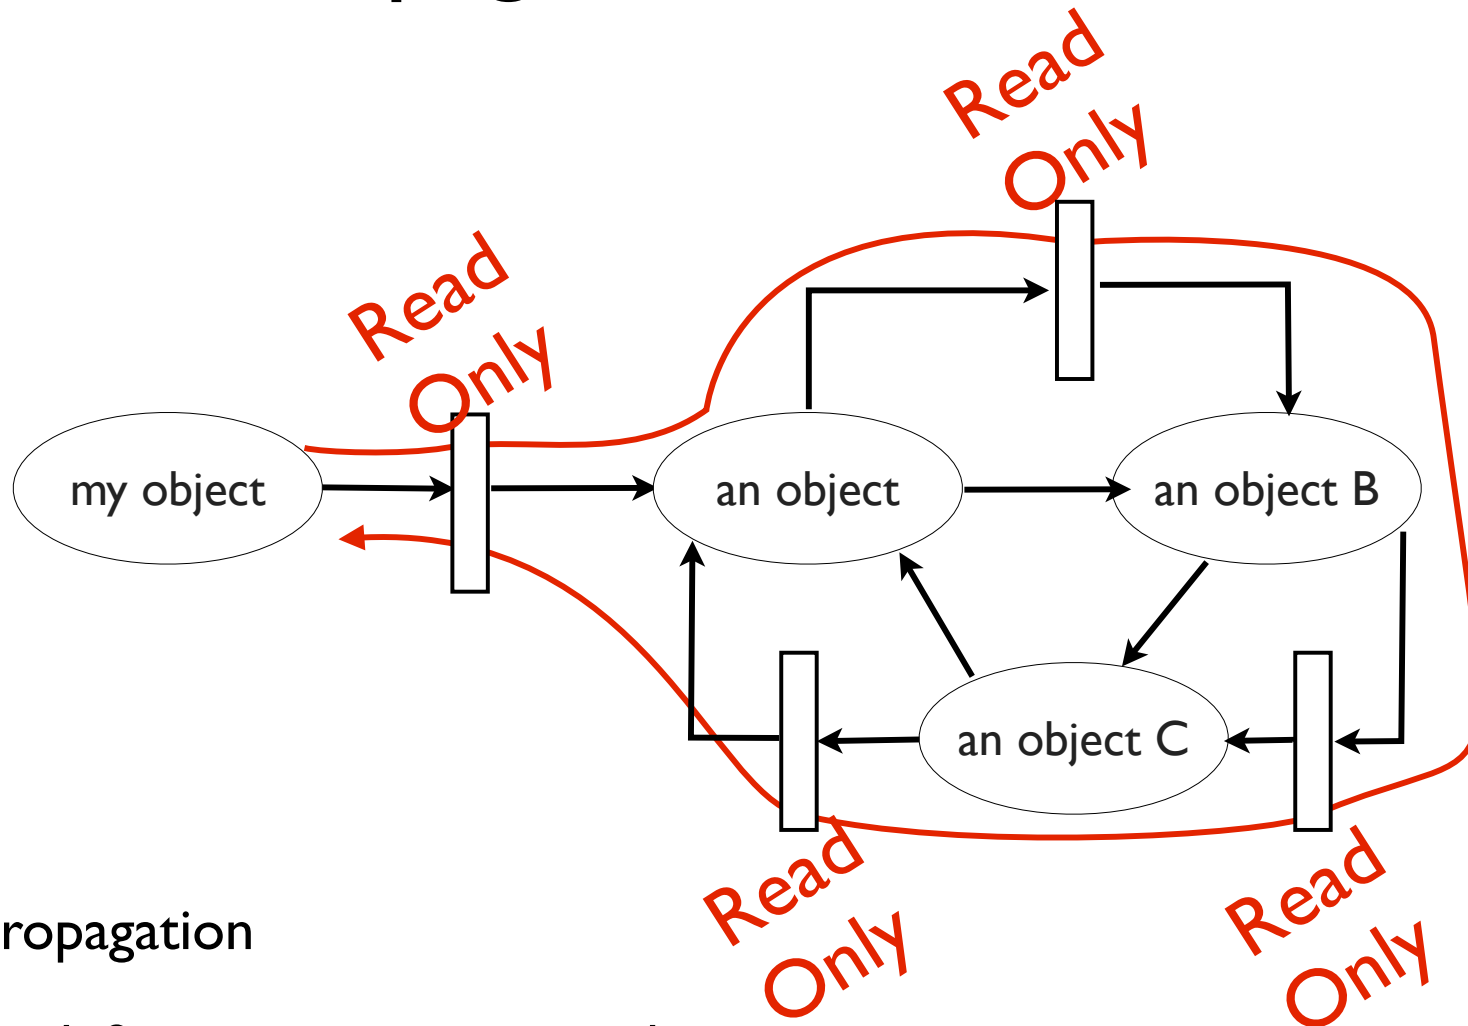
# Propagation Semantics



# Propagation Semantics



# Propagation Semantics

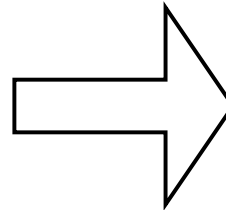


Handle propagation

- Handles define a propagation mechanism
- Propagation automatically performed by the Virtual Machine

## Core

- ✓ Reifying references
  - ▶ ✓ Controlling behavior
  - ▶ ✓ Isolating state



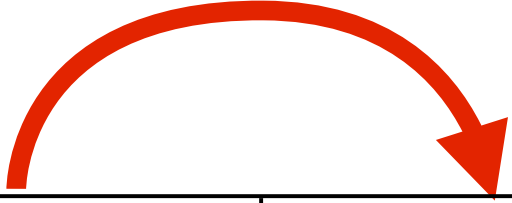
## Consequences

- ✓ Identity is preserved
- ✓ Propagated semantics
- ✗ Handle creation

# Handle: life cycle

Phases	1 Initialization	2 Execution
Representation	Object	Handle
Possibles actions	Configuration	Apply the semantics defined

# Handle: life cycle



Phases	1 Initialization	2 Execution
Representation	Object	Handle
Possibles actions	Configuration	Apply the semantics defined

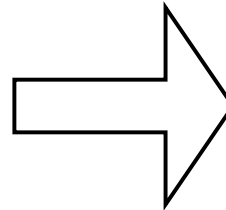
# Handle: life cycle

Phases	1 Initialization	2 Execution
Representation	Object	Handle
Possibles actions	Configuration	Apply the semantics defined



## Core

- ✓ Reifying references
  - ▶ ✓ Controlling behavior
  - ▶ ✓ Isolating state

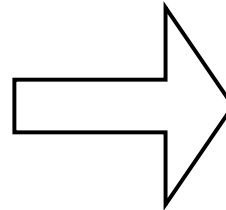


## Consequences

- ✓ Identity is preserved
- ✓ Propagated semantics
- ✓ Handle creation

## Core

- ✓ Reifying references
  - ▶ ✓ Controlling behavior
  - ▶ ✓ Isolating state



## Consequences

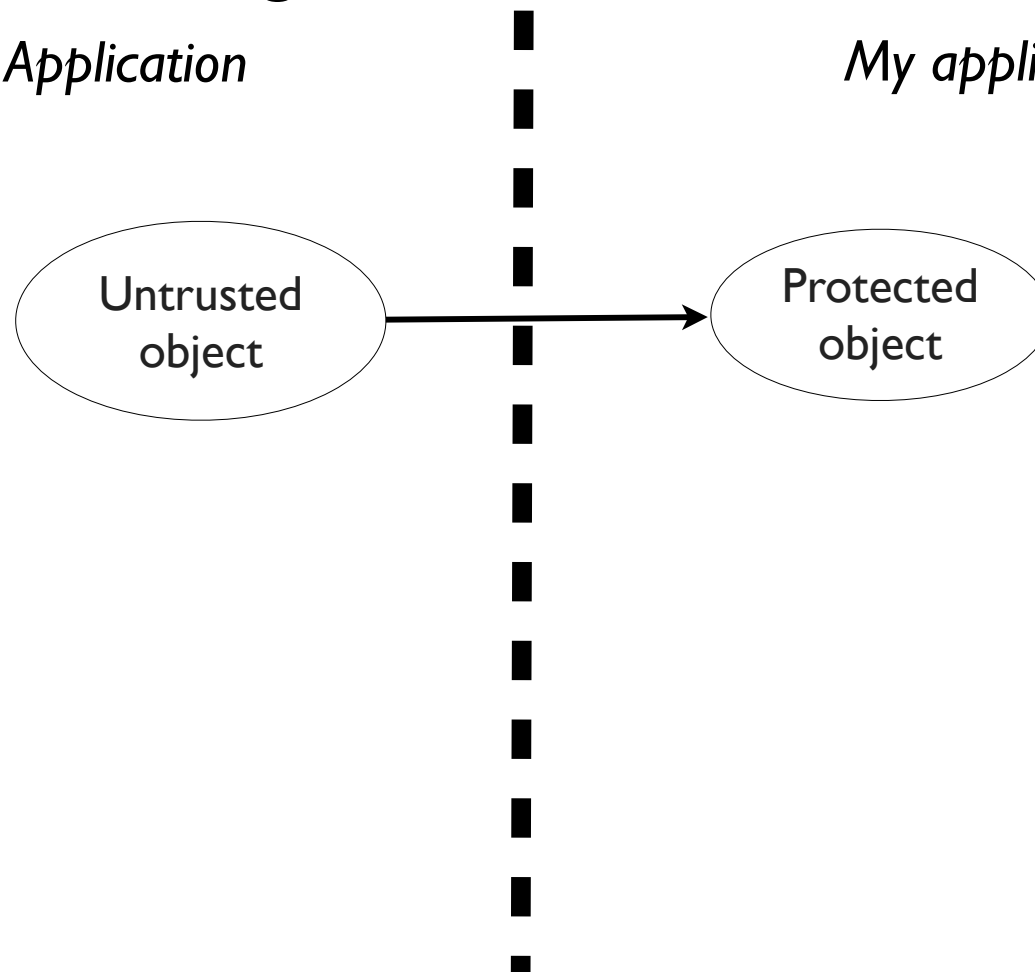
- ✓ Identity is preserved
- ✓ Propagated semantics
- ✓ Handle creation
  - ▶ ✗ Handle control ?

# Handle: control ?

If we want to change the semantics:

*Foreign Application*

*My application*

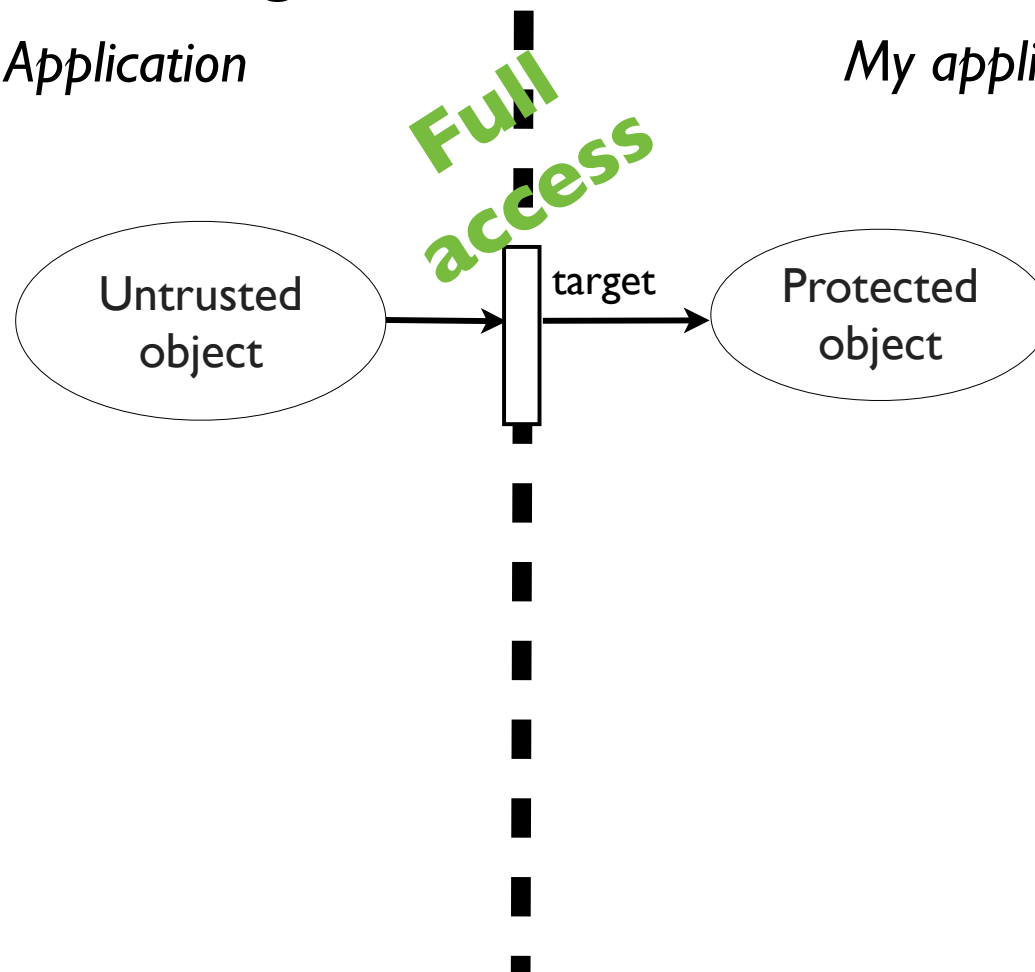


# Handle: control ?

If we want to change the semantics:

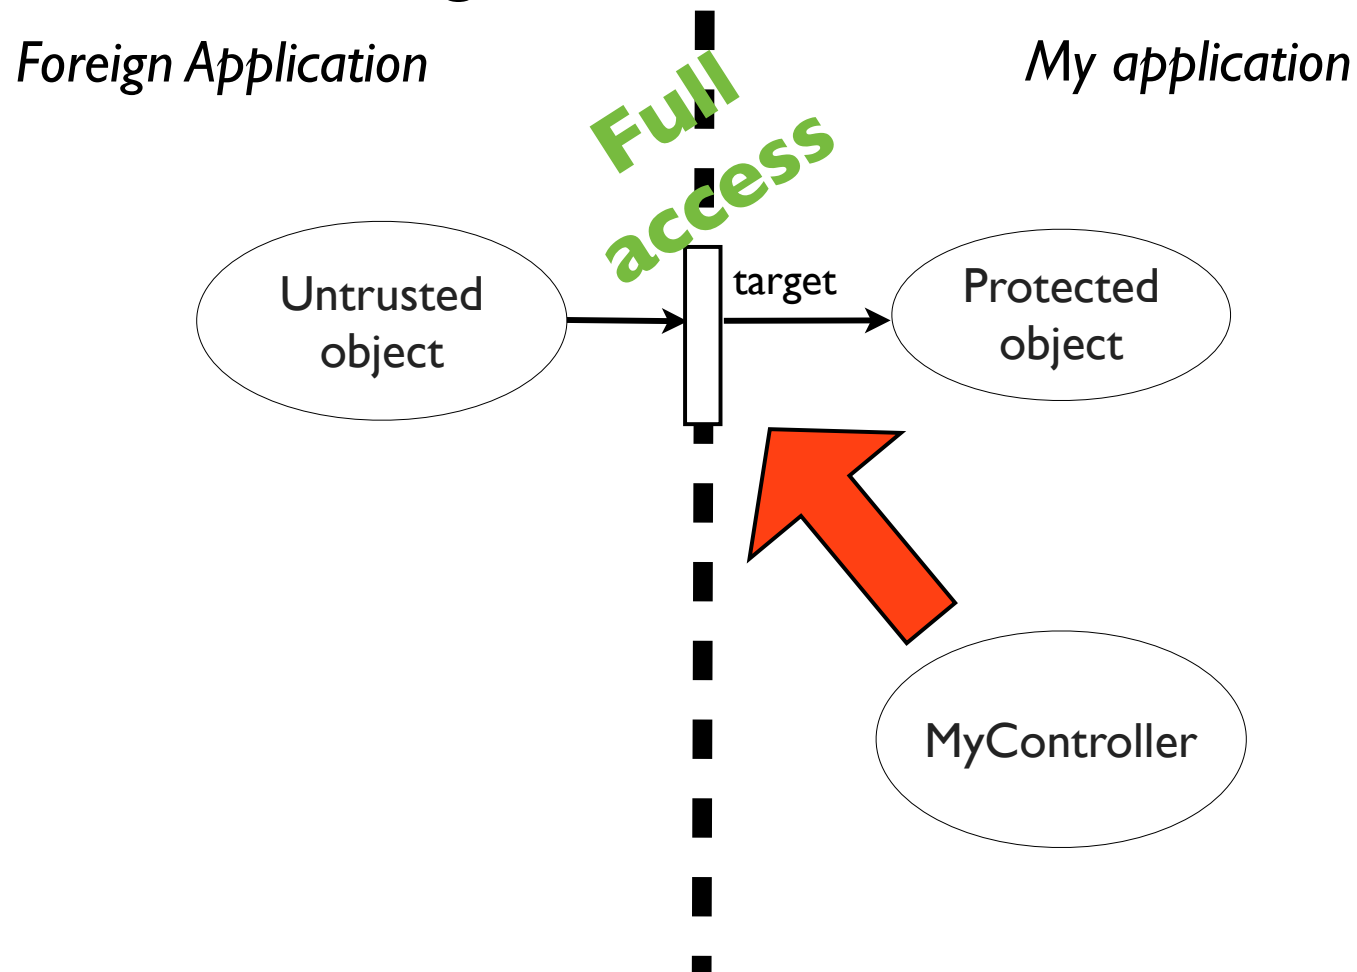
*Foreign Application*

*My application*



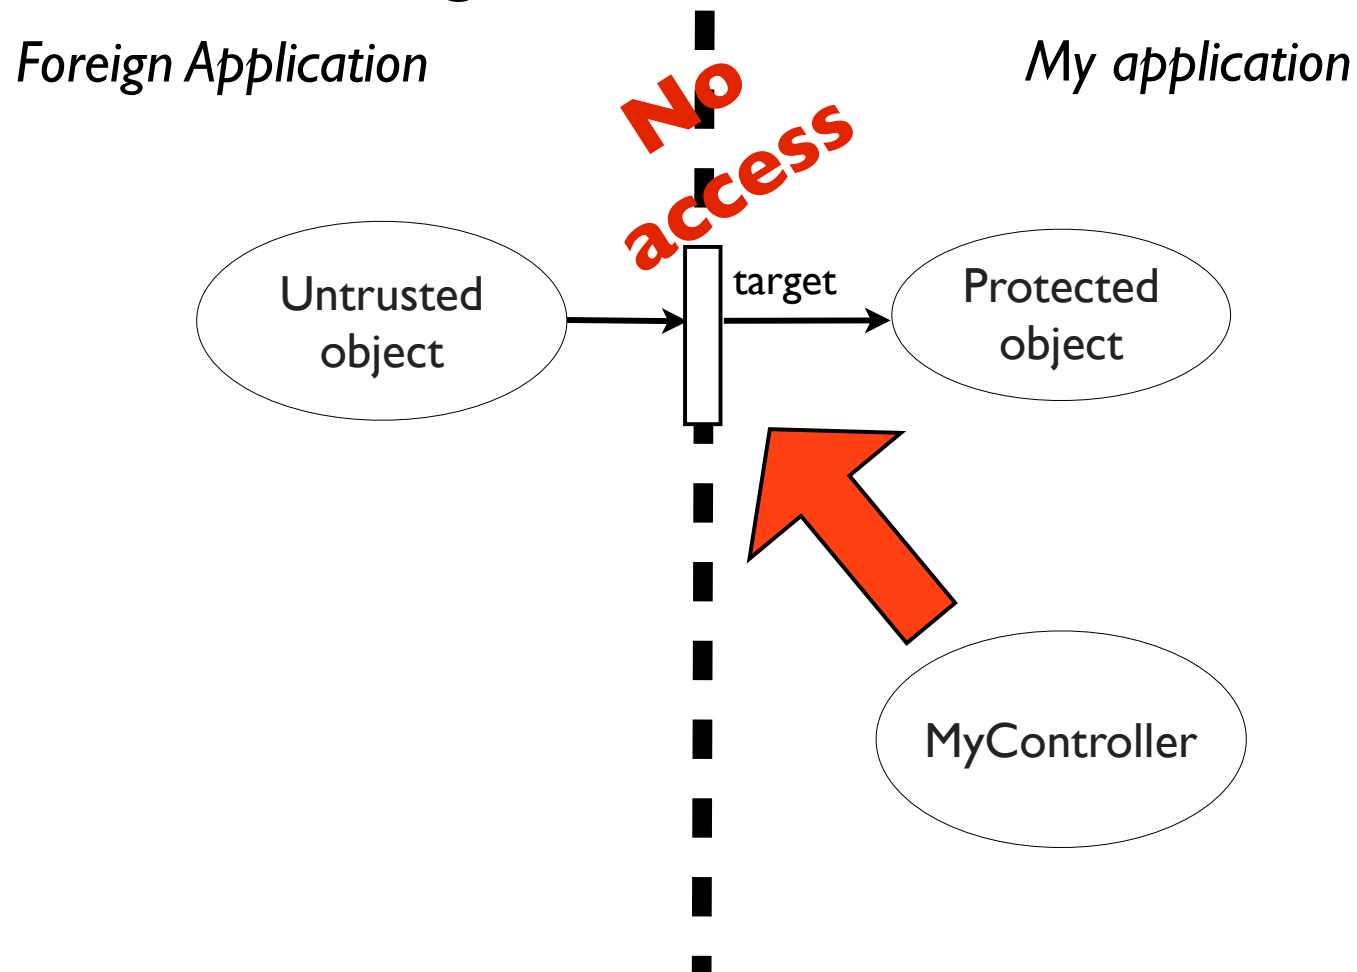
# Handle: control ?

If we want to change the semantics:



# Handle: control ?

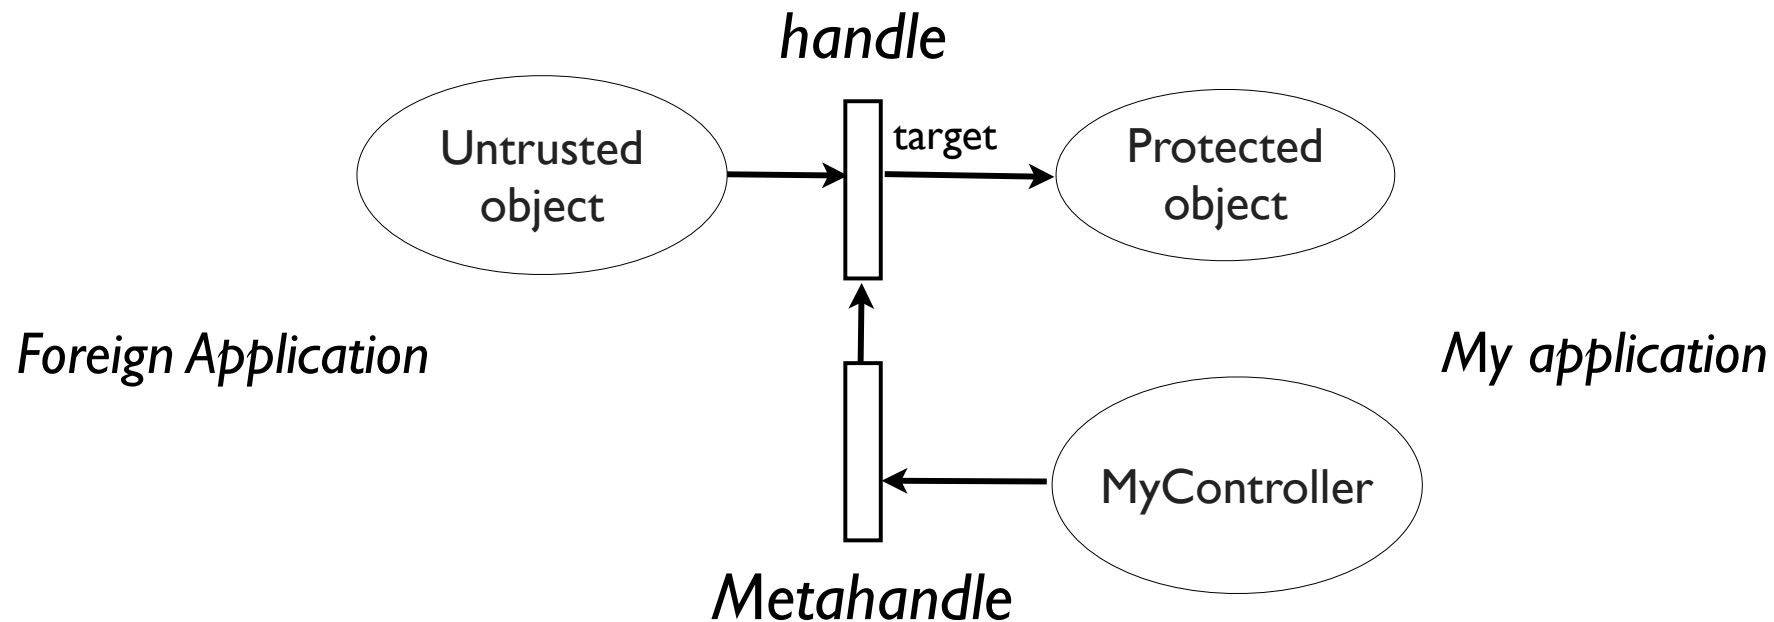
If we want to change the semantics:



# Handle: Metahandle

## Metahandle

- Handle on Handle
- Metahandle can be created during initialization phase
- Enable message passing to activated handle



# Handle: Reflection

Reflection:

Possibility to examine, change and execute behavior of objects at runtime

- Using reflection for meta-programming
- Using reflection for debugging

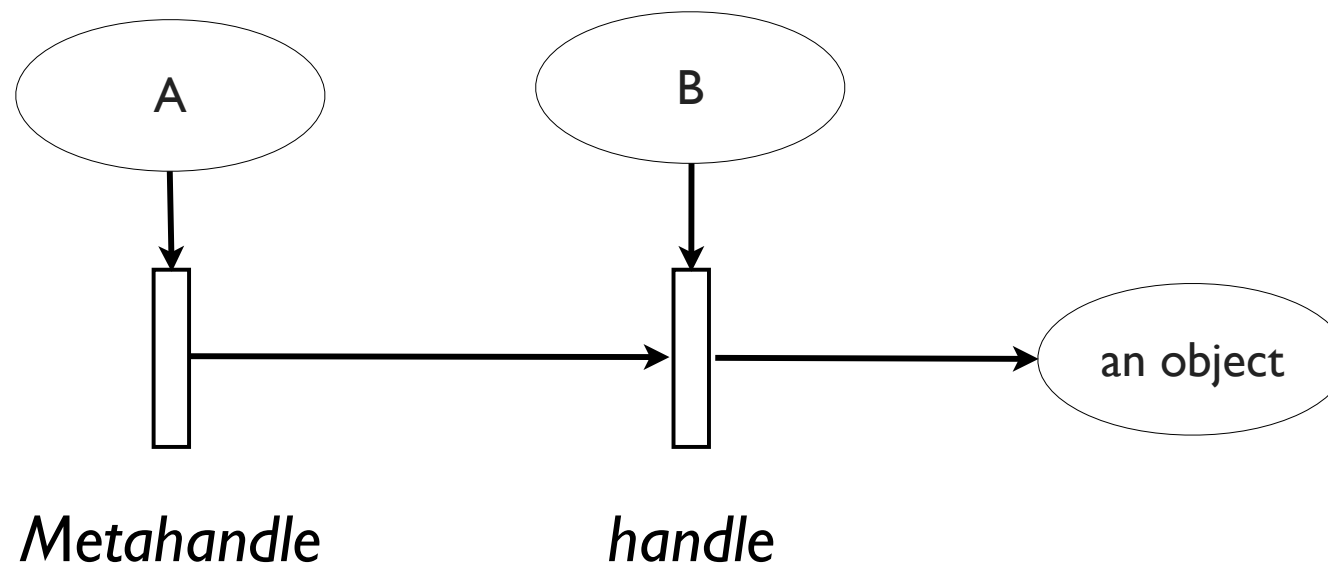


# Handle: Reflection

Reflection:

Possibility to examine, change and execute behavior of objects at runtime

- Using reflection for meta-programming
- Using reflection for debugging



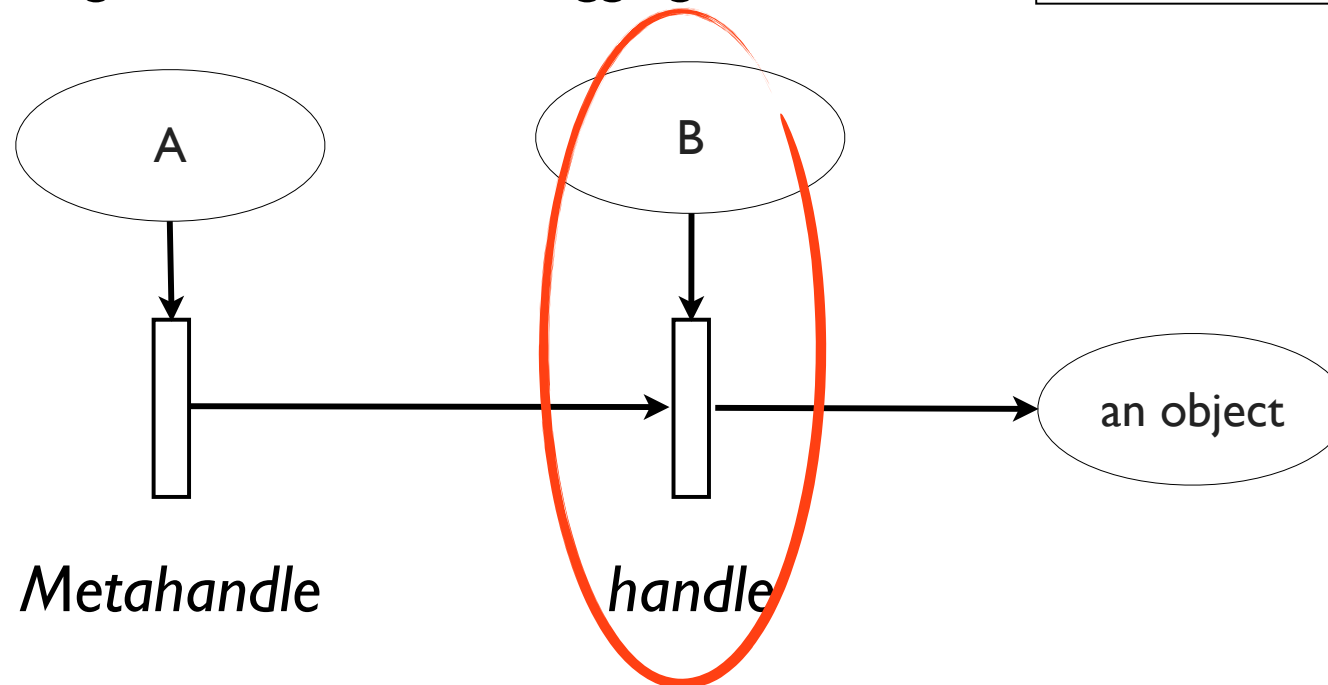
# Handle: Reflection

Reflection:

Possibility to examine, change and execute behavior of objects at runtime

- Using reflection for meta-programming
- Using reflection for debugging

Apply reflective  
API to target object



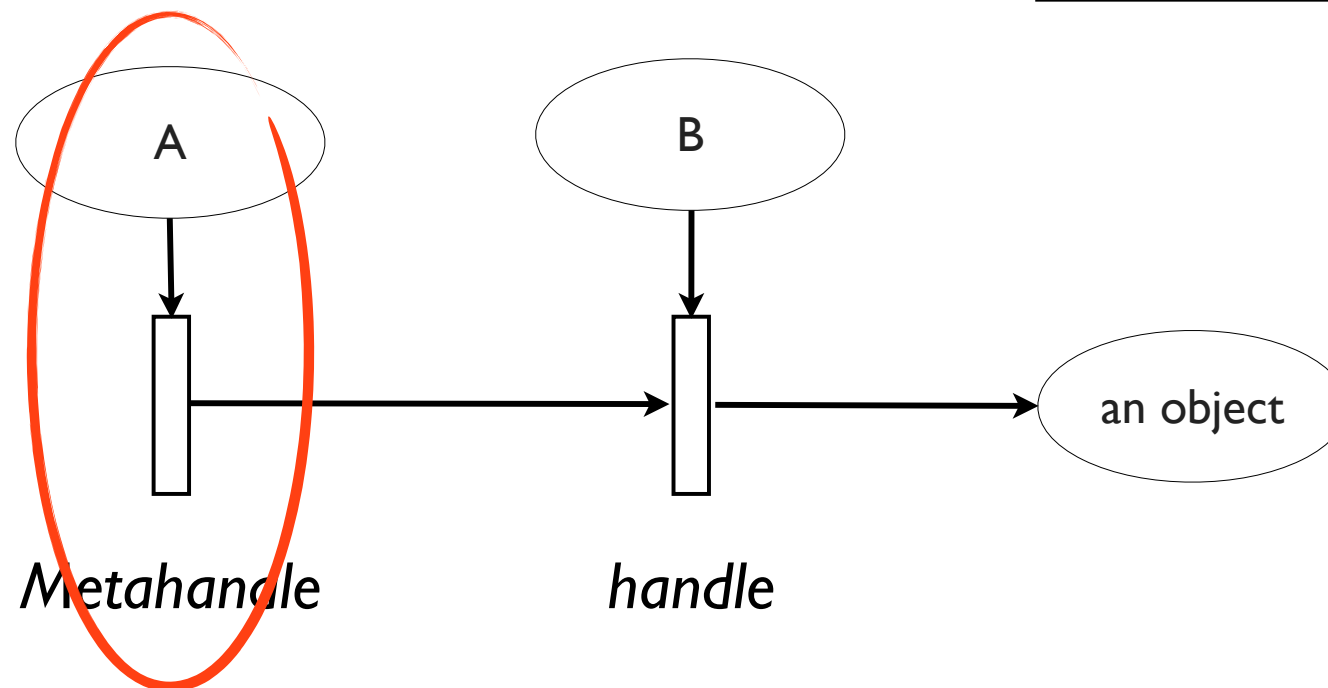
# Handle: Reflection

Reflection:

Possibility to examine, change and execute behavior of objects at runtime

- Using reflection for meta-programming
- Using reflection for debugging

Apply reflective  
API to target object



## The basics

- ✓ Reifying references
- ✓ Controlling behavior
- ✓ Isolating state

## Consequences

- ✓ Identity is preserved
- ✓ Propagated semantics
- ✓ Real time Control
- ✓ Reflection

# Outline

I. Motivation

II. Approaches

III. Solution

IV. Consequences

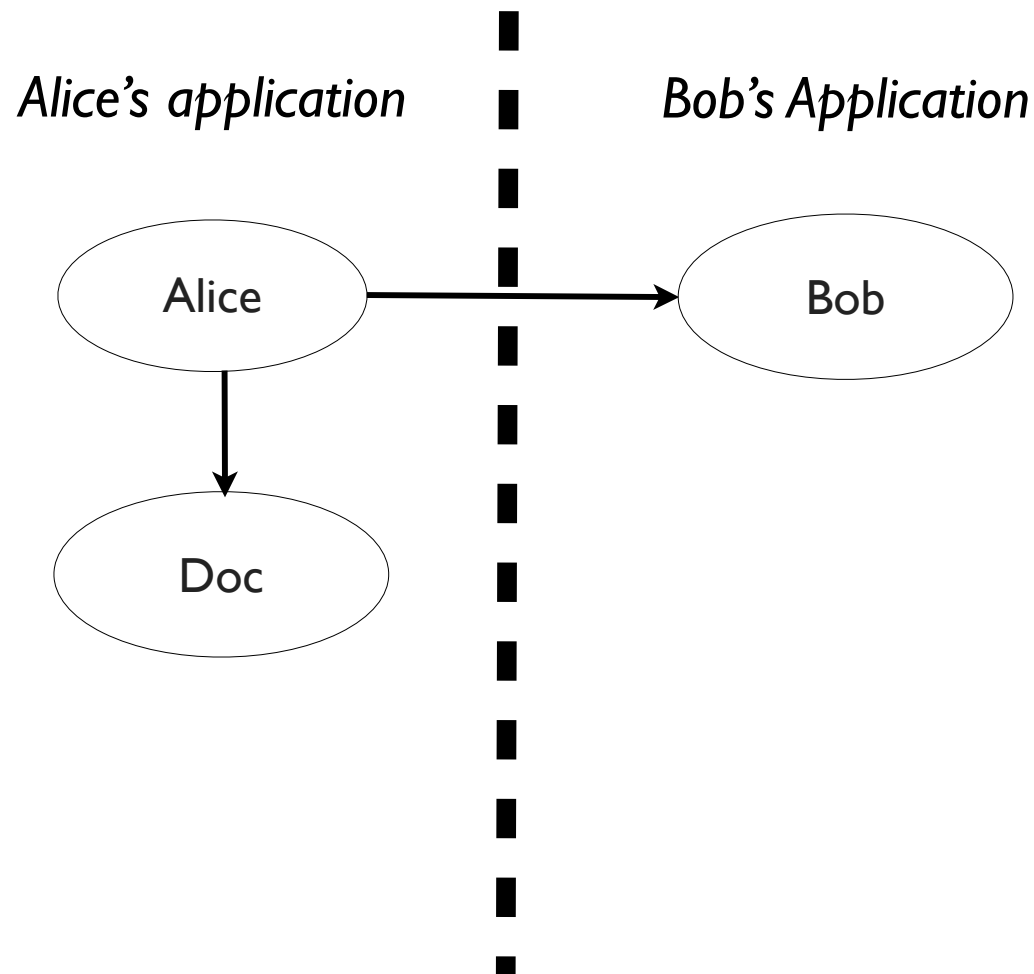
**V. Usages**

VI. Validation

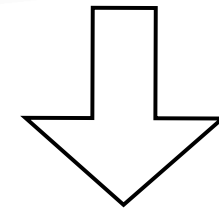
VII. Conclusion

# Handles at work

## Revocable references



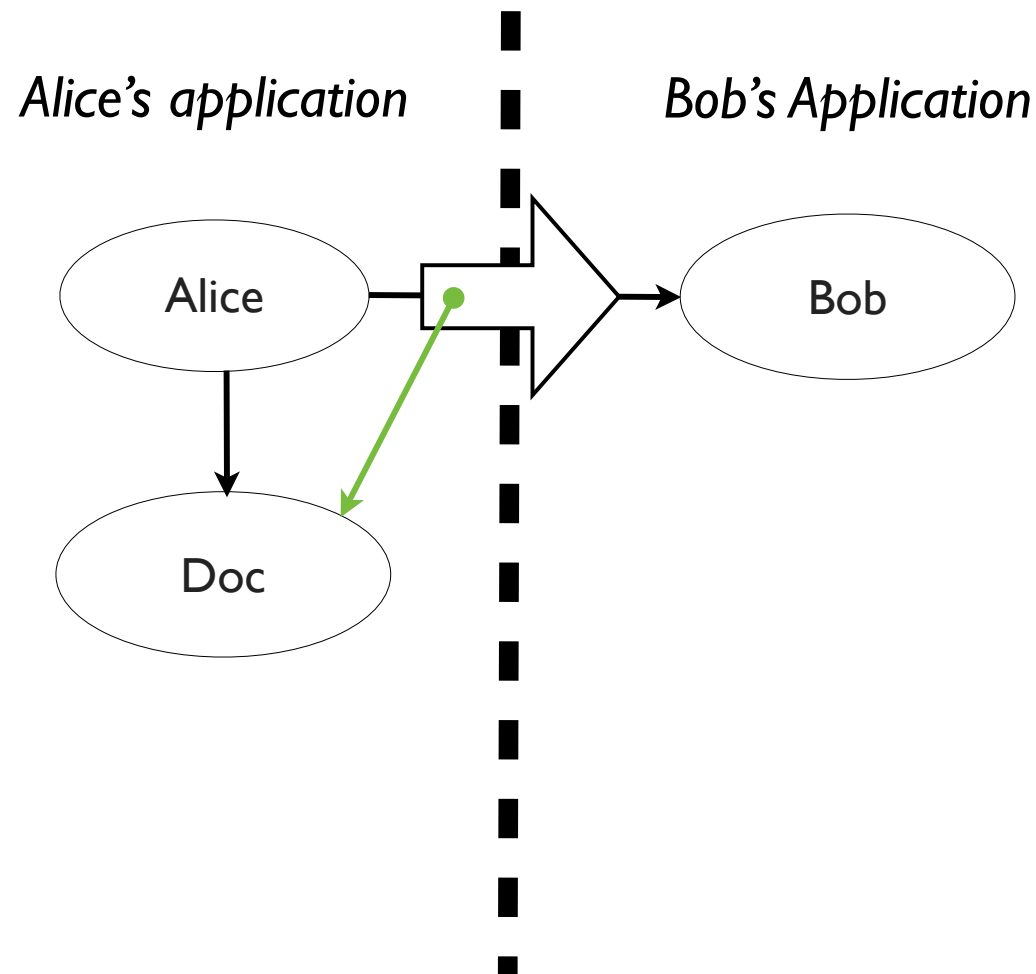
Redell  
1974



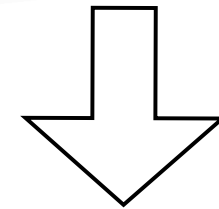
Tom Van  
Cutsem et al.  
DLS'2010

# Handles at work

## Revocable references



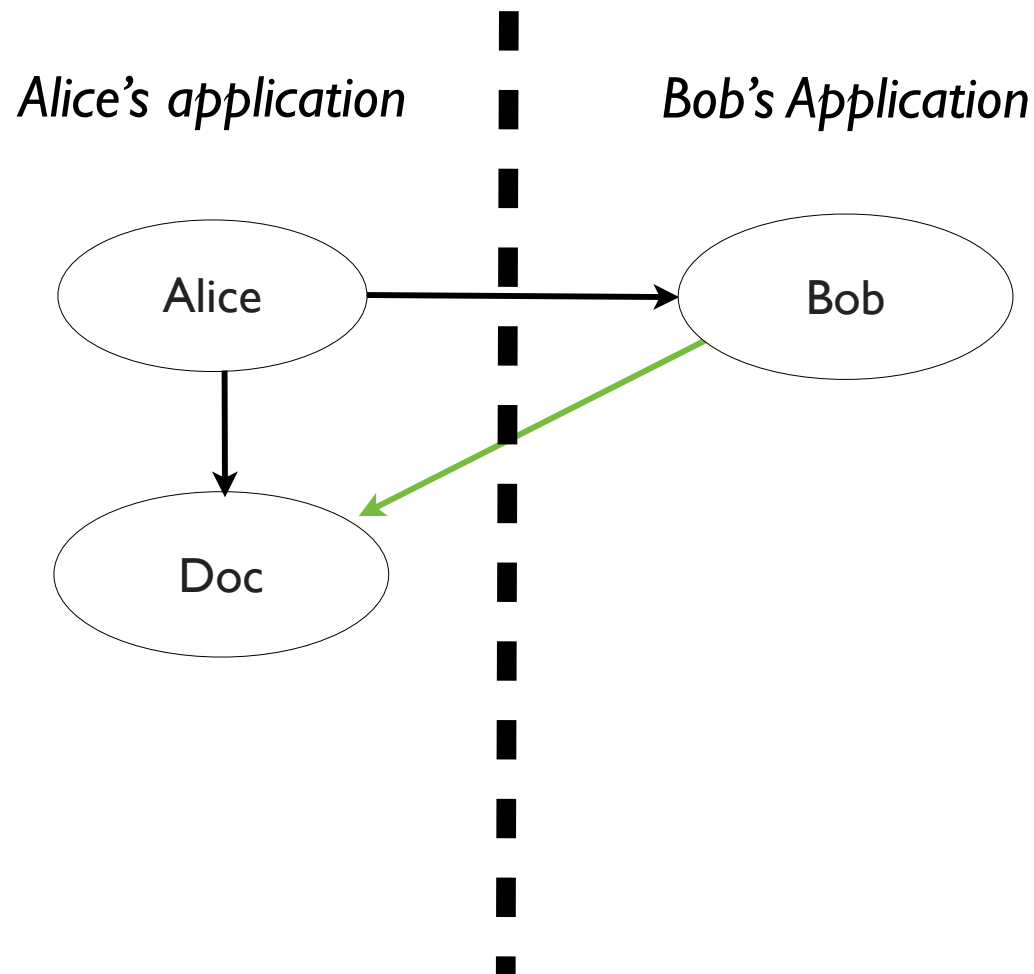
Redell  
1974



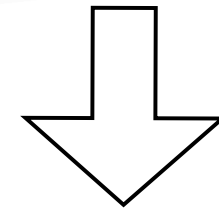
Tom Van  
Cutsem et al.  
DLS'2010

# Handles at work

## Revocable references



Redell  
1974

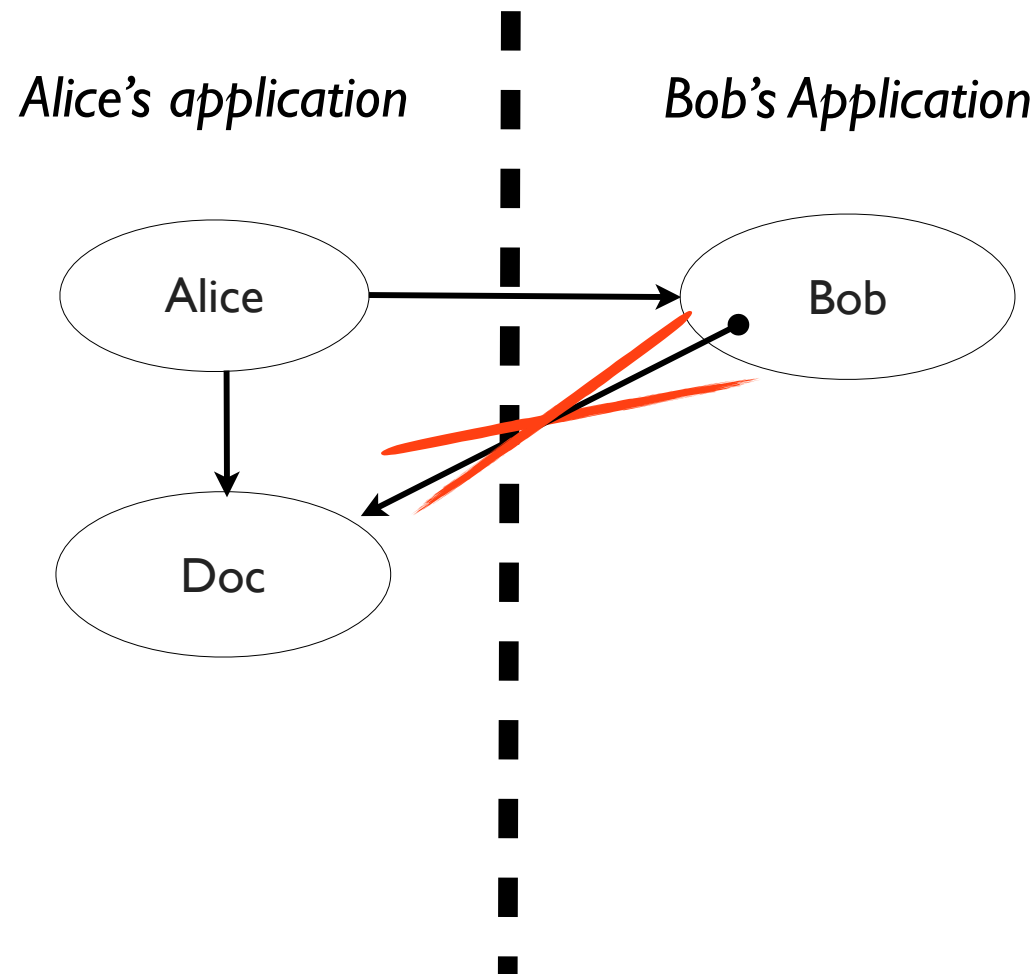


Tom Van  
Cutsem et al.  
DLS'2010

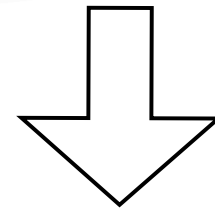


# Handles at work

## Revocable references



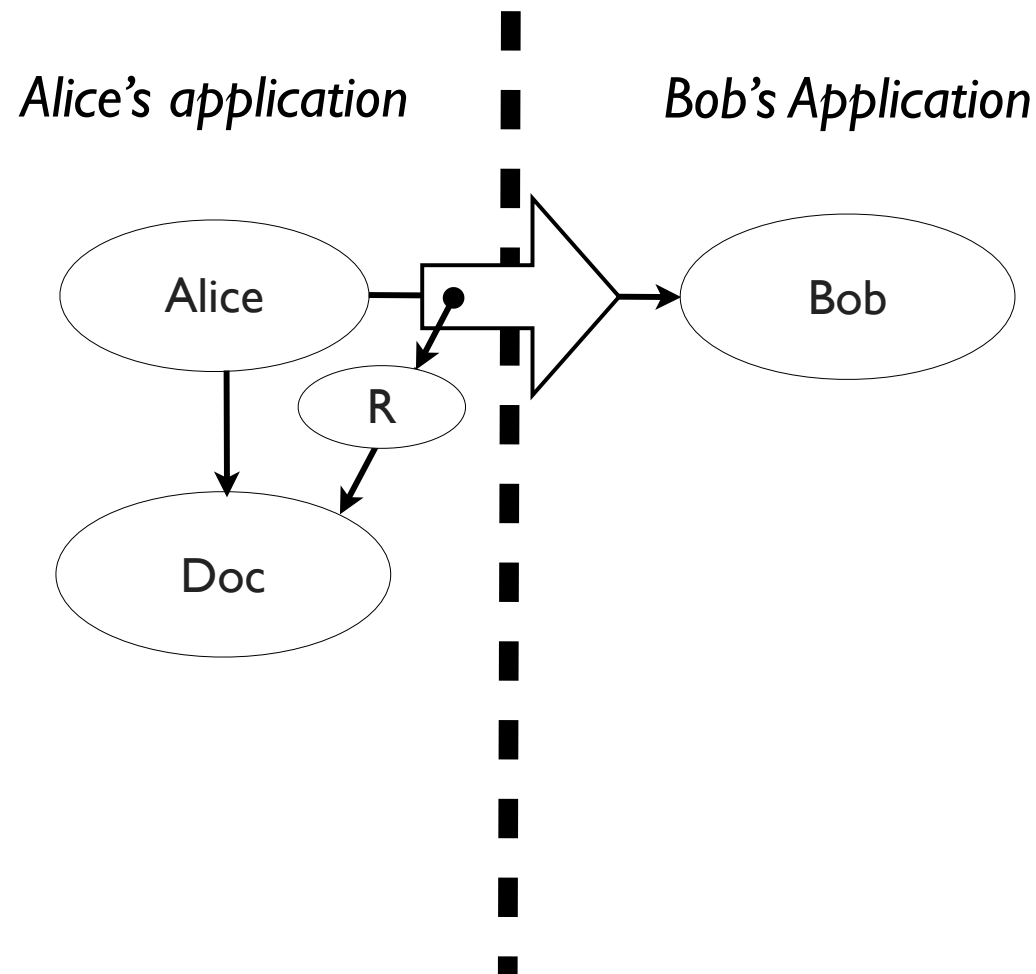
Redell  
1974



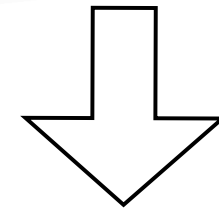
Tom Van  
Cutsem et al.  
DLS'2010

# Handles at work

## Revocable references



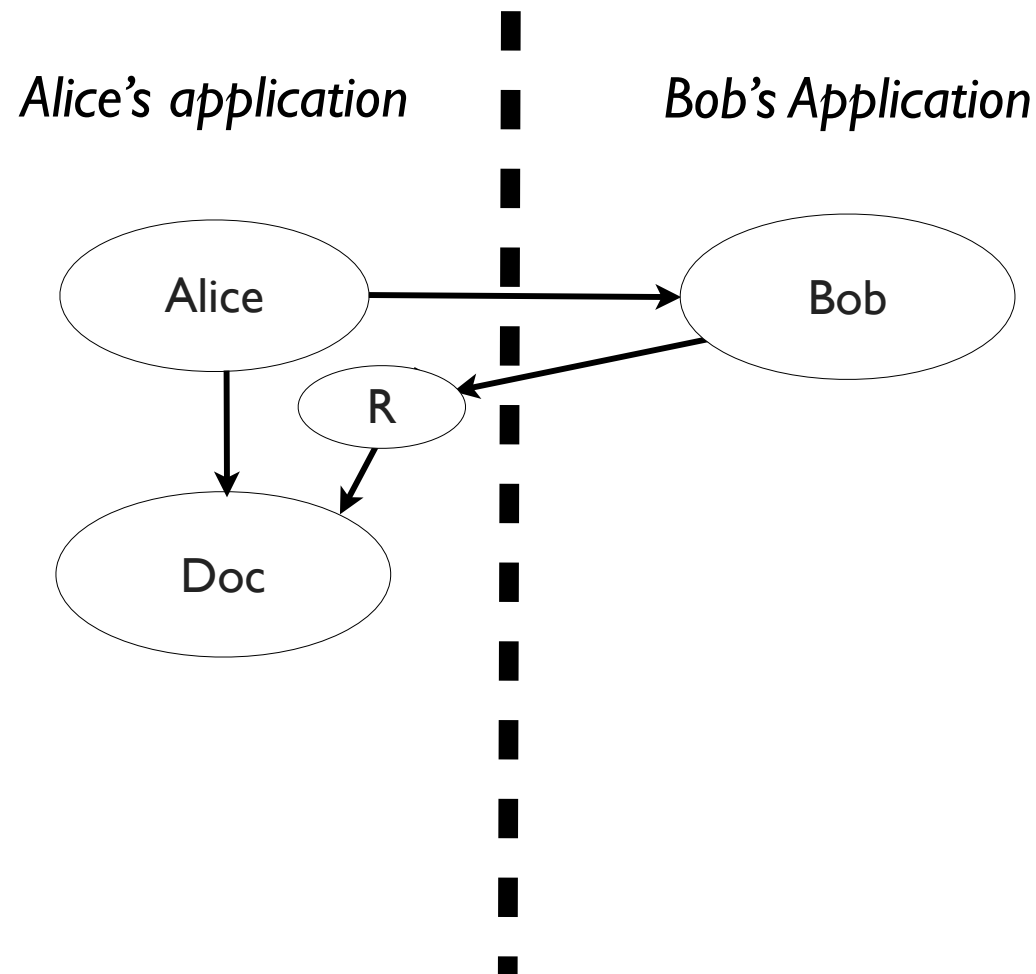
Redell  
1974



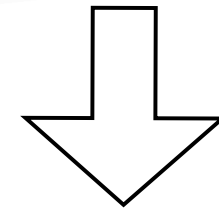
Tom Van  
Cutsem et al.  
DLS'2010

# Handles at work

## Revocable references



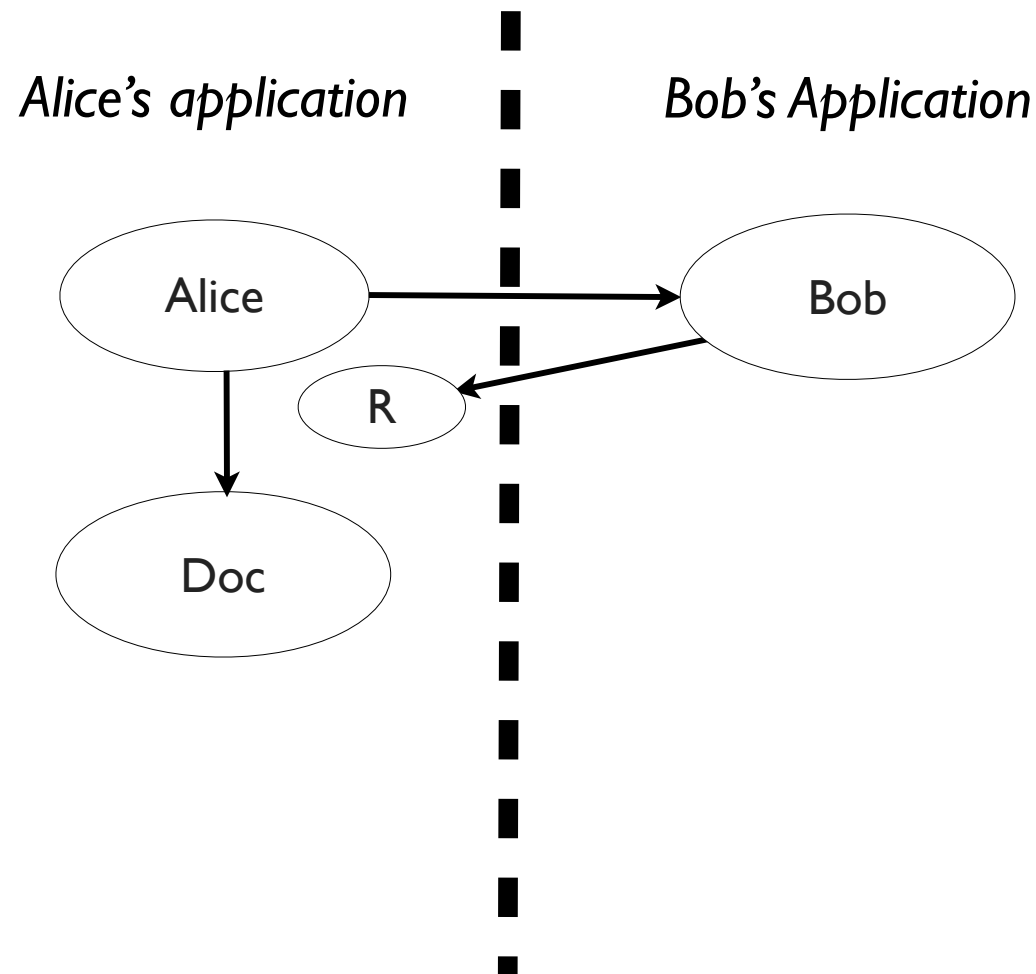
Redell  
1974



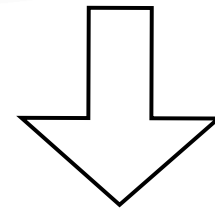
Tom Van  
Cutsem et al.  
DLS'2010

# Handles at work

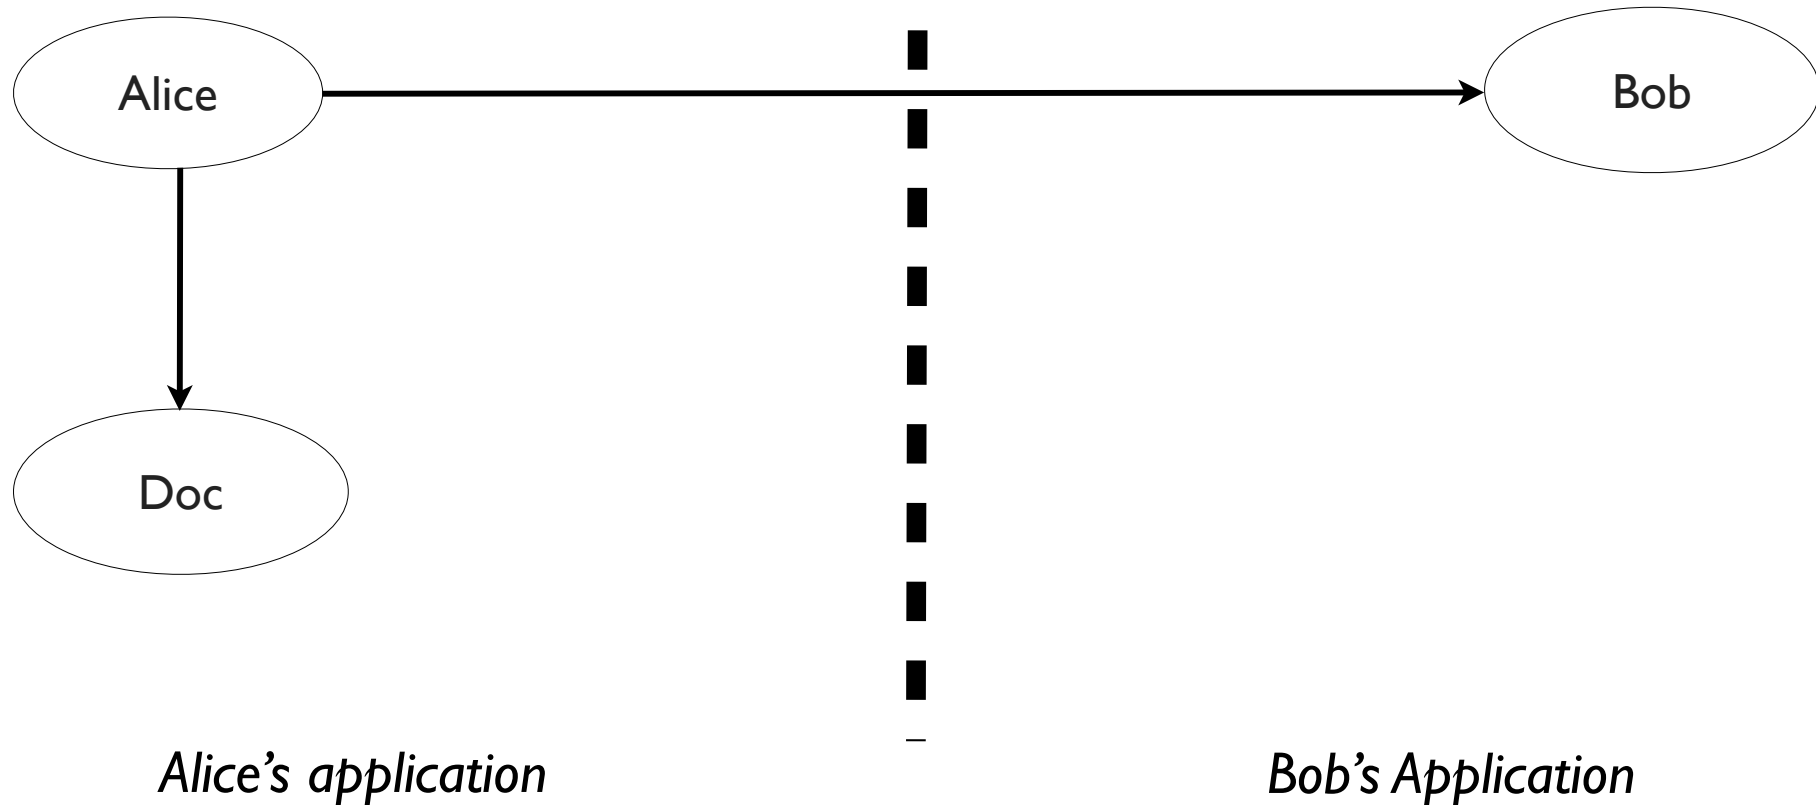
## Revocable references



Redell  
1974



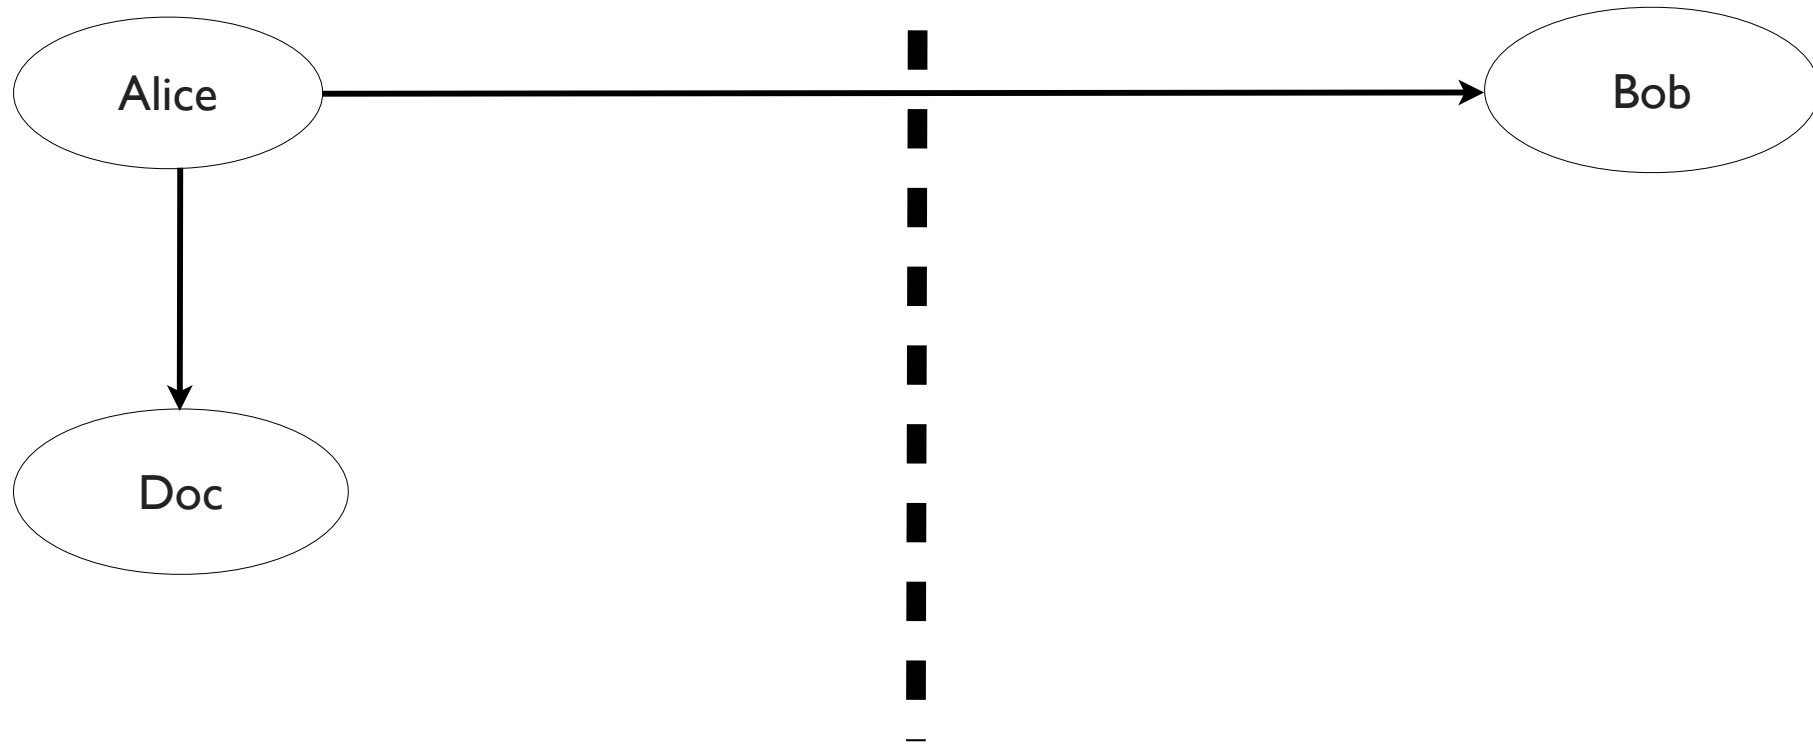
Tom Van  
Cutsem et al.  
DLS'2010



## Specific semantics:

All messages received by doc  
should raise an exception

**RevokingBehavior**  
doesNotUnderstand

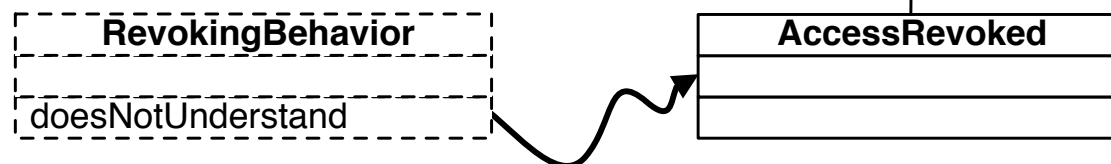


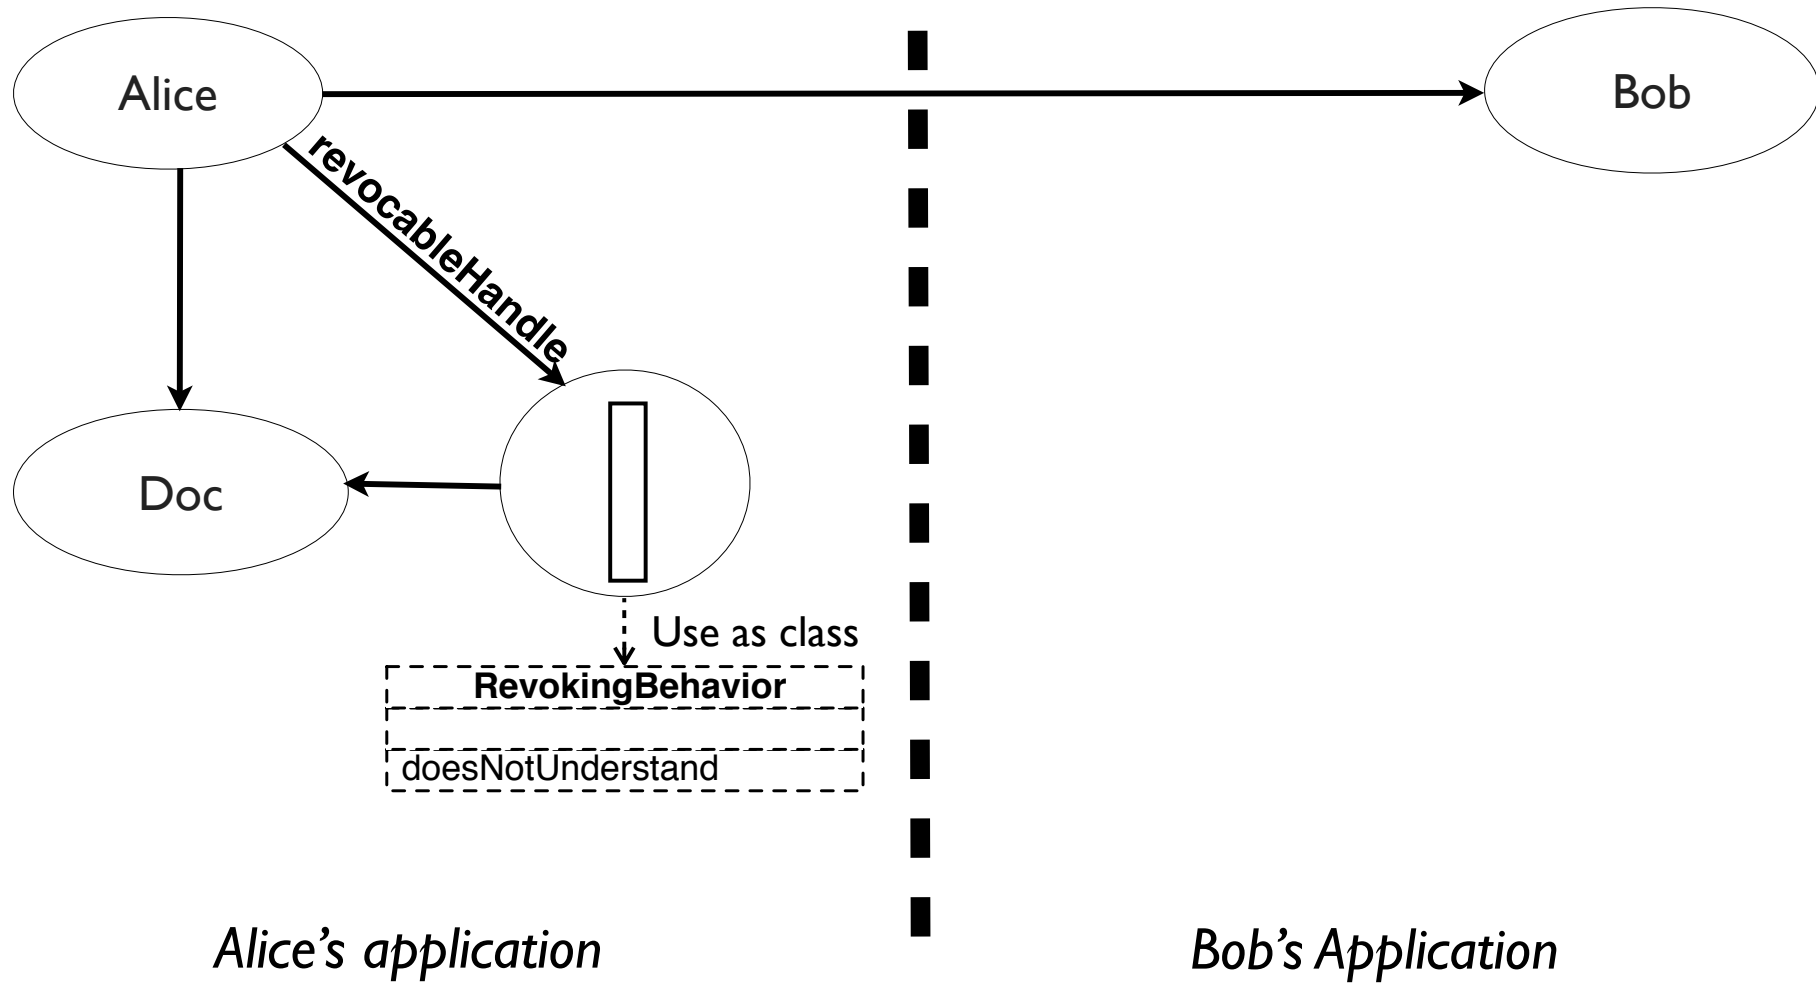
*Alice's application*

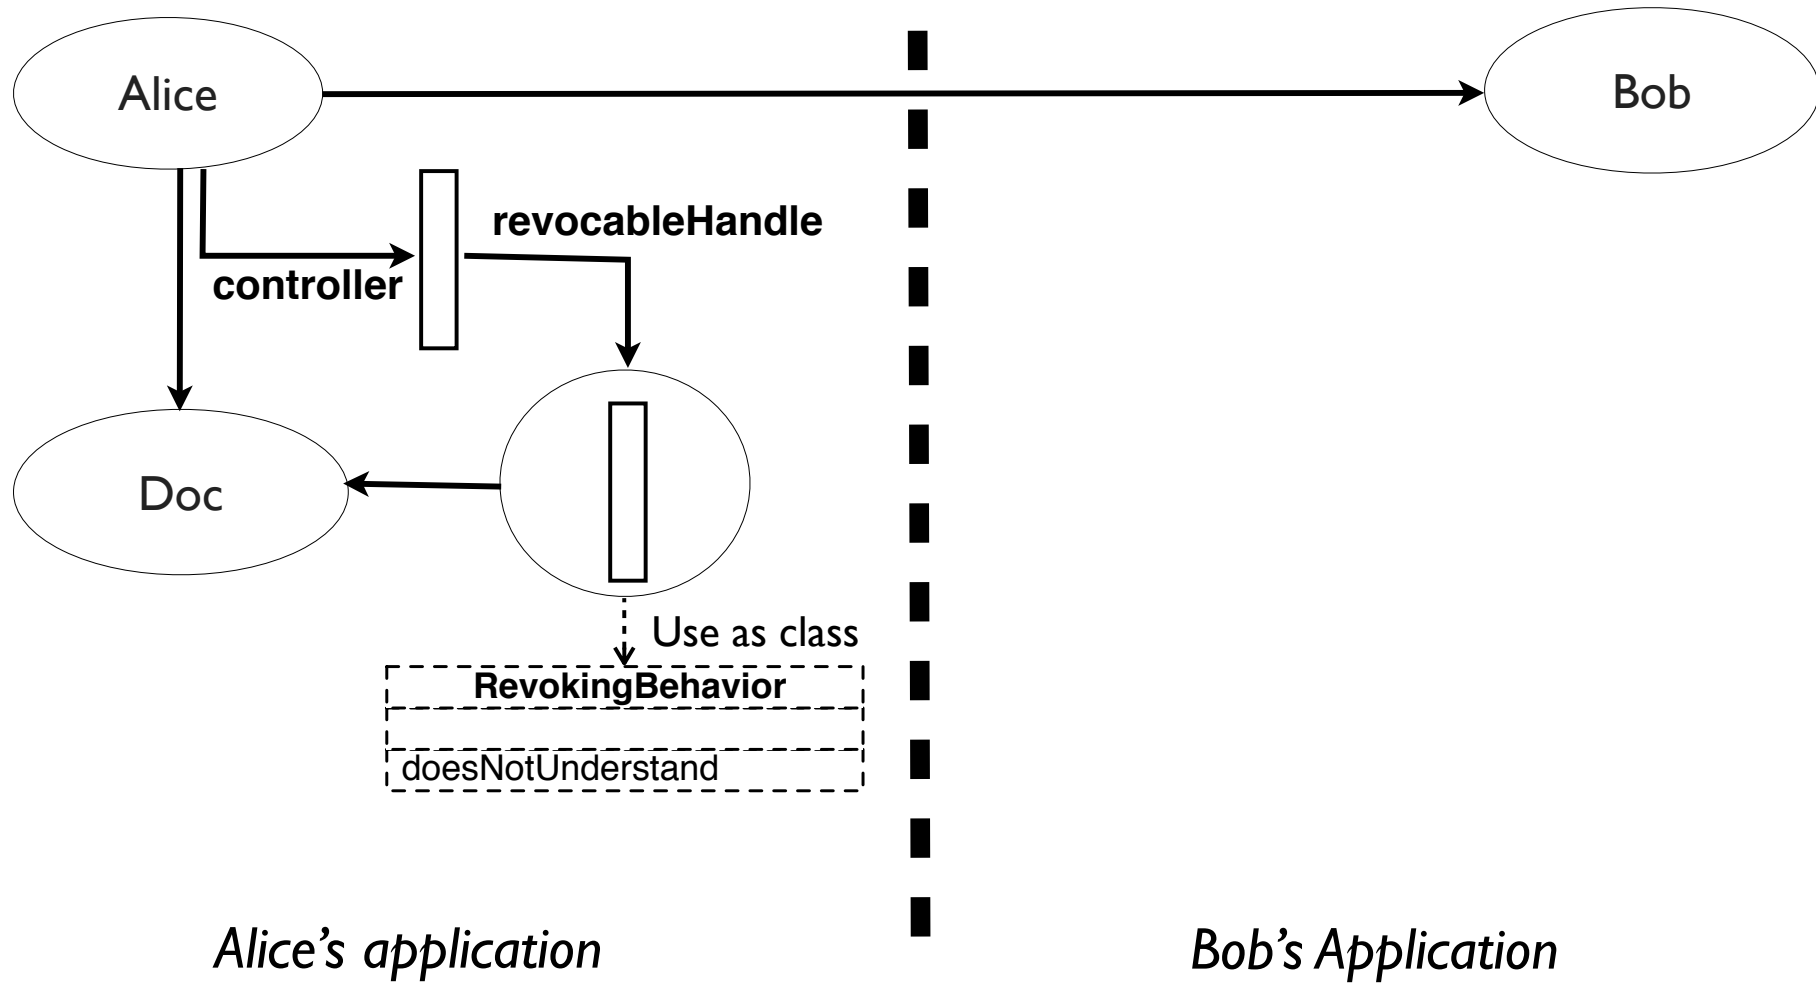
*Bob's Application*

## Specific semantics:

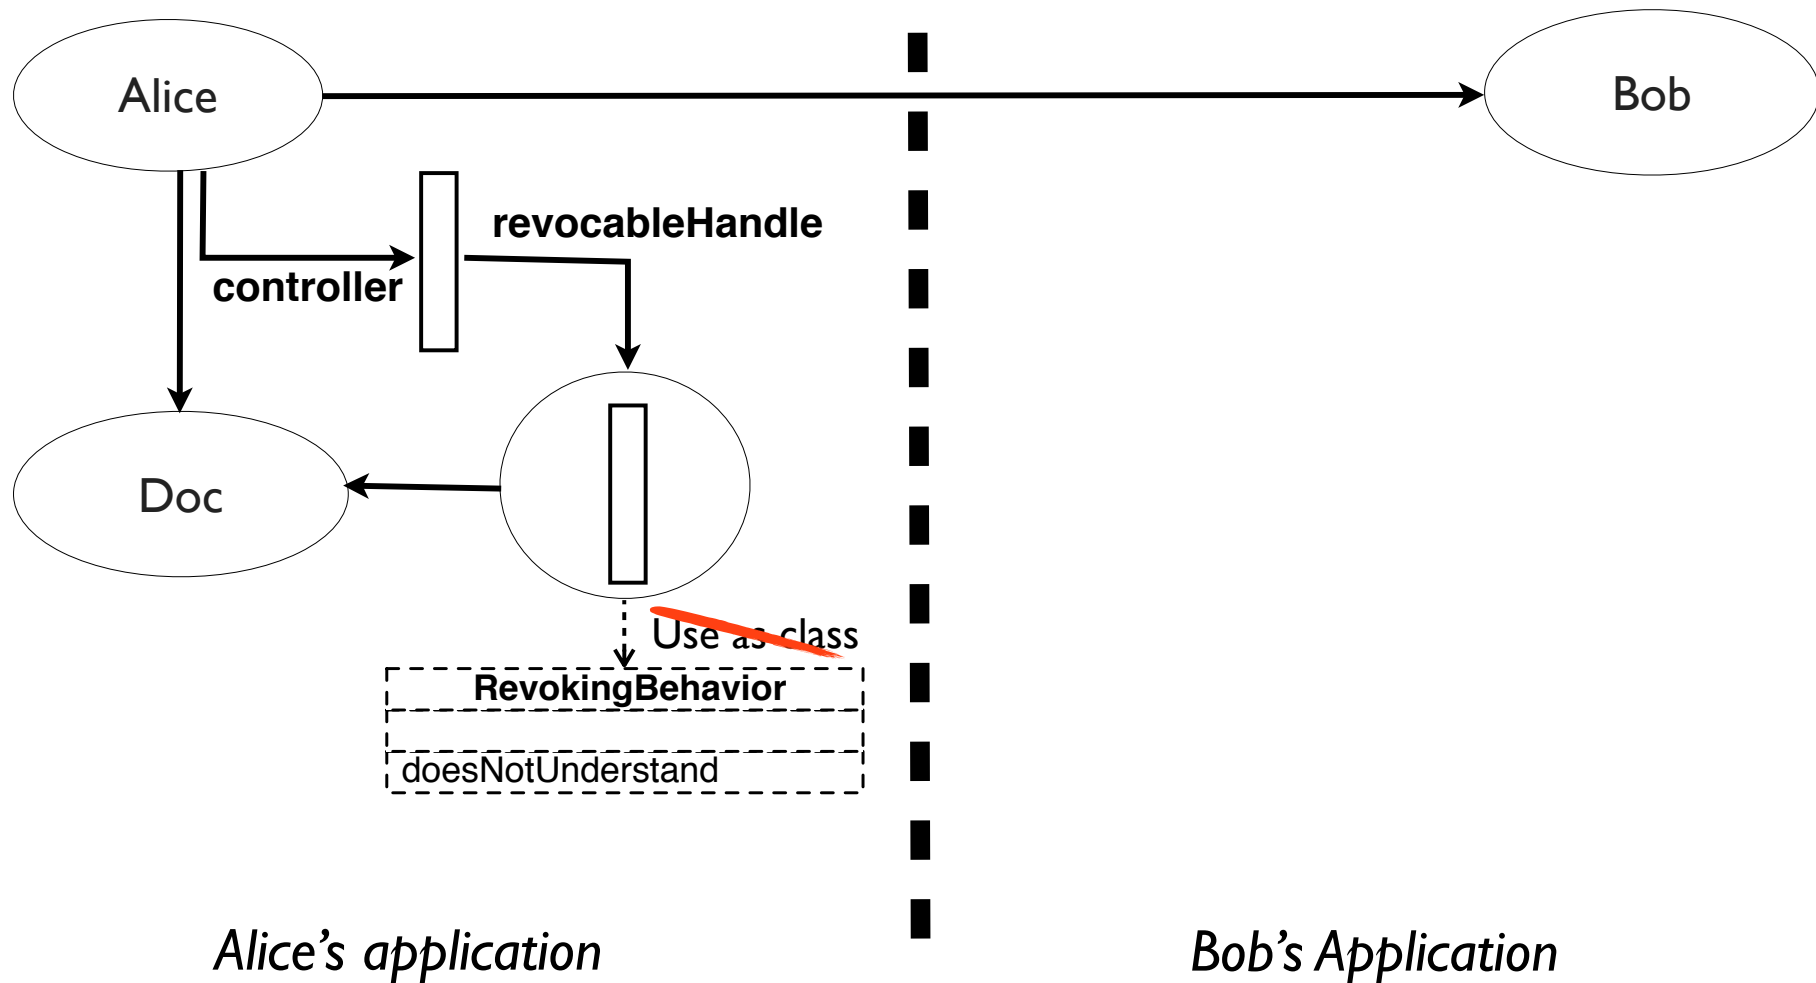
All messages received by doc should raise an exception

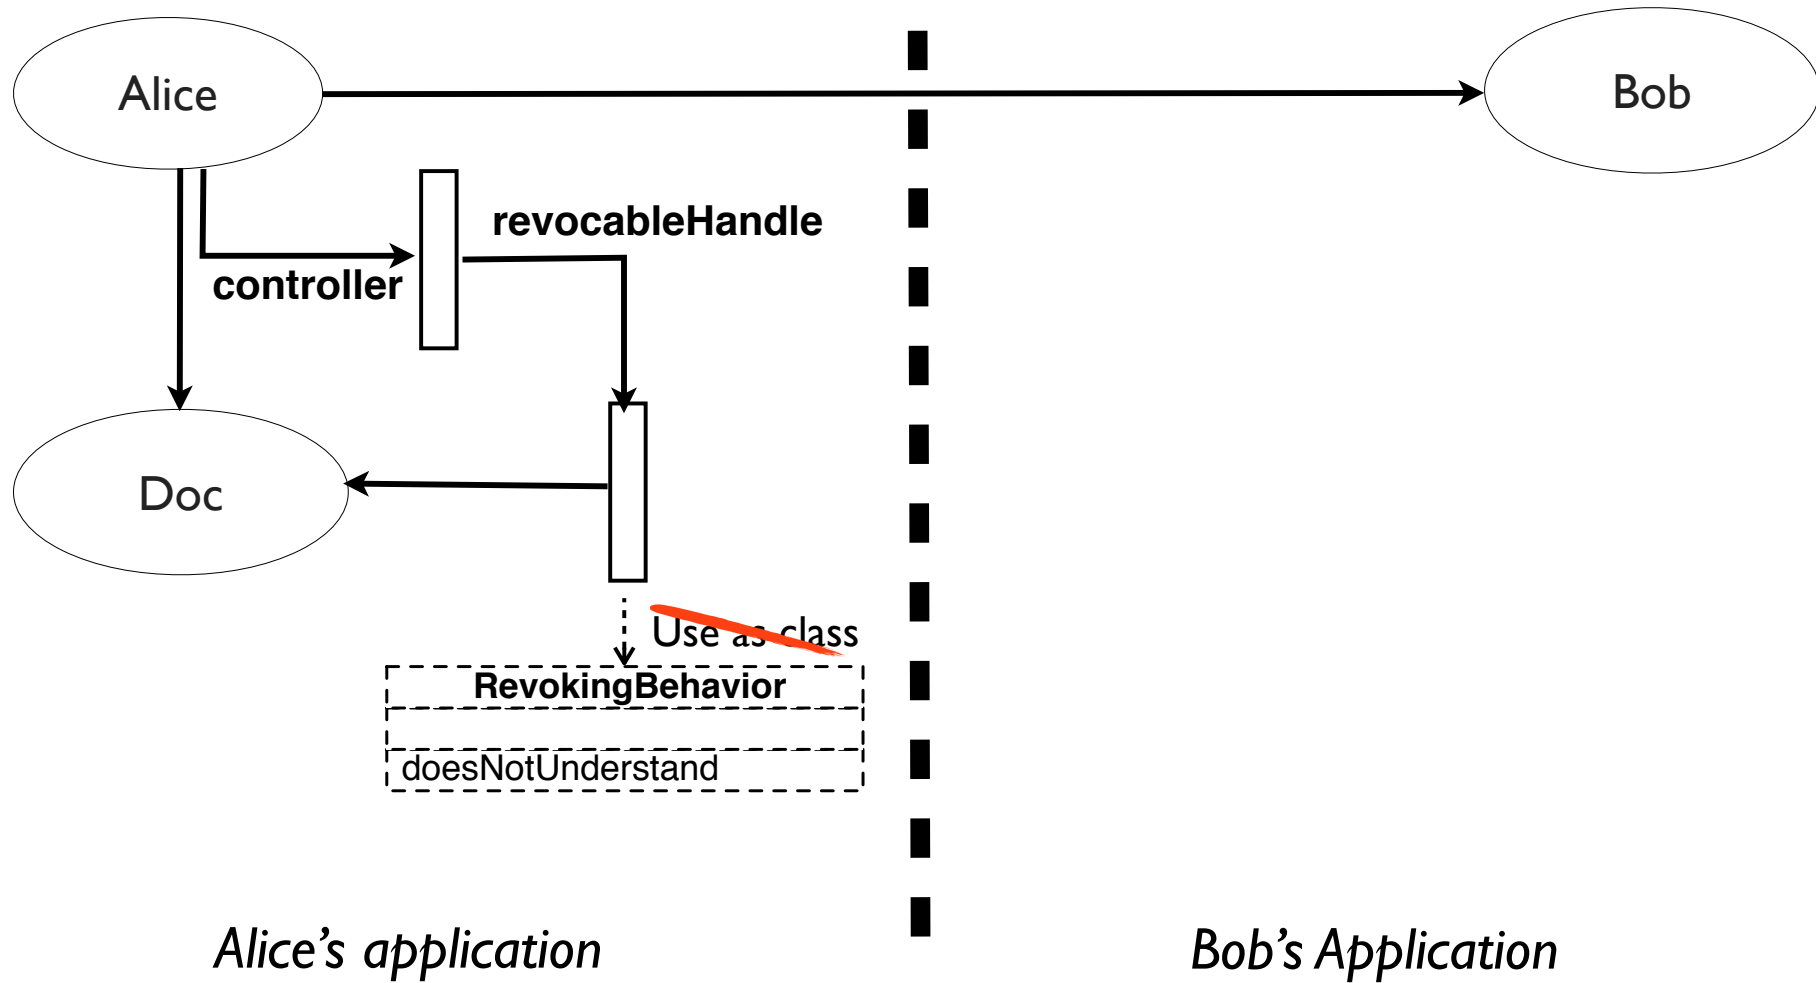


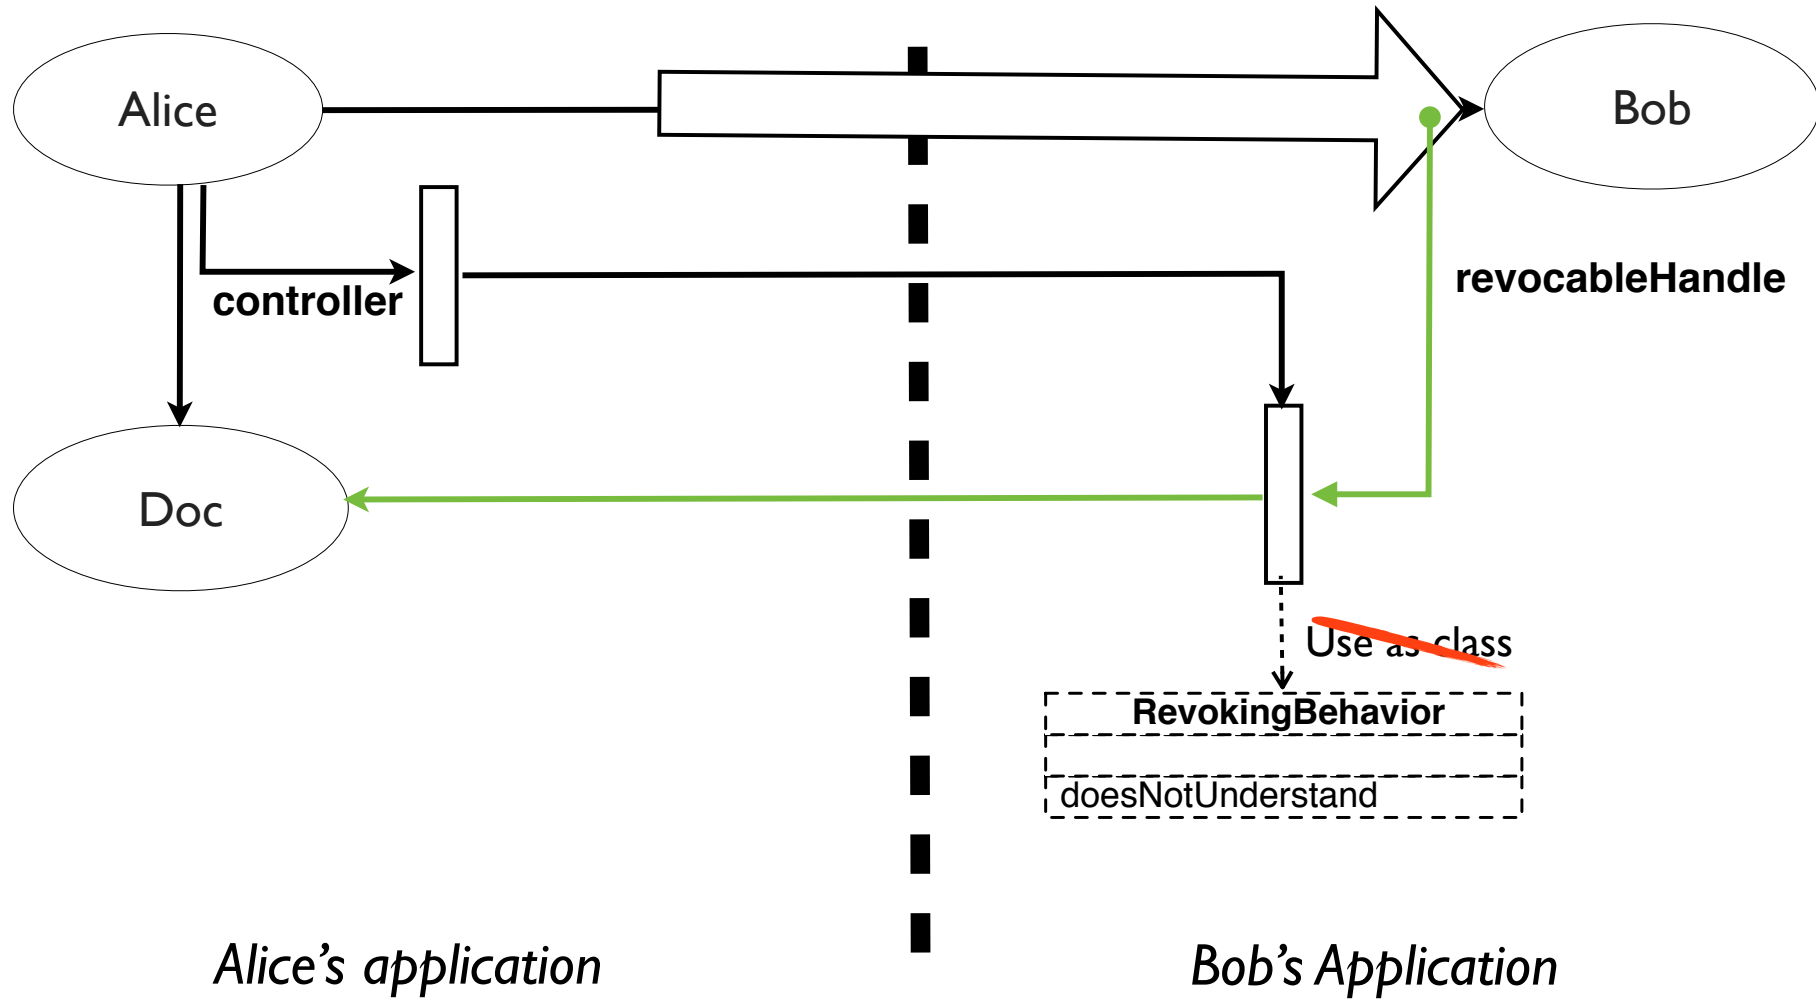


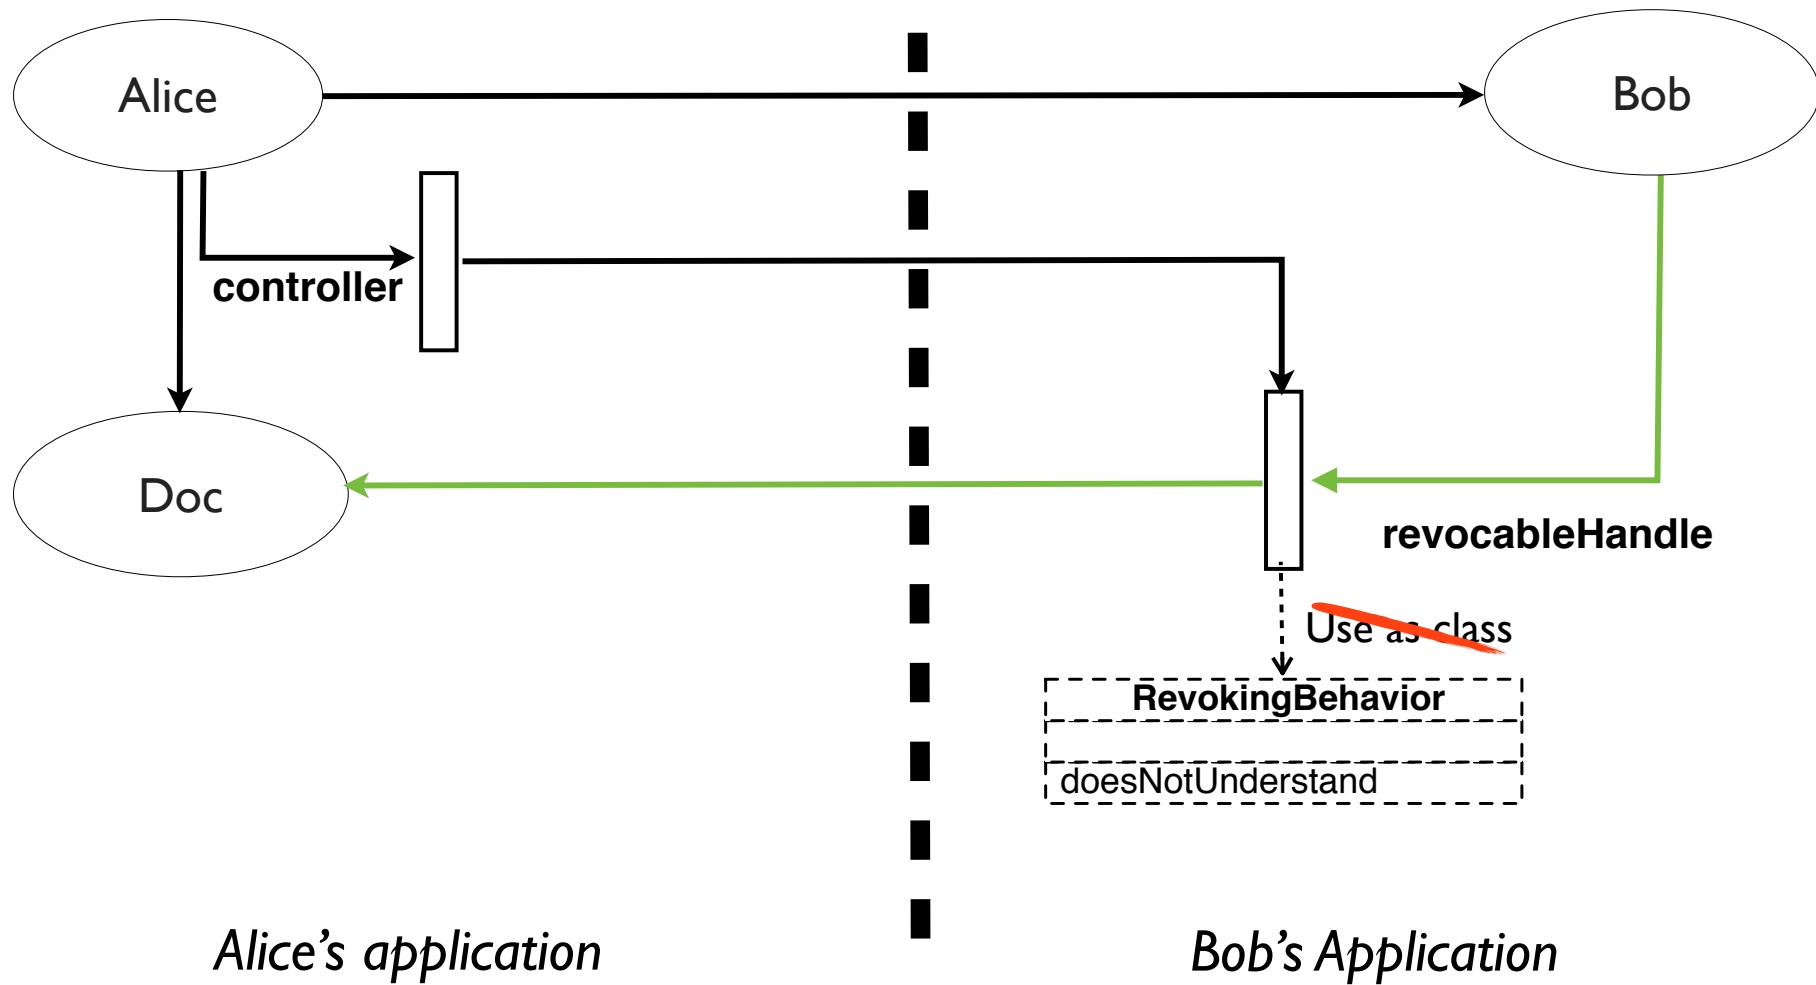


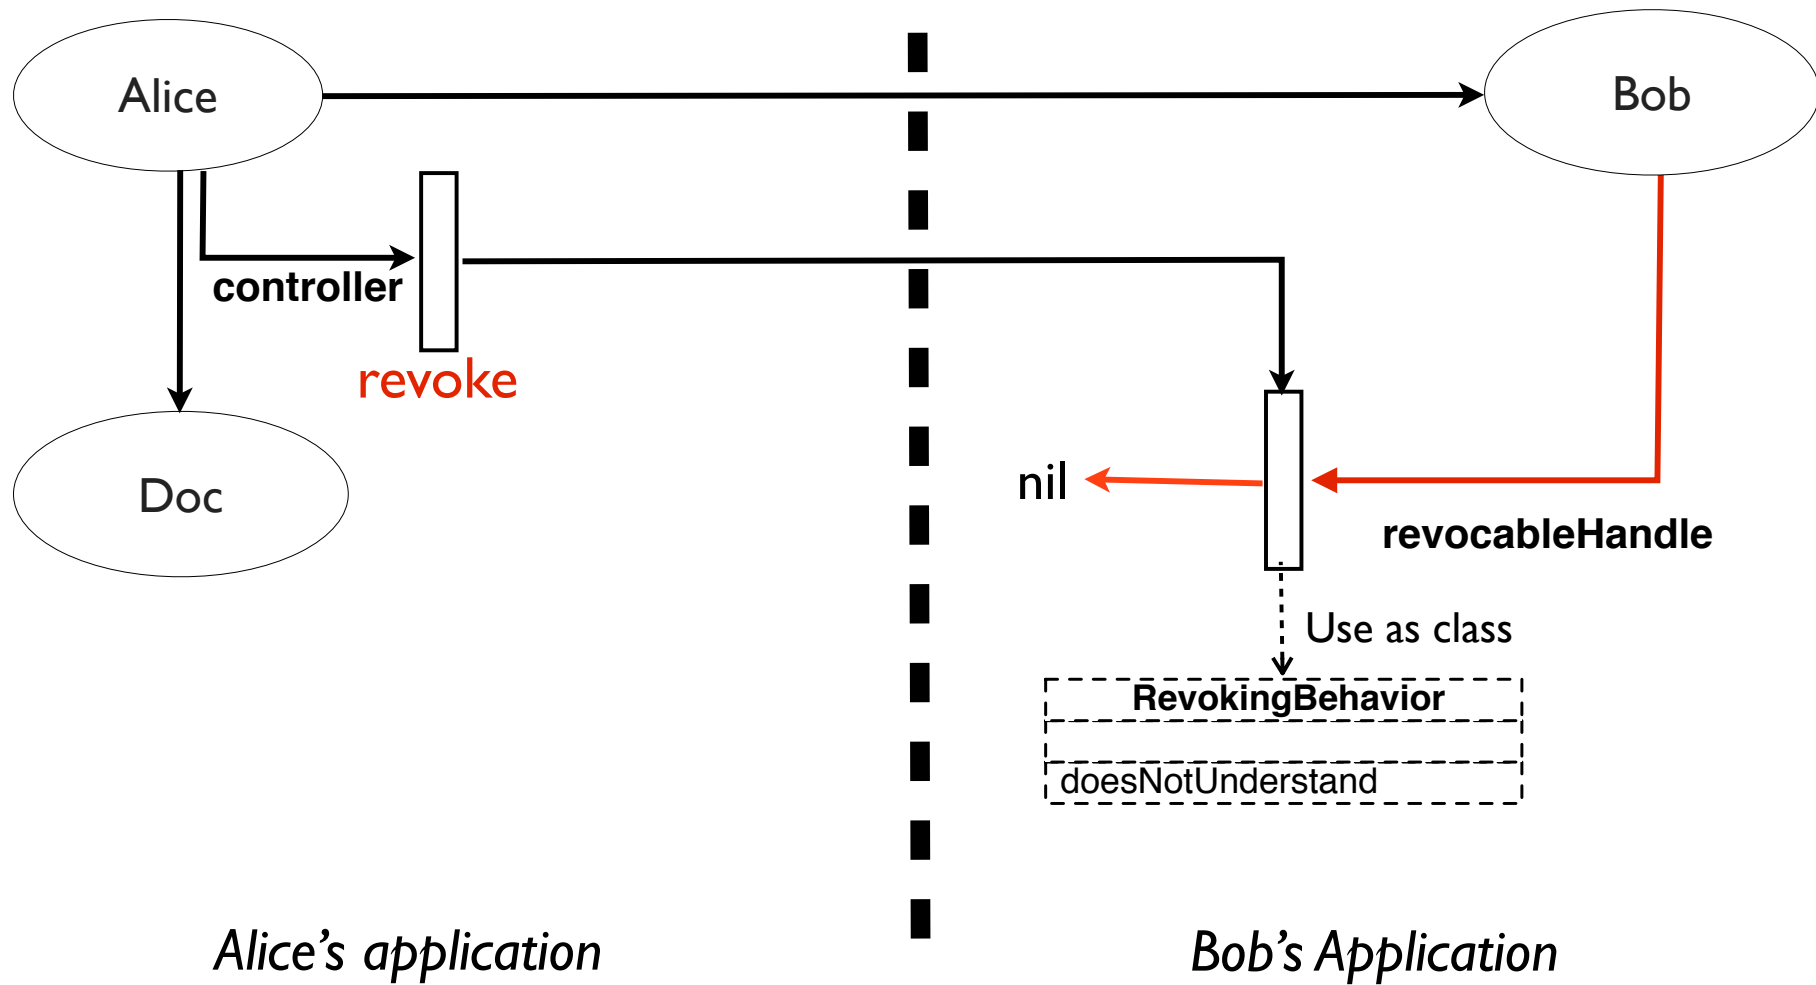












I. Motivation

II. Approaches

III. Solution

IV. Consequences

V. Usages

**VI. Validation**

VII. Conclusion

# Validation I

We formally described the Handle model:

Using the Felleisen and Hieb's notation

define properties at formal level

$$p(c) = c'$$

Where  $c'$  respects the properties  $p$ .

And  $c$  is a class.

Add handle to redex

$$\varepsilon = [\dots] \mid h_p^o$$

Handle send

$$P \vdash \langle E[h_p^o.m(v^*)], \mathcal{S} \rangle \hookrightarrow \langle E[h_p^o[e[v^*/x^*]]_{c''}], \mathcal{S} \rangle \quad [\textit{handled send}]$$

Where  $\langle c', m, x^*, e \rangle \in_P^* c''$

And  $c'$  is a class supporting the property  $p$  (via  $p(c) = c'$ )

And  $c$  is the class of the object  $o$

# Validation 2

## Implementation of the Handle's model on Pharo



### Virtual machine:

- ✓ Send / State manage
- ✓ Propagated semantics
- ✓ Primitives

### Language:

- ✓ Language interfaces
- ✓ Semantics



# Validation 3

*N-bodies* benchmark: Stress the message send

	Means	Standard deviation
normal VM	4444.50ms	38.36ms
handle VM (Normal reference)	4772.80ms	20.68ms

7,36%

*Binaries tree* benchmark: Stress the state accesses

	Means	Standard deviation
normal VM	21167.00ms	106.26ms
handle VM (Normal reference)	22321.57ms	66.35ms

5,45%

# Validation 3

## *N-bodies* benchmark: Stress the message send

cost of the  
semantics

71%

	Means	Standard deviation
normal VM	4444.50ms	38.36ms
handle VM (Normal reference)	4772.80ms	20.68ms
revocable nbody	8172.12ms	31.01ms

## *Binaries tree* benchmark: Stress the state accesses

1,747,535 graphs between 4 and 65,536 nodes

cost of the  
semantics

205%

	Means	Standard deviation
normal VM	21167.00ms	106.26ms
handle VM (Normal reference)	22321.57ms	66.35ms
revocable Binaries tree	68094.23ms	70.06ms

# Validation 3

## Behavior semantics changes

- Read-Only executions
- Revocable references

## State semantics changes

- Worlds
- Software transactional memory

# Validation 4

## Behavior semantics changes

- Read-Only executions
- Revocable references

## State semantics changes

- Worlds
- Software transactional memory

### Implemented with:

- 2 Classes
- 14 methods
- Less than 50 lines of codes


# Validation 4

## Behavior semantics changes

- Read-Only executions
- **Revocable references**

## State semantics changes

- Worlds
- Software transactional memory



Tom Van  
Cutsem et al.  
DLS'2010

## Implemented with:

- 5 Classes
- 19 methods
- Less than 100 lines of codes

# Validation 4

## Behavior semantics changes

- Read-Only executions
- Revocable references

## State semantics changes

- **Worlds**
- Software transactional memory



Warth et al.  
ECOOP'11

## Implemented with:

- 2 Classes
- 16 methods
- Less than 50 lines of codes


# Validation 4

## Behavior semantics changes

- Read-Only executions
- Revocable references

## State semantics changes

- Worlds
- **Software transactional memory**



Shavit et al.  
PODC'1995

## Implemented with:

- 4 Classes
- 35 methods
- Less than 200 lines of codes



# Publications

## International conference

Read-Only Execution for Dynamic Languages. In Proceedings of the 48th International Conference Objects, Models, Components, Patterns (TOOLS'10), Malaga, Spain, June 2010.

## Journal paper

Behavior-Propagating First Class References For Dynamically-Typed Languages. In Science of Computer Programming, 2013. **Under-revision.**



# Conclusion

## Our Solution

- ✓ Reifying references
- ✓ Controlling behavior
- ✓ Isolating state
- ✓ Identity is preserved
- ✓ Propagated semantics
- ✓ Real time Control
- ✓ Reflection

## Validation

- ✓ Operational semantic
- ✓ Implementation
- ✓ Acceptable performance
- ✓ Semantics

## Future work:

- Using Handle on a real world application
- Composition model (ReadOnly + Revocable)
- State control (add instance variables, deltas, etc.)

# Towards First Class References as a Security Infrastructure in Dynamically-Typed Languages

Jean-Baptiste Arnaud

Supervisor: Stéphane Ducasse  
Co-Supervisor: Marcus Denker  
RMoD team



February 18, 2013